Evaluating Robustness of Predictive Uncertainty Estimation: Are Dirichlet-based Models Reliable?

Anna-Kathrin Kopetzki * 1 Bertrand Charpentier * 1 Daniel Zügner 1 Sandhya Giri 1 Stephan Günnemann 1

Abstract

Dirichlet-based uncertainty (DBU) models are a recent and promising class of uncertainty-aware models. DBU models predict the parameters of a Dirichlet distribution to provide fast, high-quality uncertainty estimates alongside with class predictions. In this work, we present the first largescale, in-depth study of the robustness of DBU models under adversarial attacks. Our results suggest that uncertainty estimates of DBU models are not robust w.r.t. three important tasks: (1) indicating correctly and wrongly classified samples; (2) detecting adversarial examples; and (3) distinguishing between in-distribution (ID) and out-of-distribution (OOD) data. Additionally, we explore the first approaches to make DBU models more robust. While adversarial training has a minor effect, our median smoothing based approach significantly increases robustness of DBU models.

1. Introduction

Neural networks achieve high predictive accuracy in many tasks, but they are known to have two substantial weaknesses: First, neural networks are not robust against adversarial perturbations, i.e., semantically meaningless input changes that lead to wrong predictions (Szegedy et al., 2014; Goodfellow et al., 2015). Second, standard neural networks are unable to identify samples that are different from the ones they were trained on and tend to make over-confident predictions at test time (Lakshminarayanan et al., 2017). These weaknesses make them impracticable in sensitive domains like financial, autonomous driving or medical areas which require trust in predictions.

To increase trust in neural networks, models that provide pre-

Proceedings of the 38^{th} International Conference on Machine Learning, PMLR 139, 2021. Copyright 2021 by the author(s).

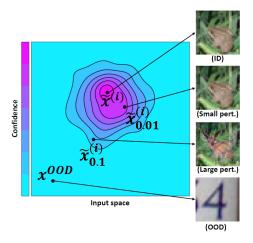


Figure 1. Visualization of the desired uncertainty estimates.

dictions along with the corresponding uncertainty have been proposed. An important, quickly growing family of models is the Dirichlet-based uncertainty (DBU) family (Malinin & Gales, 2018a; 2019; Sensoy et al., 2018; Malinin et al., 2019; Charpentier et al., 2020; Zhao et al., 2020; Nandy et al., 2020; Shi et al., 2020; Sensoy et al., 2020). Contrary to other approaches such as Bayesian Neural Networks (Blundell et al., 2015; Osawa et al., 2019; Maddox et al., 2019), drop out (Gal & Ghahramani, 2016) or ensembles (Lakshminarayanan et al., 2017), DBU models provide efficient uncertainty estimates at test time in a single forward pass by directly predicting the parameters of a Dirichlet distribution over categorical probability distributions. DBU models have the advantage that they provide both, aleatoric uncertainty estimates resulting from irreducible uncertainty (e.g. class overlap or noise) and epistemic uncertainty estimates resulting from the lack of knowledge about unseen data (e.g. an unknown object is presented to the model). Both uncertainty types can be quantified from Dirichlet distributions using different uncertainty measures such as differential entropy, mutual information, or pseudo-counts. These uncertainty measures show outstanding performance in, e.g., the detection of OOD samples and thus are superior to softmax based confidence (Malinin & Gales, 2018a; 2019; Charpentier et al., 2020).

^{*}Equal contribution ¹Technical University of Munich, Germany; Department of Informatics. Correspondence to: Anna-Kathrin Kopetzki kopetzki@in.tum.de, Bertrand Charpentier kopetzki@in.tum.de), Bertrand Charpentier kopetzki@in.tum.d

Neural networks from the families outlined above are expected to know what they don't know, i.e. they are supposed to notice when they are unsure about a prediction. This raises questions with regards to adversarial examples: should uncertainty estimation methods detect these corrupted samples by indicating a high uncertainty on them and abstain from making a prediction? Or should uncertainty estimation be robust to adversarial examples and assign the correct label even under perturbations? We argue that being robust to adversarial perturbations is the best option (see Figure 1) for two reasons. First, in image classification a human is usually not able to observe any difference between an adversarial example and an unperturbed image. Second, the size of the perturbation corresponding to a good adversarial example is typically small w.r.t. the L_p -norm and thus assumed to be semantically meaningless. Importantly, robustness should not only be required for the class predictions, but also for the uncertainty estimates. This means that DBU models should be able to distinguish robustly between ID and OOD data even if those are perturbed.

In this work, we focus on DBU models and analyze their robustness capacity w.r.t. class predictions as well as uncertainty predictions. In doing so, we go beyond simple softmax output confidence by investigating advanced uncertainty measures such as differential entropy. Specifically, we study the following questions:

- 1. Is low uncertainty a reliable indicator of correct predictions?
- 2. Can we use uncertainty estimates to detect label attacks on the class prediction?
- 3. Are uncertainty estimates such as differential entropy a robust feature for OOD detection?

In addressing these questions we place particular focus on adversarial perturbations of the input to evaluate the worst case performance of the models on increasing complex data sets and attacks. We evaluate robustness of DBU models w.r.t. to these three questions by comparing their performance on unperturbed and perturbed inputs. Perturbed inputs are obtained by computing *label attacks* and uncertainty attacks, which are a new type of attacks we propose. While label attacks aim at changing the class prediction, uncertainty attacks aim at changing the uncertainty estimate such that ID data is marked as OOD data and vice versa. In total, we performed more than 138, 800 attack settings to explore the robustness landscape of DBU models. Those settings cover different data sets, attack types, attack losses, attack radii, DBU model types and initialisation seeds. Finally, we propose and evaluate median smoothing and adversarial training based on label attacks and uncertainty attacks to make DBU models more

robust. Our median smoothing approach provides certificates on epistemic uncertainty measures such as differential entropy and allows to certify uncertainty estimation. The code and further supplementary material is available online (www.daml.in.tum.de/dbu-robustness).

2. Related work

The existence of adversarial examples is a problematic property of neural networks (Szegedy et al., 2014; Goodfellow et al., 2015). Previous works have study this phenomena by proposing adversarial attacks (Carlini & Wagner, 2017; Brendel et al., 2018; Zügner et al., 2018), defenses (Cisse et al., 2017; Gu & Rigazio, 2015) and verification techniques (Wong & Kolter, 2018; Singh et al., 2019; Cohen et al., 2019; Bojchevski et al., 2020; Kopetzki & Günnemann, 2021). This includes the study of different settings such as i.i.d. inputs, sequential inputs and graphs (Zheng et al., 2016; Bojchevski & Günnemann, 2019; Cheng et al., 2020; Schuchardt et al., 2021).

In the context of uncertainty estimation, robustness of the class prediction has been studied in previous works for Bayesian Neural Networks (Blundell et al., 2015; Osawa et al., 2019; Maddox et al., 2019), drop out (Gal & Ghahramani, 2016) or ensembles (Lakshminarayanan et al., 2017) focusing on data set shifts (Ovadia et al., 2019) or adversarial attacks (Carbone et al., 2020; Cardelli et al., 2019; Wicker et al., 2020). Despite their efficient and high quality uncertainty estimates, the robustness of DBU models has not been investigated in detail yet — indeed only for one single DBU model, (Malinin & Gales, 2019) has briefly performed attacks aiming to change the label. In contrast, our work focuses on a large variety of DBU models and analyzes two robustness properties: robustness of the class prediction w.r.t. adversarial perturbations and robustness of uncertainty estimation w.r.t. our newly proposed attacks against uncertainty measures.

This so called *uncertainty attack* directly targets uncertainty estimation and are different from traditional *label attacks*, which target the class prediction (Madry et al., 2018; Dang-Nhu et al., 2020). They allow us to jointly evaluate robustness of the class prediction and robustness of uncertainty estimation. This goes beyond previous attack defenses that were either focused on evaluating *robustness w.r.t. class predictions* (Carlini & Wagner, 2017; Weng et al., 2018) or detecting attacks against the class prediction (Carlini & Wagner, 2017).

Different models have been proposed to account for uncertainty while being robust. (Smith & Gal, 2018) and (Lee et al., 2018) have tried to improve label attack detection based on uncertainty using drop-out or density estimation. In addition to improving label attack detection for large un-

seen perturbations, (Stutz et al., 2020) aimed at improving robustness w.r.t. class label predictions on small input perturbations. They used adversarial training and soft labels for adversarial samples further from the original input. (Qin et al., 2020) suggested a similar adversarial training procedure, that softens labels depending on the input robustness. These previous works consider the aleatoric uncertainty that is contained in the predicted categorical probabilities, but in contrast to DBU models they are not capable of taking epistemic uncertainty into account.

Recently, four studies tried to obtain certificates on aleatoric uncertainty estimates. (Tagasovska & Lopez-Paz, 2019) and (Kumar et al., 2020) compute confidence intervals and certificates on softmax predictions. (Bitterwolf et al., 2020) uses interval bound propagation to compute bounds on softmax predictions within the L_{∞} -ball around an OOD sample using ReLU networks. (Meinke & Hein, 2020) focuses on obtaining certifiably low confidence for OOD data. These four studies estimate confidence based on softmax predictions, which accounts for aleatoric uncertainty only. In this paper, we provide certificates which apply for all uncertainty measures. In particular, we use our certificates on epistemic uncertainty measures such as differential entropy which are well suited for OOD detection.

3. Dirichlet-based uncertainty models

Standard (softmax) neural networks predict the parameters of a categorical distribution $\boldsymbol{p}^{(i)} = [p_1^{(i)}, \dots, p_C^{(i)}]$ for a given input $\boldsymbol{x}^{(i)} \in \mathbb{R}^d$, where C is the number of classes. Given the parameters of a categorical distribution, the *aleatoric uncertainty* can be evaluated. The aleatoric uncertainty is the uncertainty on the class label prediction $\boldsymbol{y}^{(i)} \in \{1,\dots,C\}$. For example if we predict the outcome of an unbiased coin flip, the model is expected to have high aleatoric uncertainty and predict $\boldsymbol{p}(\text{head}) = 0.5$.

In contrast to standard (softmax) neural networks, DBU models predict the parameters of a Dirichlet distribution – the natural prior of categorical distributions – given input $\boldsymbol{x}^{(i)}$ (i.e. $q^{(i)} = \mathrm{Dir}(\boldsymbol{\alpha}^{(i)})$ where $f_{\theta}(\boldsymbol{x}^{(i)}) = \boldsymbol{\alpha}^{(i)} \in \mathbb{R}_{+}^{C}$). Hence, the *epistemic distribution* $q^{(i)}$ expresses the *epistemic* uncertainty on $\boldsymbol{x}^{(i)}$, i.e. the uncertainty on the categorical distribution prediction $p^{(i)}$. From the epistemic distribution, follows an estimate of the *aleatoric distribution* of the class label prediction $\mathrm{Cat}(\bar{p}^{(i)})$ where $\mathbb{E}_{q^{(i)}}[p^{(i)}] = \bar{p}^{(i)}$. An advantage of DBU models is that one pass through the neural network is sufficient to compute epistemic distribution, aleatoric distribution, and predict the class label:

$$q^{(i)} = \text{Dir}(\boldsymbol{\alpha}^{(i)}), \ \bar{p}_c^{(i)} = \frac{\alpha_c^{(i)}}{\alpha_0^{(i)}}, \ y^{(i)} = \arg\max_c \ [\bar{p}_c^{(i)}]$$

where $\alpha_0^{(i)} = \sum_{c=1}^C \alpha_c^{(i)}$. This parametrization allows to compute classic uncertainty measures in closed-form such as the total pseudo-count $m_{\alpha_0}^{(i)} = \sum_c \alpha_c^{(i)}$, the differential entropy of the Dirichlet distribution $m_{\text{diffE}}^{(i)} = h(\text{Dir}(\boldsymbol{\alpha}^{(i)}))$ or the mutual information $m_{\text{MI}}^{(i)} = I(y^{(i)}, \boldsymbol{p}^{(i)})$ (App. 6.1, (Malinin & Gales, 2018a)). Hence, these measure can efficiently be used to assign high uncertainty to unknown data, which makes DBU models specifically suited for detection of OOD samples.

Several recently proposed models for uncertainty estimations belong to the family of DBU models, such as PriorNet, EvNet, DDNet and PostNet. These models differ in terms of their parametrization of the Dirichlet distribution, the training, and density estimation. An overview of theses differences is provided in Table 1. In our study we evaluate all recent versions of these models.

Contrary to the other models, Prior Networks (**PriorNet**) (Malinin & Gales, 2018a; 2019) require OOD data for training to "teach" the neural network the difference between ID and OOD data. PriorNet is trained with a loss function consisting of two KL-divergence terms. The fist term is designed to learn Dirichlet parameters for ID data, while the second one is used to learn a flat Dirichlet distribution for OOD data:

$$L_{\text{PriorNet}} = \frac{1}{N} \left[\sum_{\boldsymbol{x}^{(i)} \in \text{ID data}} [\text{KL}[\text{Dir}(\alpha^{\text{ID}})||q^{(i)}]] + \sum_{\boldsymbol{x}^{(i)} \in OOD data} [\text{KL}[\text{Dir}(\alpha^{\text{OOD}})||q^{(i)}]] \right]$$
(2)

where $\alpha^{\rm ID}$ and $\alpha^{\rm OOD}$ are hyper-parameters. Usually $\alpha^{\rm ID}$ is set to $1e^1$ for the correct class and 1 for all other classes, while $\alpha^{\rm OOD}$ is set to 1 for all classes. There a two variants of PriorNet. The first one is trained based on reverse KL-divergence (Malinin & Gales, 2019), while the second one is trained with KL-divergence (Malinin & Gales, 2018a). In our experiments, we include the most recent reverse version of PriorNet, as it shows superior performance (Malinin & Gales, 2019).

Evidential Networks (**EvNet**) (Sensoy et al., 2018) are trained with a loss that computes the sum of squares between the on-hot encoded true label $y^{*(i)}$ and the predicted categorical $p^{(i)}$ under the Dirichlet distribution:

$$L_{\text{EvNet}} = \frac{1}{N} \sum_{i} \mathbb{E}_{\boldsymbol{p}^{(i)} \sim \text{Dir}(\boldsymbol{\alpha}^{(i)})} ||\boldsymbol{y} *^{(i)} - \boldsymbol{p}^{(i)}||^{2}$$
(3)

Ensemble Distribution Distillation (**DDNet**) (Malinin et al., 2019) is trained in two steps. First, an ensemble of M classic neural networks needs to be trained. Then, the softlabels $\{p_m^{(i)}\}_{m=1}^M$ provided by the ensemble of networks are

	$\alpha^{(i)}$ -parametrization	Loss	OOD training data	Ensemble training	Density estimation
PostNet	$f_{ heta}(oldsymbol{x}^{(i)}) = oldsymbol{1} + oldsymbol{lpha}^{(i)}$	Bayesian loss	No	No	Yes
PriorNet	$f_{ heta}(oldsymbol{x}^{(i)}) = oldsymbol{lpha}^{(i)}$	Reverse KL	Yes	No	No
DDNet	$f_{ heta}(oldsymbol{x}^{(i)}) = oldsymbol{lpha}^{(i)}$	Dir. Likelihood	No	Yes	No
EvNet	$f_{ heta}(oldsymbol{x}^{(i)}) = oldsymbol{1} + oldsymbol{lpha}^{(i)}$	Expected MSE	No	No	No

Table 1. Summary of DBU models. Further details on the loss functions are provided in the appendix.

distilled into a Dirichlet-based network by fitting them with the maximum likelihood under the Dirichlet distribution:

$$L_{\text{DDNet}} = -\frac{1}{N} \sum_{i} \sum_{m=1}^{M} [\ln q^{(i)}(\pi^{im})]$$
 (4)

where π^{im} denotes the soft-label of mth neural network.

Posterior Network (**PostNet**) (Charpentier et al., 2020) performs density estimation for ID data with normalizing flows and uses a Bayesian loss formulation:

$$L_{\text{PostNet}} = \frac{1}{N} \sum_{i} \mathbb{E}_{q(p^{(i)})} [\text{CE}(p^{(i)}, y^{(i)})] - H(q^{(i)})$$

where CE denotes the cross-entropy. All loss functions can be computed in closed-form. For more details please have a look at the original paper on PriorNet (Malinin & Gales, 2018a), PostNet (Charpentier et al., 2020), DDNet (Malinin & Gales, 2019) and EvNet (Sensoy et al., 2018). Note that EvNet and PostNet model the Dirichlet parameters as $f_{\theta}(\boldsymbol{x}^{(i)}) = 1 + \boldsymbol{\alpha}^{(i)}$ while PriorNet, RevPriorNet and DDNet compute them as $f_{\theta}(\boldsymbol{x}^{(i)}) = \boldsymbol{\alpha}^{(i)}$.

4. Robustness of Dirichlet-based uncertainty models

We analyze robustness of DBU models on tasks in connection with uncertainty estimation w.r.t. the following four aspects: *accuracy*, *confidence calibration*, *label attack detection* and *OOD detection*. Uncertainty is quantified by differential entropy, mutual information or pseudo counts. A formal definition of all uncertainty estimation measures is provided in the appendix (see Section 6.1).

Robustness of Dirichlet-based uncertainty models is evaluated based on *label attacks* and a newly proposed type of attacks called *uncertainty attacks*. While label attacks aim at changing the predicted class, uncertainty attacks aim at changing the uncertainty assigned to a prediction. All previous works are based on label attacks and focus on robustness w.r.t. the class prediction. Thus, we are the first to propose attacks targeting uncertainty estimates such as differential entropy and analyze desirable robustness properties of DBU models beyond the class prediction. Label attacks and uncertainty attacks both compute a perturbed input $\tilde{x}^{(i)}$ close to the original input $x^{(i)}$ i.e. $||x^{(i)} - \tilde{x}^{(i)}||_2 < r$ where

r is the attack radius. This perturbed input is obtained by optimizing a loss function $l(\boldsymbol{x})$ using Fast Gradient Sign Method (FGSM) or Projected Gradient Descent (PGD). Furthermore, we include a black box attack setting (Noise) which generates 10 noise samples from a Gaussian distribution, which is centered at the original input. From these 10 perturbed samples we choose the one with the greatest effect on the loss function and use it as attack. To complement attacks, we compute certificates on uncertainty estimates using median smoothing (yeh Chiang et al., 2020).

The following questions we address by our experiments have a common assessment metric and can be treated as binary classification problems: distinguishing between correctly and wrongly classified samples, discriminating between non-attacked input and attacked inputs or differentiating between ID data and OOD data. To quantify the performance of the models on these binary classification problems, we compute the area under the precision recall curve (AUC-PR).

Experiments are performed on two image data sets (MNIST (LeCun & Cortes, 2010) and CIFAR10 (Krizhevsky et al., 2009)), which contain bounded inputs and two tabular data sets (Segment (Dua & Graff, 2017) and Sensorless drive (Dua & Graff, 2017)), consisting of unbounded inputs. Note that unbounded inputs are challenging since it is impossible to describe the infinitely large OOD distribution. As Prior-Net requires OOD training data, we use two further image data sets (FashionMNIST (Xiao et al., 2017) and CIFAR100 (Krizhevsky et al., 2009)) for training on MNIST and CI-FAR10, respectively. All other models are trained without OOD data. To obtain OOD data for the tabular data sets, we remove classes from the ID data set (class window for the Segment data set and class 9 for Sensorless drive) and use them as the OOD data. Further details on the experimental setup are provided in the appendix (see Section 6.2).

4.1. Uncertainty estimation under label attacks

Label attacks aim at changing the predicted class. To obtain a perturbed input with a different label, we maximize the cross-entropy loss $\tilde{\boldsymbol{x}}^{(i)} \approx \arg\max_{\boldsymbol{x}} l(\boldsymbol{x}) = \mathrm{CE}(\boldsymbol{p}^{(i)}, \boldsymbol{y}^{(i)})$ under the radius constraint. For the sake of completeness we additionally analyze label attacks w.r.t. to their performance of changing the class prediction and the accuracy of the neural network under label attacks constraint by different radii (see Appendix, Table 8). As expected and partially

shown by previous works, none of the DBU models is robust against label attacks. However, we note that PriorNet is slightly more robust than the other DBU models. This might be explained by the use of OOD data during training, which can be seen as some kind of robust training. From now on, we switch to the core focus of this work and analyze robustness properties of uncertainty estimation.

Is low uncertainty a reliable indicator of correct predictions?

Expected behavior: Predictions with low uncertainty are more likely to be correct than high uncertainty predictions. Assessment metric: We distinguish between correctly classified samples (label 0) and wrongly classified ones (label 1) based on the differential entropy scores produced by the DBU models (Malinin & Gales, 2018a). Correctly classified samples are expected to have low differential entropy, reflecting the model's confidence, and analogously wrongly predicted samples are expected to have higher differential entropy. Observed behavior: Note that the positive and negative class are not balanced, thus, the use of AUC-PR scores (Saito & Rehmsmeier, 2015) are important to enable meaningful measures. While uncertainty estimates are indeed an indicator of correctly classified samples on unperturbed data, none of the models maintains its high performance on perturbed data computed by PGD, FGSM or Noise label attacks (see. Table 2, 14 and 15). Thus, using uncertainty estimates as indicator for correctly labeled inputs is not robust to adversarial perturbations. This result is notable, since the used attacks do not target uncertainty.

Can uncertainty estimates be used to detect label attacks against the class prediction?

Expected behavior: Adversarial examples are not from the natural data distribution. Therefore, DBU models are expected to detect them as OOD data by assigning them a higher uncertainty. We expect that perturbations computed based on a bigger attack radius r are easier to detect as their distance from the data distribution is larger. Assessment metric: The goal of attack-detection is to distinguish between unperturbed samples (label 0) and perturbed samples (label 1). Uncertainty on samples is quantified by the differential uncertainty (Malinin & Gales, 2018a). Unperturbed samples are expected to have low differential entropy, because they are from the same distribution as the training data, while perturbed samples are expected to have a high differential entropy. Observed behavior: Table 8 shows that the accuracy of all models decreases significantly under PGD label attacks, but none of the models is able to provide an equivalently increasing attack detection rate (see Table 3). Even larger perturbations are hard to detect for DBU models.

Similar results are obtained when we use mutual information or the precision α_0 to quantify uncertainty (see ap-

pendix Table 13 and 12). Although PGD label attacks do not explicitly consider uncertainty, they seem to generate adversarial examples with similar uncertainty as the original input. Such high-certainty adversarial examples are illustrated in Figure 2, where certainty is visualized based on the precision α_0 , which is supposed to be high for ID data and low for OOD data. While the original input (perturbation size 0.0) is correctly classified as frog and ID data, there exist adversarial examples that are classified as deer or bird. The certainty (α_0 -score) on the prediction of these adversarial examples has a similar or even higher value than on the prediction of the original input. Using the differential entropy to distinguish between ID and OOD data results in the same ID/OOD assignment since the differential entropy of the three right-most adversarial examples is similar or even smaller than on the unperturbed input.

Under the less powerful FGSM and Noise attacks (see Appendix), DBU models achieve mostly higher attack detection rates than under PGD attacks. This suggests that uncertainty estimation is able to detect weak attacks, which is consistent with the observations in (Malinin & Gales, 2018b) but fails under stronger PGD attacks.

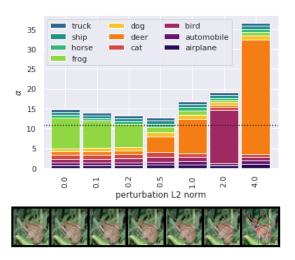


Figure 2. Input and predicted Dirichlet-parameters under label attacks (dotted line: threshold to distinguish ID and OOD data).

On tabular data sets, PostNet shows a better label attack detection rate for large perturbations. This observation might be explained by the fact that the density estimation of the ID samples has been shown to work better for tabular data sets (Charpentier et al., 2020). Overall, none of the DBU models provides a reliable indicator for adversarial inputs that target the class prediction.

4.2. Attacking uncertainty estimation

DBU models are designed to provide sophisticated uncertainty estimates (beyond softmax scores) alongside predic-

-ttl (m-tri ALIC DD)	label
attacks (metric: AUC-PR).	

			CIFA	AR10				,	Sensorl	ess		
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	0.0	0.1	0.2	0.5	1.0	2.0
PostNet	98.7	88.6	56.2	7.8	1.2	0.4	99.7	8.3	3.9	3.6	7.0	9.8
PriorNet	92.9	77.7	60.5	37.6	24.9	11.3	99.8	10.5	3.2	0.7	0.2	0.2
DDNet	97.6	91.8	78.3	18.1	0.8	0.0	99.7	11.9	1.6	0.4	0.2	0.1
EvNet	97.9	85.9	57.2	10.2	4.0	2.4	99.9	22.9	13.0	6.0	3.7	3.2

Table 3. Label Attack-Detection by normally trained DBU models based on differential entropy under PGD label attacks (AUC-PR).

		(CIFAR1	0			Se	nsorles	S	
Att. Rad.	0.1	0.2	0.5	1.0	2.0	0.1	0.2	0.5	1.0	2.0
PostNet	63.4	66.9	42.1	32.9	31.6	47.7	42.3	36.9	48.5	85.0
PriorNet	53.3	56.0	55.6	49.2	42.2	38.8	33.6	31.4	33.1	40.9
DDNet	55.8	60.5	57.3	38.7	32.3	53.5	42.2	35.0	32.8	32.6
EvNet	48.4	46.9	46.3	46.3	44.5	48.2	42.6	38.2	36.0	37.2

tions and use them to detect OOD samples. In this section, we propose and analyze a new attack type that targets these uncertainty estimates. DBU models enable us to compute uncertainty measures i.e. differential entropy, mutual information and precision α_0 in closed from (see (Malinin & Gales, 2018a) for a derivation). Uncertainty attacks use this closed form solution as loss function for PGD, FGSM or Noise attacks. Since differential entropy is the most widely used metric for ID-OOD-differentiation, we present results based on the differential entropy loss function $\tilde{x}^{(i)} \approx \arg\max_{\boldsymbol{x}} l(\boldsymbol{x}) = \text{Diff-E}(\text{Dir}(\alpha^{(i)}))$:

$$\begin{split} \text{Diff-E}(\text{Dir}(\alpha^{(i)})) &= \sum_{c}^{K} \ln \Gamma(\alpha_{c}^{(i)}) - \ln \Gamma(\alpha_{0}^{(i)}) \\ &- \sum_{c}^{K} (\alpha_{c}^{(i)} - 1) \cdot (\Psi(\alpha_{c}^{(i)}) - \Psi(\alpha_{0}^{(i)})) \end{split} \tag{6}$$

where $\alpha_0^{(i)} = \sum_c \alpha_c^{(i)}$. Result based on further uncertainty measures, loss functions and more details on attacks are provided in the appendix.

We analyze the performance of DBU models under uncertainty attacks w.r.t. two tasks. First, uncertainty attacks are computed on ID data aiming to indicate it as OOD data, while OOD data is left non-attacked. Second, we attack OOD data aiming to indicate it as ID data, while ID data is not attacked. Hence, uncertainty attacks target at posing ID data as OOD data and vice versa.

Are uncertainty estimates a robust feature for OOD detection?

Expected behavior: We expect DBU models to be able to

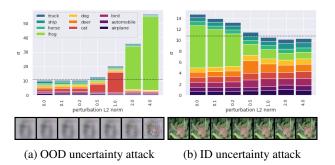


Figure 3. ID and OOD input with corresponding Dirichletparameters under uncertainty attacks (dotted line: threshold to distinguish ID and OOD).

distinguish between ID and OOD data by providing reliable uncertainty estimates, even under small perturbations. Thus, we expect uncertainty estimates of DBU models to be robust under attacks. Assessment metric: We distinguish between ID data (label 0) and OOD data (label 1) based on the differential entropy as uncertainty scoring function (Malinin & Gales, 2018a). Differential entropy is expected to be small on ID samples and high on OOD samples. Experiments on further uncertainty measure and results on the AUROC metric are provided in the appendix. Observed behavior: OOD samples are perturbed as illustrated in Figure 3. Part (a) of the figure illustrates an OOD-samples, that is correctly identified as OOD. Adding adversarial perturbations ≥ 0.5 changes the Dirichlet parameters such that the resulting images are identified as ID, based on precision or differential entropy as uncertainty measure. Perturbing an ID sample (part (b)) results in images that are marked as

Table 4. OOD detection based on differential entropy under PGD uncertainty attacks against differential entropy computed on ID data and
OOD data (metric: AUC-PR).

	I	D-Atta	ck (non	-attacke	ed OOD))		C	OD-At	tack (no	on-attac	ked ID))
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0		0.0	0.1	0.2	0.5	1.0	2.0
						CIFAR10) – S	SVHN					
PostNet	81.8	64.3	47.2	22.4	17.6	16.9		81.8	60.5	40.7	23.3	21.8	19.8
PriorNet	54.4	40.1	30.0	17.9	15.6	15.4		54.4	40.7	30.7	19.5	16.5	15.7
DDNet	82.8	71.4	59.2	28.9	16.0	15.4		82.8	72.0	57.2	20.8	15.6	15.4
EvNet	80.3	62.4	45.4	21.7	17.9	16.5		80.3	58.2	46.5	34.6	28.0	23.9
					Ser	ıs. – Sens	. cla	ass 10,	11				
PostNet	74.5	39.8	36.1	36.0	45.9	46.0		74.5	43.3	42.0	32.1	35.1	82.6
PriorNet	32.3	26.6	26.5	26.5	26.6	28.3		32.3	26.7	26.6	26.6	27.0	30.4
DDNet	31.7	26.8	26.6	26.5	26.6	27.1		31.7	27.1	26.7	26.7	26.8	26.9
EvNet	66.5	30.5	28.2	27.1	28.1	31.8		66.5	38.7	36.1	30.2	28.2	28.8

OOD samples. OOD detection performance of all DBU models rapidly decreases with the size of the perturbation, regardless of whether attacks are computed on ID or OOD data (see Table 4). This performance decrease is also observed with AUROC as metric, attacks based on FGSM, Noise, when we use mutual information or precision α_0 to distinguish between ID samples and OOD samples (see appendix Table 21 - 28). Thus, using uncertainty estimation to distinguish between ID and OOD data is not robust.

4.3. How to make DBU models more robust?

Our robustness analysis based on label attacks and uncertainty attacks shows that predictions, uncertainty estimation and the differentiation between ID and OOD data are not robust. Next, we explore approaches to improve robustness properties of DBU models w.r.t. these tasks based on randomized smoothing and adversarial training.

Randomized smoothing was originally proposed for certification of classifiers (Cohen et al., 2019). The core idea is to draw multiple samples $x_s^{(i)} \sim \mathcal{N}(x^{(i)}, \sigma)$ around the input data $x^{(i)}$, to feed all these samples through the neural network, and to aggregate the resulting set of predictions (e.g. by taking their mean), to get a smoothed prediction. Besides allowing certification, as a side effect, the smoothed model is more robust. Our idea is to use randomized smoothing to improve robustness of DBU models, particularly w.r.t. uncertainty estimation. In contrast to discrete class predictions, however, certifying uncertainty estimates such as differential entropy scores requires a smoothing approach that is able to handle continuous values as in regression tasks. So far, only few works for randomized smoothing for regression models have been proposed (Kumar et al., 2020;

yeh Chiang et al., 2020). We choose median smoothing (yeh Chiang et al., 2020), because it is applicable to unbounded domains as required for the uncertainty estimates covered in this work. In simple words: The set of uncertainty scores obtained from the $\boldsymbol{x}_s^{(i)} \sim \mathcal{N}(\boldsymbol{x}^{(i)}, \sigma)$ is aggregated by taking their median.

In the following experiments we focus on differential entropy as the uncertainty score. We denote the resulting smoothed differential entropy, i.e. the median output, as $m(\boldsymbol{x}^{(i)})$. Intuitively, we expect that the random sampling around a data point as well as the outlier-insensitivity of the median to improve the robustness of the uncertainty estimates w.r.t. adversarial examples.

To measure the performance and robustness of our smoothed DBU models, we apply median smoothing on the same tasks as in the previous sections, i.e., distinguishing between correctly and wrongly labeled inputs, attack detection, OOD detection and compute the corresponding AUC-PR score under label attacks and uncertainty attacks. The bold, middle part of the columns in Tables 5, 6, and 7 show the AUC-PR scores on CIFAR10, which we call empirical performance of the smoothed models. To facilitate the comparison with the base model of Section 4, we highlight the AUC-PR scores in blue in cases where the smooth model is more robust. The highlighting clearly shows that randomized smoothing increases the robustness of the empirical performance on OOD detection. OOD detection under strong PGD attacks (attack radius ≥ 0.5) performs comparable to random guessing (i.e. AUC-PR scores around 50% whith 50% ID and 50% OOD data). This shows that DBU models are not reliably efficient w.r.t. this task. In attack detection and distinguishing between correctly and wrongly predicted

Table 5. Distinguishing between correctly and wrongly predicted labels based on differential entropy under PGD label attacks. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$80.5 \cdot 91.5 \cdot 94.5$ $81.9 \cdot 86.8 \cdot 88.0$ $65.9 \cdot 81.2 \cdot 83.0$ $76.3 \cdot 90.2 \cdot 91.7$	52.8 · 71.6 · 95.2 69.6 · 78.0 · 90.1 55.8 · 70.5 · 87.2 54.7 · 74.3 · 95.7	$31.9 \cdot 51.0 \cdot 96.8 \\ 50.9 \cdot 65.8 \cdot 89.4 \\ 37.8 \cdot 56.8 \cdot 88.1 \\ 31.6 \cdot 51.5 \cdot 94.5$	$\begin{array}{c} 5.6 \cdot 11.7 \cdot 100.0 \\ 36.5 \cdot 59.9 \cdot 97.0 \\ 10.1 \cdot 21.9 \cdot 94.3 \\ 5.8 \cdot 11.9 \cdot 86.9 \end{array}$	$0.3 \cdot 0.6 \cdot 100.0$ $24.3 \cdot 39.3 \cdot 100.0$ $0.9 \cdot 1.6 \cdot 99.6$ $1.9 \cdot 7.0 \cdot 100.0$	$0.0 \cdot 0.0 \cdot 100.0$ $9.2 \cdot 17.9 \cdot 100.0$ $0.0 \cdot 0.0 \cdot 100.0$ $1.1 \cdot 4.0 \cdot 100.0$
Smoothed + adv. w. label attacks	PostNet PriorNet DDNet EvNet	- - -	52.1 · 71.8 · 95.6 57.6 · 71.7 · 88.9 58.6 · 78.4 · 92.2 24.3 · 34.2 · 51.8	31.2 · 47.9 · 96.1 46.1 · 64.5 · 90.1 49.4 · 66.0 · 90.5 32.6 · 49.5 · 95.5	$7.8 \cdot 14.7 \cdot 98.6$ $38.1 \cdot 59.3 \cdot 99.5$ $12.0 \cdot 21.4 \cdot 98.1$ $5.9 \cdot 13.0 \cdot 100.0$	1.8 · 4.4 · 100.0 32.3 · 51.7 · 100.0 0.8 · 1.0 · 96.6 2.6 · 5.2 · 99.9	$\begin{array}{c} 0.3 \cdot 0.5 \cdot 100.0 \\ 22.1 \cdot 41.6 \cdot 97.4 \\ 0.0 \cdot 0.0 \cdot 100.0 \\ 2.9 \cdot 5.9 \cdot 100.0 \end{array}$
Smoothed + adv. w. uncert. attacks	PostNet PriorNet DDNet EvNet	- - -	52.8 · 74.2 · 94.6 50.6 · 68.1 · 88.6 68.8 · 84.4 · 93.2 54.2 · 73.7 · 96.1	33.0 · 49.4 · 87.5 44.4 · 66.1 · 96.0 45.1 · 60.8 · 86.8 30.5 · 50.0 · 99.5	$7.7 \cdot 14.2 \cdot 99.0 \\ 35.1 \cdot 57.4 \cdot 98.4 \\ 12.3 \cdot 22.0 \cdot 91.0 \\ 7.1 \cdot 13.9 \cdot 100.0$	0.6 · 1.2 · 100.0 18.4 · 32.2 · 100.0 0.8 · 1.7 · 87.0 3.7 · 8.7 · 75.2	$0.7 \cdot 1.1 \cdot 100.0$ $15.2 \cdot 29.3 \cdot 100.0$ $0.0 \cdot 0.0 \cdot 100.0$ $3.3 \cdot 5.8 \cdot 100.0$

Table 6. Attack detection (PGD label attacks) based on differential entropy. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	33.1 · 50.4 · 89.9 35.9 · 50.6 · 74.5 36.3 · 50.3 · 76.4 32.9 · 50.4 · 89.8	31.0 · 50.2 · 96.9 33.0 · 50.3 · 82.8 32.8 · 49.9 · 84.6 31.4 · 50.1 · 94.0	$30.7 \cdot 50.2 \cdot 100.0$ $31.2 \cdot 50.0 \cdot 95.7$ $30.8 \cdot 50.1 \cdot 98.0$ $30.8 \cdot 50.0 \cdot 98.0$	$30.7 \cdot 50.0 \cdot 100.0$ $30.7 \cdot 50.4 \cdot 99.9$ $30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 50.3 \cdot 100.0$	$30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 50.4 \cdot 100.0$ $30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 49.6 \cdot 100.0$
Smoothed	PostNet	32.7 · 50.1 · 90.4	31.1 · 50.2 · 96.5	30.7 · 50.2 · 99.7	30.7 · 50.3 · 100.0	30.7 · 50.2 · 100.0
+ adv. w.	PriorNet	35.2 · 51.8 · 78.6	32.8 · 51.1 · 84.4	30.8 · 50.2 · 98.7	30.7 · 50.5 · 100.0	30.8 · 50.1 · 98.2
label	DDNet	35.5 · 50.6 · 79.2	33.4 · 50.3 · 84.1	30.8 · 50.1 · 99.2	30.7 · 50.0 · 100.0	30.7 · 50.5 · 100.0
attacks	EvNet	40.3 · 50.4 · 66.8	31.4 · 50.3 · 95.8	30.7 · 50.3 · 100.0	30.7 · 50.1 · 100.0	30.7 · 50.0 · 100.0
Smoothed	PostNet	$33.3 \cdot 50.6 \cdot 88.7$	$32.5 \cdot 50.1 \cdot 87.9$	30.7 · 49.9 · 99.8	30.7 · 50.1 · 100.0	$30.7 \cdot 50.0 \cdot 100.0$
+ adv. w.	PriorNet	$34.5 \cdot 51.0 \cdot 80.1$	$31.4 \cdot 50.6 \cdot 92.8$	30.9 · 50.0 · 97.7	30.7 · 50.1 · 100.0	$30.7 \cdot 50.0 \cdot 100.0$
uncert.	DDNet	$37.4 \cdot 50.8 \cdot 74.5$	$33.4 \cdot 50.2 \cdot 83.0$	30.9 · 50.1 · 96.8	30.8 · 49.9 · 98.1	$30.7 \cdot 49.9 \cdot 100.0$
attacks	EvNet	$32.8 \cdot 50.1 \cdot 92.0$	$30.8 \cdot 50.0 \cdot 99.6$	30.7 · 50.1 · 100.0	31.2 · 50.2 · 96.1	$31.0 \cdot 50.0 \cdot 100.0$

labels the smoothed DBU model are mostly more robust than the base models for attack radii ≥ 0.5 .

Certified performance. Using the median based on smoothing improves the empirical robustness, but it does not provide formal guarantees how low/high the performance might actually get under perturbed data (since any attack is only a heuristic). Here, we propose novel guarantees by exploiting the individual certificates we obtain via randomized smoothing. Note that the certification procedure (yeh Chiang et al., 2020) enables us to derive lower and upper bounds $\underline{m}(\boldsymbol{x}^{(i)}) \leq m(\boldsymbol{x}^{(i)}) \leq \overline{m}(\boldsymbol{x}^{(i)})$ which hold with high probability and indicate how much the median might change in the worst-case when $\boldsymbol{x}^{(i)}$ gets perturbed subject to a specific (attack) radius.

These bounds allow us to compute certificates that bound the performance of the smooth models, which we refer to as the *guaranteed lowest performance* and *guaranteed highest performance*. More precisely, for the guaranteed lowest performance of the model we take the pessimistic view that all ID data points realize their individual upper bounds $\overline{m}(x^{(i)})$, i.e. have their highest possible uncertainty (worst case). On the other hand, we assume all OOD samples realize their lower bounds $\underline{m}(\boldsymbol{x}_s^{(i)})$. Using these values as the uncertainty scores for all data points we obtain the guaranteed lowest performance of the model. A guaranteed lowest performance of e.g. 35.0 means that even under the worst case conditions an attack is not able to decrease the performance below 35.0. Analogously, we can take the optimistic view to obtain the guaranteed highest performance of the smoothed models. Tables 5, 6 and 7 show the guaranteed lowest/highest performance (non-bold, left/right of the empirical performance). Our results show that the difference between guaranteed highest and guaranteed lowest performance increases with the attack radius, which might be explained by the underlying lower/upper bounds on the median being tighter for smaller perturbations.

Table 7. OOD detection based on differential entropy under PGD uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
				ID-	Attack		
	PostNet	$72.1 \cdot 82.7 \cdot 88.0$	$35.0\cdot56.6\cdot97.4$	$31.9 \cdot 65.6 \cdot 99.8$	$30.7 \cdot 50.6 \cdot 100.0$	$30.7 \cdot 46.9 \cdot 100.0$	$30.7 \cdot 51.6 \cdot 100.0$
Smoothed	PriorNet	$50.2 \cdot 53.1 \cdot 55.9$	$33.5 \cdot 43.3 \cdot 65.3$	$31.3 \cdot 39.7 \cdot 69.1$	$31.3 \cdot 48.3 \cdot 98.2$	$30.7 \cdot 44.4 \cdot 99.9$	$30.7 \cdot 45.4 \cdot 100.0$
models	DDNet	$72.0 \cdot 75.8 \cdot 79.8$	$35.6 \cdot 46.2 \cdot 69.8$	$32.9 \cdot 50.3 \cdot 87.1$	$31.1 \cdot 58.7 \cdot 98.6$	$30.7 \cdot 59.3 \cdot 100.0$	$30.7 \cdot 44.5 \cdot 100.0$
	EvNet	$79.5 \cdot 87.1 \cdot 92.8$	$34.1 \cdot \textbf{58.6} \cdot 95.1$	$32.5 \cdot 61.2 \cdot 96.9$	$31.7 \cdot 60.6 \cdot 98.7$	$30.7 \cdot 62.4 \cdot 100.0$	$30.7 \cdot 57.3 \cdot 100.0$
Smoothed	PostNet	-	$35.0 \cdot \textbf{58.5} \cdot 97.7$	$31.2\cdot46.6\cdot97.4$	$30.8 \cdot 57.7 \cdot 99.7$	$30.7 \cdot 49.8 \cdot 100.0$	$30.7 \cdot $ 50.9 $\cdot 100.0$
+ adv. w.	PriorNet	-	$31.5 \cdot 36.7 \cdot 57.2$	$33.1 \cdot 51.8 \cdot 84.8$	$30.7 \cdot 57.7 \cdot 98.7$	$30.7 \cdot 40.0 \cdot 99.9$	$30.9 \cdot 53.6 \cdot 96.7$
label	DDNet	-	$36.2 \cdot 50.0 \cdot 78.6$	$32.1 \cdot 41.3 \cdot 70.2$	$30.8 \cdot 56.4 \cdot 100.0$	$30.7 \cdot 49.4 \cdot 100.0$	$30.7 \cdot 54.8 \cdot 100.0$
attacks	EvNet	-	$46.8 \cdot 61.0 \cdot 79.7$	$32.3 \cdot 58.9 \cdot 99.1$	$30.7 \cdot 45.0 \cdot 100.0$	$30.7 \cdot 63.3 \cdot 100.0$	$30.8 \cdot 38.1 \cdot 100.0$
Smoothed	PostNet	-	$35.2 \cdot \textbf{55.9} \cdot 96.0$	$34.5 \cdot 59.2 \cdot 94.9$	$30.7 \cdot 47.0 \cdot 100.0$	$30.7 \cdot 58.2 \cdot 100.0$	$30.7 \cdot 42.9 \cdot 100.0$
+ adv. w.	PriorNet	-	$31.8 \cdot 38.9 \cdot 64.1$	$31.0 \cdot 41.8 \cdot 87.9$	$30.7 \cdot 42.9 \cdot 99.2$	$30.7 \cdot 48.6 \cdot 100.0$	$30.7 \cdot \textbf{46.6} \cdot 100.0$
uncert.	DDNet	-	$39.7 \cdot 52.1 \cdot 75.7$	$36.4 \cdot \textbf{56.8} \cdot 83.8$	$31.0 \cdot 51.5 \cdot 97.4$	$31.0 \cdot 56.8 \cdot 97.8$	$30.7 \cdot 49.1 \cdot 100.0$
attacks	EvNet	-	$34.8 \cdot 64.9 \cdot 99.6$	$30.8 \cdot 48.9 \cdot 99.8$	$30.7 \cdot 66.8 \cdot 100.0$	$30.9 \cdot 41.5 \cdot 93.8$	$31.1 \cdot 55.1 \cdot 100.0$
				001)-Attack		
	PostNet	$72.0 \cdot 82.7 \cdot 88.0$	$35.1\cdot \textbf{56.8}\cdot 97.3$	$32.0 \cdot 65.8 \cdot 99.8$	$30.7 \cdot 50.7 \cdot 100.0$	$30.7 \cdot \textbf{46.5} \cdot 100.0$	$30.7 \cdot 51.7 \cdot 100.0$
Smoothed	PriorNet	$50.3 \cdot 53.1 \cdot 55.9$	$33.6 \cdot 43.7 \cdot 65.9$	$31.3 \cdot 39.8 \cdot 69.4$	$31.3 \cdot 48.3 \cdot 98.2$	$30.7 \cdot 44.5 \cdot 99.9$	$30.7 \cdot \textbf{46.4} \cdot 100.0$
models	DDNet	$72.0 \cdot 75.8 \cdot 79.8$	$35.6 \cdot 46.2 \cdot 70.0$	$32.9 \cdot 50.1 \cdot 86.7$	$31.1 \cdot 58.8 \cdot 98.6$	$30.7 \cdot 59.3 \cdot 100.0$	$30.7 \cdot 44.6 \cdot 100.0$
	EvNet	$79.5 \cdot 87.1 \cdot 92.8$	$34.1 \cdot 58.8 \cdot 95.2$	$32.6 \cdot 61.2 \cdot 96.9$	$31.7 \cdot 60.5 \cdot 98.7$	$30.7 \cdot 62.4 \cdot 100.0$	$30.7 \cdot 57.6 \cdot 100.0$
Smoothed	PostNet	-	$35.0 \cdot \textbf{58.5} \cdot 97.8$	$31.2 \cdot \textbf{46.6} \cdot 97.2$	$30.8 \cdot 57.7 \cdot 99.7$	$30.7 \cdot \textbf{50.2} \cdot 100.0$	$30.7 \cdot 51.5 \cdot 100.0$
+ adv. w.	PriorNet	-	$31.6 \cdot 37.3 \cdot 59.3$	$33.2 \cdot 52.7 \cdot 85.8$	$30.7 \cdot 57.8 \cdot 98.7$	$30.7 \cdot 40.1 \cdot 99.9$	$30.9 \cdot 53.8 \cdot 96.8$
label	DDNet	-	$36.4 \cdot 50.2 \cdot 78.9$	$32.1 \cdot 41.5 \cdot 70.4$	$30.9 \cdot 56.2 \cdot 100.0$	$30.7 \cdot 49.3 \cdot 100.0$	$30.7 \cdot 55.1 \cdot 100.0$
attacks	EvNet	-	$47.2 \cdot 61.1 \cdot 80.0$	$32.4 \cdot 59.1 \cdot 99.1$	$30.7 \cdot 45.0 \cdot 100.0$	$30.7 \cdot 63.2 \cdot 100.0$	$30.8 \cdot 38.0 \cdot 100.0$
Smoothed	PostNet	-	$35.3 \cdot \textbf{56.4} \cdot 96.1$	$34.5 \cdot \textbf{59.0} \cdot 94.9$	$30.7 \cdot \textbf{46.8} \cdot 100.0$	$30.7 \cdot $ 57.8 $\cdot 100.0$	$30.7 \cdot \textbf{43.2} \cdot 100.0$
+ adv. w.	PriorNet	-	$31.9 \cdot 39.4 \cdot 65.5$	$31.0 \cdot 42.0 \cdot 88.6$	$30.7 \cdot 42.9 \cdot 99.2$	$30.7 \cdot \textbf{48.4} \cdot 100.0$	$30.7 \cdot 47.1 \cdot 100.0$
uncert.	DDNet	-	$40.2 \cdot 52.9 \cdot 76.5$	$36.4 \cdot 56.9 \cdot 83.9$	$31.1 \cdot 51.5 \cdot 97.3$	$31.0 \cdot 57.0 \cdot 97.8$	$30.7 \cdot 49.1 \cdot 100.0$
attacks	EvNet		$34.9 \cdot 64.8 \cdot 99.6$	$30.8 \cdot 48.8 \cdot 99.8$	$30.7 \cdot 66.1 \cdot 100.0$	$30.9 \cdot 41.6 \cdot 93.6$	$31.1 \cdot 54.7 \cdot 100.0$

Adversarial training. Randomized smoothing improves robustness of DBU models and allows us to compute performance guarantees. However, an open question is whether it is possible to increase robustness even further by combining it with adversarial training. To obtain adversarially trained models we augment the data set using perturbed samples that are computed by PGD attacks against the cross-entropy loss (label attacks) or the differential entropy (uncertainty attacks). These perturbed samples $\tilde{x}^{(i)}$ are computed during each epoch of the training based on inputs $x^{(i)}$ and added to the training data (with the label $y^{(i)}$ of the original input). Tables 5, 6, and 7 illustrate the results. We choose the attack radius used during training and the σ used for smoothing to be equal. To facilitate comparison, we highlight the empirical performance of the adversarially trained models in blue if it is better than the performance of the base model. Our results show that the additional use of adversarial training has a minor effect on the robustness and does not result in a significant further increase of the robustness.

We conclude that median smoothing is a promising technique to increase robustness w.r.t. distinguishing between correctly labeled samples and wrongly labeled samples, attack detection and differentiation between in-distribution data and out-of-distribution data of all Dirichlet-based un-

certainty models, while additional adversarial training has a minor positive effect on robustness.

5. Conclusion

This work analyzes robustness of uncertainty estimation by DBU models and answers multiple questions in this context. Our results show: (1) While uncertainty estimates are a good indicator to identify correctly classified samples on unperturbed data, performance decrease drastically on perturbed data-points. (2) None of the Dirichlet-based uncertainty models is able to detect PGD label attacks against the class prediction by uncertainty estimation, regardless of the used uncertainty measure. (3) Detecting OOD samples and distinguishing between ID-data and OOD-data is not robust. (4) Applying median smoothing to uncertainty estimates increases robustness of DBU models w.r.t. all analyzed tasks, while adversarial training based on label or uncertainty attacks resulted in minor improvements.

Acknowledgments

This research was supported by BMW AG.

References

- Bitterwolf, J., Meinke, A., and Hein, M. Provable worst case guarantees for the detection of out-of-distribution data. *Neural Information Processing Systems*, 2020.
- Blundell, C., Cornebise, J., Kavukcuoglu, K., and Wierstra, D. Weight uncertainty in neural networks. *International Conference on Machine Learning*, 2015.
- Bojchevski, A. and Günnemann, S. Certifiable robustness to graph perturbations. *Neural Information Processing Systems*, 2019.
- Bojchevski, A., Klicpera, J., and Günnemann, S. Efficient robustness certificates for discrete data: Sparsity-aware randomized smoothing for graphs, images and more. *International Conference on Machine Learning*, 2020.
- Brendel, W., Rauber, J., and Bethge, M. Decision-based adversarial attacks: Reliable attacks against black-box machine learning models. *International Conference on Learning Representations*, 2018.
- Carbone, G., Wicker, M., Laurenti, L., Patane, A., Bortolussi, L., and Sanguinetti, G. Robustness of bayesian neural networks to gradient-based attacks. *Neural Information Processing Systems*, 2020.
- Cardelli, L., Kwiatkowska, M., Laurenti, L., Paoletti, N., Patane, A., and Wicker, M. Statistical guarantees for the robustness of bayesian neural networks. *Conference on Artificial Intelligence*, 2019.
- Carlini, N. and Wagner, D. Adversarial examples are not easily detected: Bypassing ten detection methods. *ACM Workshop on Artificial Intelligence and Security*, 2017.
- Carlini, N. and Wagner, D. Towards evaluating the robustness of neural networks. Symposium on Security and Privacy, 2017.
- Charpentier, B., Zügner, D., and Günnemann, S. Posterior network: Uncertainty estimation without ood samples via density-based pseudo-counts. *Neural Information Processing Systems*, 2020.
- Cheng, M., Yi, J., Chen, P.-Y., Zhang, H., and Hsieh, C.-J. Seq2sick: Evaluating the robustness of sequence-to-sequence models with adversarial examples. *AAAI Conference on Artificial Intelligence*, 2020.
- Cisse, M., Bojanowski, P., Grave, E., Dauphin, Y., and Usunier, N. Parseval networks: Improving robustness to adversarial examples. *International Conference on Machine Learning*, 2017.

- Clanuwat, T., Bober-Irizar, M., Kitamoto, A., Lamb, A., Yamamoto, K., and Ha, D. Deep learning for classical japanese literature. *Neural Information Processing Systems, Machine Learning for Creativity and Design Workshop*, 2018.
- Cohen, J. M., Rosenfeld, E., and Kolter, J. Z. Certified adversarial robustness via randomized smoothing. *International Conference on Machine Learning*, 2019.
- Dang-Nhu, R., Singh, G., Bielik, P., and Vechev, M. Adversarial attacks on probabilistic autoregressive forecasting models. *International Conference on Machine Learning*, 2020.
- Davis, J. and Goadrich, M. The relationship between precision-recall and roc curves. *International Conference on Machine Learning*, 2006.
- Dua, D. and Graff, C. UCI machine learning repository. *University of California*, 2017.
- Gal, Y. and Ghahramani, Z. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. *International Conference on Machine Learning*, 2016.
- Goodfellow, I. J., Shlens, J., and Szegedy, C. Explaining and harnessing adversarial examples. *International Conference on Learning Representations*, 2015.
- Gu, S. and Rigazio, L. Towards deep neural network architectures robust to adversarial examples. *International Conference on Learning Representations*, 2015.
- Kopetzki, A.-K. and Günnemann, S. Reachable sets of classifiers and regression models: (non-)robustness analysis and robust training. *Machine Learning Journal*, 2021.
- Krizhevsky, A., Nair, V., and Hinton, G. Cifar-10. *Canadian Institute for Advanced Research*, 2009.
- Kumar, A., Levine, A., Feizi, S., and Goldstein, T. Certifying confidence via randomized smoothing. *Neural Information Processing Systems*, 2020.
- Lakshminarayanan, B., Pritzel, A., and Blundell, C. Simple and scalable predictive uncertainty estimation using deep ensembles. *Neural Information Processing Systems*, 2017.
- LeCun, Y. and Cortes, C. MNIST handwritten digit database. *National Institute of Standards and Technology*, 2010.
- Lee, K., Lee, K., Lee, H., and Shin, J. A simple unified framework for detecting out-of-distribution samples and adversarial attacks. *Neural Information Processing Sys*tems, 2018.

- Maddox, W. J., Izmailov, P., Garipov, T., Vetrov, D. P., and Wilson, A. G. A simple baseline for bayesian uncertainty in deep learning. *Neural Information Processing Systems*, 2019.
- Madry, A., Makelov, A., Schmidt, L., Tsipras, D., and Vladu, A. Towards deep learning models resistant to adversarial attacks. *International Conference on Learn*ing Representations, 2018.
- Malinin, A. and Gales, M. Predictive uncertainty estimation via prior networks. *Neural Information Processing Systems*, 2018a.
- Malinin, A. and Gales, M. Prior networks for detection of adversarial attacks. *arXiv:1812.02575 [stat.ML]*, 2018b.
- Malinin, A. and Gales, M. Reverse kl-divergence training of prior networks: Improved uncertainty and adversarial robustness. *Neural Information Processing Systems*, 2019.
- Malinin, A., Mlodozeniec, B., and Gales, M. Ensemble distribution distillation. *International Conference on Learn*ing Representations, 2019.
- Meinke, A. and Hein, M. Towards neural networks that provably know when they don't know. *International Conference on Learning Representations*, 2020.
- Nandy, J., Hsu, W., and Lee, M. Towards maximizing the representation gap between in-domain & out-of-distribution examples. *International Conference on Machine Learning, Uncertainty and Robustness in Deep Learning Workshop*, 2020.
- Netzer, Y., Wang, T., Coates, A., Bissacco, A., Wu, B., and Ng, A. Y. Reading digits in natural images with unsupervised feature learning. *Neural Information Processing Systems, Deep Learning and Unsupervised Feature Learning Workshop*, 2011.
- Osawa, K., Swaroop, S., Khan, M. E. E., Jain, A., Eschenhagen, R., Turner, R. E., and Yokota, R. Practical deep learning with bayesian principles. *Neural Information Processing Systems*, 2019.
- Ovadia, Y., Fertig, E., Ren, J., Nado, Z., Sculley, D., Nowozin, S., Dillon, J. V., Lakshminarayanan, B., and Snoek, J. Can you trust your model's uncertainty? evaluating predictive uncertainty under dataset shift. *Neural Information Processing Systems*, 2019.
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J.,
 Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga,
 L., Desmaison, A., Kopf, A., Yang, E., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy, S., Steiner, B., Fang,
 L., Bai, J., and Chintala, S. Pytorch: An imperative

- style, high-performance deep learning library. *Neural Information Processing Systems*, 2019.
- Qin, Y., Wang, X., Beutel, A., and Chi, E. H. Improving uncertainty estimates through the relationship with adversarial robustness. *arXiv:2006.16375 [cs.LG]*, 2020.
- Saito, T. and Rehmsmeier, M. The precision-recall plot is more informative than the roc plot when evaluating binary classifiers on imbalanced datasets. *PLOS ONE*, 2015.
- Schuchardt, J., Bojchevski, A., Klicpera, J., and Günnemann, S. Collective robustness certificates: Exploiting interdependence in graph neural networks. *International Conference on Learning Representations*, 2021.
- Sensoy, M., Kaplan, L., and Kandemir, M. Evidential deep learning to quantify classification uncertainty. *Neural Information Processing Systems*, 2018.
- Sensoy, M., Kaplan, L., Cerutti, F., and Saleki, M. Uncertainty-aware deep classifiers using generative models. *AAAI Conference on Artificial Intelligence*, 2020.
- Shi, W., Zhao, X., Chen, F., and Yu, Q. Multifaceted uncertainty estimation for label-efficient deep learning. *Neural Information Processing Systems*, 2020.
- Simonyan, K. and Zisserman, A. Very deep convolutional networks for large-scale image recognition. *International Conference on Learning Representations*, 2015.
- Singh, G., Ganvir, R., Püschel, M., and Vechev, M. Beyond the single neuron convex barrier for neural network certification. *Neural Information Processing Systems*, 2019.
- Smith, L. and Gal, Y. Understanding measures of uncertainty for adversarial example detection. *Uncertainty in Artificial Intelligence*, 2018.
- Stutz, D., Hein, M., and Schiele, B. Confidence-calibrated adversarial training: Generalizing to unseen attacks. *International Conference on Machine Learning*, 2020.
- Szegedy, C., Zaremba, W., Sutskever, I., Bruna, J., Erhan, D., Goodfellow, I., and Fergus, R. Intriguing properties of neural networks. *International Conference on Learning Representations*, 2014.
- Tagasovska, N. and Lopez-Paz, D. Single-model uncertainties for deep learning. *Neural Information Processing Systems*, 2019.
- Weng, T.-W., Zhang, H., Chen, P.-Y., Yi, J., Su, D., Gao, Y., Hsieh, C.-J., and Daniel, L. Evaluating the robustness of neural networks: An extreme value theory approach. *International Conference on Learning Representations*, 2018.

- Wicker, M., Laurenti, L., Patane, A., and Kwiatkowska, M. Probabilistic safety for bayesian neural networks. *Uncertainty in Artificial Intelligence*, 2020.
- Wong, E. and Kolter, J. Z. Provable defenses against adversarial examples via the convex outer adversarial polytope. *International Conference on Machine Learning*, 2018.
- Xiao, H., Rasul, K., and Vollgraf, R. Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms. *Zalando SE*, 2017.
- yeh Chiang, P., Curry, M. J., Abdelkader, A., Kumar, A., Dickerson, J., and Goldstein, T. Detection as regression: Certified object detection by median smoothing. *Neural Information Processing Systems*, 2020.
- Zhao, X., Chen, F., Hu, S., and Cho, J.-H. Uncertainty aware semi-supervised learning on graph data. *Neural Information Processing Systems*, 2020.
- Zheng, S., Song, Y., Leung, T., and Goodfellow, I. Improving the robustness of deep neural networks via stability training. *Conference on Computer Vision and Pattern Recognition*, 2016.
- Zügner, D., Akbarnejad, A., and Günnemann, S. Adversarial attacks on neural networks for graph data. *ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2018.

6. Appendix

6.1. Closed-form computation of uncertainty measures & Uncertainty attacks

Dirichlet-based uncertainty models allow to compute several uncertainty measures in closed form (see (Malinin & Gales, 2018a) for a derivation). As proposed by (Malinin & Gales, 2018a), we use precision m_{α_0} , differential entropy $m_{\rm diffE}$ and mutual information $m_{\rm MI}$ to estimate uncertainty on predictions.

The differential entropy $m_{\rm diffE}$ of a DBU model reaches its maximum value for equally probable categorical distributions and thus, a on flat Dirichlet distribution. It is a measure for distributional uncertainty and expected to be low on ID data, but high on OOD data.

$$m_{\text{diffE}} = \sum_{c}^{K} \ln \Gamma(\alpha_c) - \ln \Gamma(\alpha_0)$$

$$- \sum_{c}^{K} (\alpha_c - 1) \cdot (\Psi(\alpha_c) - \Psi(\alpha_0))$$
(7)

where α are the parameters of the Dirichlet-distribution, Γ is the Gamma function and Ψ is the Digamma function.

The mutual information $m_{\rm MI}$ is the difference between the total uncertainty (entropy of the expected distribution) and the expected uncertainty on the data (expected entropy of the distribution). This uncertainty is expected to be low on ID data and high on OOD data.

$$m_{\rm MI} = -\sum_{c=1}^{K} \frac{\alpha_c}{\alpha_0} \left(\ln \frac{\alpha_c}{\alpha_0} - \Psi(\alpha_c + 1) + \Psi(\alpha_0 + 1) \right)$$
(8)

Furthermore, we use the precision α_0 to measure uncertainty, which is expected to be high on ID data and low on OOD data.

$$m_{\alpha_0} = \alpha_0 = \sum_{c=1}^K \alpha_c \tag{9}$$

As these uncertainty measures are computed in closed form and it is possible to obtain their gradients, we use them (i.e. $m_{\rm diffE}$, $m_{\rm MI}$, m_{α_0}) are target function of our uncertainty attacks. Changing the attacked target function allows us to use a wide range of gradient-based attacks such as FGSM attacks, PGD attacks, but also more sophisticated attacks such as Carlini-Wagner attacks.

6.2. Details of the Experimental setup

Models. We trained all models with a similar based architecture. We used namely 3 linear layers for vector data sets, 3 convolutional layers with size of 5 + 3 linear layers for

MNIST and the VGG16 (Simonyan & Zisserman, 2015) architecture with batch normalization for CIFAR10. All the implementation are performed using Pytorch (Paszke et al., 2019). We optimized all models using Adam optimizer. We performed early stopping by checking for loss improvement every 2 epochs and a patience of 10. The models were trained on GPUs (1 TB SSD).

We performed a grid-search for hyper-parameters for all models. The learning rate grid search was done in $[1e^{-5}, 1e^{-3}]$. For PostNet, we used Radial Flows with a depth of 6 and a latent space equal to 6. Further, we performed a grid search for the regularizing factor in $[1e^{-7}, 1e^{-4}]$. For PriorNet, we performed a grid search for the OOD loss weight in [1, 10]. For DDNet, we distilled the knowledge of 5 neural networks after a grid search in [2, 5, 10, 20] neural networks. Note that it already implied a significant overhead at training compare to other models.

Metrics. For all experiments, we focused on using AUC-PR scores since it is well suited to imbalance tasks (Saito & Rehmsmeier, 2015) while bringing theoretically similar information than AUC-ROC scores (Davis & Goadrich, 2006). We scaled all scores from [0, 1] to [0, 100]. All results are average over 5 training runs using the best hyper-parameters found after the grid search.

Data sets. For vector data sets, we use 5 different random splits to train all models. We split the data in training, validation and test sets (60%, 20%, 20%).

We use the segment vector data set (Dua & Graff, 2017), where the goal is to classify areas of images into 7 classes (window, foliage, grass, brickface, path, cement, sky). We remove class window from ID training data to provide OOD training data to PriorNet. Further, We remove the class 'sky' from training and instead use it as the OOD data set for OOD detection experiments. Each input is composed of 18 attributes describing the image area. The data set contains 2, 310 samples in total.

We further use the Sensorless Drive vector data set (Dua & Graff, 2017), where the goal is to classify extracted motor current measurements into 11 different classes. We remove class 9 from ID training data to provide OOD training data to PriorNet. We remove classes 10 and 11 from training and use them as the OOD dataset for OOD detection experiments. Each input is composed of 49 attributes describing motor behaviour. The data set contains 58, 509 samples in total.

Additionally, we use the MNIST image data set (LeCun & Cortes, 2010) where the goal is to classify pictures of hand-drawn digits into 10 classes (from digit 0 to digit 9). Each input is composed of a $1 \times 28 \times 28$ tensor. The data set contains 70,000 samples. For OOD detection experiments, we use FashionMNIST (Xiao et al., 2017) and KMNIST

(Clanuwat et al., 2018) containing images of Japanese characters and images of clothes, respectively. FashionMNIST was used as training OOD for PriorNet while KMNIST is used as OOD at test time.

Finally, we use the CIFAR10 image data set (Krizhevsky et al., 2009) where the goal is to classify a picture of objects into 10 classes (airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck). Each input is a $3\times32\times32$ tensor. The data set contains 60,000 samples. For OOD detection experiments, we use street view house numbers (SVHN) (Netzer et al., 2011) and CIFAR100 (Krizhevsky et al., 2009) containing images of numbers and objects respectively. CIFAR100 was used as training OOD for PriorNet while SVHN is used as OOD at test time.

Perturbations. For all label and uncertainty attacks, we used Fast Gradient Sign Methods and Project Gradient Descent. We tried 6 different attack radii [0.0, 0.1, 0.2, 0.5, 1.0, 2.0, 4.0]. These radii operate on the input space after data normalization. We bound perturbations by L_{∞} -norm or by L_{2} -norm, with

$$L_{\infty}(x) = \max_{i=1,\dots,D} |x_i|$$
 and $L_2(x) = (\sum_{i=1}^D x_i^2)^{0.5}$. (10)

For L_{∞} -norm it is obvious how to relate perturbation size ε with perturbed input images, because all inputs are standardized such that the values of their features are between 0 and 1. A perturbation of size $\varepsilon=0$ corresponds to the original input, while a perturbation of size $\varepsilon=1$ corresponds to the whole input space and allows to change all features to any value.

For L_2 -norm the relation between perturbation size ε and perturbed input images is less obvious. To justify our choice for ε w.r.t. this norm, we relate perturbations size ε_2 corresponding to L_2 -norm with perturbations size ε_∞ corresponding to L_∞ -norm. First, we compute ε_2 , such that the L_2 -norm is the smallest super-set of the L_∞ -norm. Let us consider a perturbation of ε_∞ . The largest L_2 -norm would be obtained if each feature is perturbed by ε_∞ . Thus, perturbation ε_2 , such that L_2 encloses L_∞ is $\varepsilon_2 = (\sum_{i=1}^D \varepsilon_\infty^2)^{0.5} = \sqrt{D}\varepsilon_\infty$. For the MNIST-data set, with $D = 28 \times 28$ input features L_2 -norm with $\varepsilon_2 = 28$ encloses L_∞ -norm with $\varepsilon_\infty = 1$.

Alternatively, ε_2 can be computes such that the volume spanned by L_2 -norm is equivalent to the one spanned by L_∞ -norm. Using that the volume spanned by L_∞ -norm is ε_∞^D and the volume spanned by L_2 -norm is $\frac{\pi^{0.5D}\varepsilon_2^D}{\Gamma(0.5D+1)}$ (where Γ is the Gamma-function), we obtain volume equivalence if $\varepsilon_2 = \Gamma(0.5D+1)^{\frac{1}{D}}\sqrt{\pi}\varepsilon_\infty$. For the MNIST-data set, with $D=28\times28$ input features L_2 -norm with $\varepsilon_2\approx21.39$ is volume equivalent to L_∞ -norm with $\varepsilon_\infty=1$.

6.3. Additional Experiments

Table 8 and 9 illustrate that no DBU model maintains high accuracy under gradient-based label attacks. Accuracy under PGD attacks decreases more than under FGSM attacks, since PGD is stronger. Interestingly Noise attacks achieve also good performances with increasing Noise standard deviation. Note that the attack is not constraint to be with a given radius for Noise attacks.

Table 8. Accuracy under PGD label attacks.

Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
				MNIST	1					Cl	FAR10			
PostNet	99.4	99.2	98.8	96.8	89.6	53.8	13.0	89.5	73.5	51.7	13.2	2.2	0.8	0.3
PriorNet	99.3	99.1	98.8	97.4	93.9	75.3	4.8	88.2	77.8	68.4	54.0	37.9	17.5	5.1
DDNet	99.4	99.1	98.8	97.5	91.6	48.8	0.2	86.1	73.9	59.1	20.5	1.5	0.0	0.0
EvNet	99.2	98.9	98.4	96.8	92.4	73.1	40.9	89.8	71.7	48.8	11.5	2.7	1.5	0.4
			S	ensorle	ss			Segment						
PostNet	98.3	13.1	6.4	4.0	7.0	9.8	11.3	98.9	82.8	50.1	19.2	8.8	5.1	8.6
PriorNet	99.3	16.5	5.6	1.2	0.4	0.2	1.6	99.5	90.7	47.6	7.8	0.2	0.0	0.4
DDNet	99.3	12.4	2.4	0.6	0.3	0.1	0.1	99.2	90.8	45.7	6.9	0.0	0.0	0.0
EvNet	99.0	35.3	22.3	11.2	7.0	5.2	4.0	99.3	91.8	54.0	10.3	0.8	0.5	0.6

Table 9. Accuracy under FGSM label attacks.

							-							
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
				MNIST						C	IFAR10)		
PostNet	99.4	99.2	98.9	97.7	95.2	90.1	79.2	89.5	72.3	54.9	31.2	21.0	16.8	15.6
PriorNet	99.3	99.1	98.9	97.7	95.8	93.2	76.7	88.2	77.3	70.1	59.4	52.3	48.5	46.8
DDNet	99.4	99.2	98.9	97.8	94.7	79.2	25.2	86.1	73.0	60.2	32.5	14.6	7.1	6.0
EvNet	99.2	98.9	98.6	97.6	95.8	90.1	74.4	89.8	71.4	54.5	29.6	18.1	14.4	13.4
			S	ensorle	ss			Segment						
PostNet	98.3	19.6	10.9	10.9	11.9	12.4	12.5	98.9	79.6	57.3	31.5	18.4	20.6	19.9
PriorNet	99.3	24.7	11.8	8.6	8.5	8.1	8.3	99.5	85.5	40.5	8.9	0.4	0.3	0.2
DDNet	99.3	18.0	8.2	6.5	5.4	6.7	7.8	99.2	86.4	36.2	11.9	0.9	0.0	0.0
EvNet	99.0	42.0	28.0	17.5	13.7	13.6	14.9	99.3	90.6	55.2	14.2	2.4	0.5	0.1

Table 10. Accuracy under Noise label attacks.

Noise Std	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
				MNIST	1					С	IFAR10)		
PostNet	99.4	98.6	91.8	14.9	1.3	0.1	0.0	91.7	21.5	10.1	0.1	1.2	0.0	1.9
PriorNet	99.3	98.5	95.7	14.4	0.0	0.0	0.0	87.7	28.1	11.2	9.7	5.0	8.5	9.0
DDNet	99.4	98.6	92.4	13.3	0.7	0.0	0.0	81.7	23.0	11.2	11.2	11.0	7.8	6.7
EvNet	99.3	96.9	81.6	11.7	0.5	0.0	0.0	89.5	20.7	11.1	5.2	0.5	2.3	3.9
			S	ensorle	ss					S	egment			
PostNet	98.1	0.1	3.7	11.7	11.7	11.7	11.7	98.5	39.4	3.9	1.8	12.1	20.3	22.1
PriorNet	99.3	0.2	0.0	0.0	0.0	0.3	2.4	99.4	47.9	8.8	0.0	0.0	0.0	0.0
DDNet	99.0	0.4	0.1	0.0	0.0	0.0	0.0	99.1	50.0	10.3	0.0	0.0	0.3	0.0
EvNet	98.6	0.2	0.0	0.1	1.4	4.6	8.8	99.1	50.3	10.3	1.2	0.3	0.0	1.5

6.3.1. Uncertainty estimation under label attacks

Is low uncertainty a reliable indicator of correct predictions?

On non-perturbed data uncertainty estimates are an indicator of correctly classified samples, but if the input data is perturbed none of the DBU models maintains its high performance. Thus, uncertainty estimates are not a robust indicator of correctly labeled inputs.

Table 2, 11, 12, and 13 illustrate that neither differential entropy nor precision, nor mutual information are a reliable indicator of correct predictions under PGD attacks. DBU-models achieve significantly better results when they are attacked by FGSM-attacks (Table 14), but as FGSM attacks provide much weaker adversarial examples than PGD attacks, this cannot be seen as real advantage.

Table 11. Distinguishing between correctly and wrongly predicted labels based on the differential entropy under PGD label attacks (AUC-PR).

			N	INIST						Seg	gment			
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
PostNet	99.9	99.9	99.8	98.7	89.5	43.5	9.0	99.9	77.6	31.6	11.1	5.3	4.4	8.7
PriorNet	99.9	99.8	99.6	97.7	90.5	69.1	6.4	100.0	96.8	44.5	4.5	0.4	0.0	15.2
DDNet	100.0	100.0	99.9	99.7	97.6	50.2	0.1	100.0	96.8	54.0	4.3	0.0	0.0	0.0
EvNet	99.6	99.3	98.7	96.1	88.8	63.1	31.7	100.0	95.9	44.3	5.9	0.8	0.6	0.7

Table 12. Distinguishing between correctly and wrongly predicted labels based on the precision α_0 under PGD label attacks (AUC-PR).

Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
			M	INIST						CII	FAR10			
PostNet	100.0	99.9	99.7	98.2	87.9	39.1	6.9	98.7	88.6	56.2	7.8	1.2	0.4	0.3
PriorNet	99.9	99.8	99.6	97.7	90.4	69.1	6.6	92.9	77.7	60.5	37.6	24.9	11.3	3.0
DDNet	100.0	100.0	100.0	99.8	98.2	51.1	0.1	97.6	91.8	78.3	18.1	0.8	0.0	0.0
EvNet	99.6	99.2	98.6	95.7	88.6	63.6	32.6	97.9	85.9	57.2	10.2	4.0	2.4	0.3
			Ser	isorless						Se	gment			
PostNet	99.6	7.0	3.3	3.1	6.9	9.8	11.3	99.9	74.2	31.6	11.1	5.0	4.2	8.6
PriorNet	99.8	10.5	3.2	0.6	0.2	0.2	1.8	100.0	96.9	45.2	4.4	0.4	0.0	1.2
DDNet	99.8	8.7	1.3	0.3	0.2	0.1	0.2	100.0	97.1	45.0	4.1	0.0	0.0	0.0
EvNet	99.9	23.2	13.2	6.0	3.7	2.7	2.1	100.0	95.7	44.5	5.9	0.8	0.6	0.7

Table 13. Distinguishing between correctly and wrongly predicted labels based on the mutual information under PGD label attacks (AUC-PR).

Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
]	MNIST						CI	FAR10			
PostNet	99.7	99.7	99.6	99.2	92.4	40.0	6.9	97.3	84.5	56.2	12.2	2.4	0.7	0.3
PriorNet	99.9	99.8	99.6	97.7	90.3	68.9	6.4	82.7	65.6	51.4	35.5	24.4	11.0	2.9
DDNet	100.0	99.9	99.9	99.7	97.4	50.2	0.1	96.9	90.8	77.2	18.8	0.8	0.0	0.0
EvNet	97.8	97.0	95.7	92.6	86.1	62.3	28.9	91.3	72.4	47.9	11.4	1.6	0.9	1.6
			Se	ensorles	s					Se	gment			
PostNet	99.3	7.0	3.3	3.3	7.0	9.8	11.3	99.9	73.2	31.5	11.1	5.0	4.3	8.7
PriorNet	99.8	10.5	3.2	0.6	0.2	0.1	11.8	100.0	96.6	45.2	4.5	0.4	0.0	1.1
DDNet	99.6	8.6	1.3	0.3	0.2	0.1	0.1	100.0	96.5	42.4	4.1	0.0	0.0	0.0
EvNet	99.1	22.0	12.6	5.9	3.7	2.7	2.2	100.0	90.5	41.0	5.9	0.8	0.6	0.7

Table 14. Distinguishing between correctly and wrongly predicted labels based on the differential entropy under FGSM label attacks (AUC-PR).

Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
			N	MNIST						Cl	FAR10			
PostNet	99.9	99.9	99.8	99.4	97.8	92.1	83.2	98.5	88.7	68.9	31.0	18.6	15.5	16.7
PriorNet	99.9	99.9	99.7	98.3	94.1	88.5	78.6	90.1	73.6	61.6	46.1	38.5	35.6	37.3
DDNet	100.0	100.0	99.9	99.8	98.7	86.4	23.0	97.3	90.6	78.7	39.4	13.7	6.0	5.1
EvNet	99.6	99.4	99.1	97.8	95.8	90.4	76.8	98.0	86.2	67.4	32.7	19.9	18.2	19.7
			Se	nsorless	3					S	egment			
PostNet	99.7	11.7	7.3	9.3	11.8	12.5	12.5	99.9	73.6	40.6	23.7	17.2	19.8	20.2
PriorNet	99.8	21.4	10.4	8.5	9.0	9.2	10.3	100.0	93.7	37.7	5.8	1.1	0.9	0.8
DDNet	99.7	18.5	5.4	4.3	4.2	5.7	7.9	100.0	94.1	42.9	7.2	1.0	0.0	0.0
EvNet	99.9	44.8	29.2	18.2	15.1	14.9	15.5	100.0	93.7	48.7	8.7	2.4	1.6	0.5

Table 15. Distinguishing between correctly and wrongly predicted labels based on the differential entropy under Noise label attacks (AUC-PR).

Noise Std	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
]	MNIST						CI	FAR10			
PostNet	99.9	99.8	99.6	74.2	7.4	0.2	0.0	98.7	76.3	24.3	0.4	4.9	0.0	1.7
PriorNet	99.9	99.9	99.8	73.4	0.0	0.0	0.0	85.0	27.8	15.9	20.4	7.0	7.7	8.3
DDNet	100.0	99.9	99.4	51.1	0.6	0.1	0.0	96.1	61.0	39.8	14.2	11.3	6.9	6.9
EvNet	99.5	98.4	88.5	20.2	0.9	0.0	0.0	97.5	66.1	21.4	7.7	2.3	3.0	3.8
			Se	ensorles	s					Se	egment			
PostNet	99.7	0.3	3.2	13.3	12.0	11.7	11.7	99.9	53.9	4.8	1.8	11.2	21.7	21.6
PriorNet	100.0	0.3	0.0	0.0	0.0	7.8	11.5	100.0	84.5	15.6	0.0	0.0	0.0	0.0
DDNet	99.7	0.9	0.6	0.0	0.0	0.0	0.0	100.0	82.7	23.9	0.0	0.0	0.6	0.0
EvNet	99.8	0.3	0.0	0.1	1.7	5.5	10.0	100.0	78.3	19.0	3.5	0.5	0.0	1.7

Can we use uncertainty estimates to detect attacks against the class prediction?

PGD attacks do not explicitly consider uncertainty during the computation of adversarial examples, but they seem to provide perturbed inputs with similar uncertainty as the original input.

FGSM and Noise attacks are easier to detect, but also weaker thand PGD attacks. This suggests that DBU models are capable of detecting weak attacks by using uncertainty estimation.

Table 16 Attack	Detection based	on differential	entropy under PG	D label attacks ((ALIC-PR)
Table 10. Attack-	Detection Dasco	. On unicicilia	CHILODY MINGELL (iz iadei allaeks i	AUC-FIN.

			MN	IST					Segn	nent		
Att. Rad.	0.1	0.2	0.5	1.0	2.0	4.0	0.1	0.2	0.5	1.0	2.0	4.0
PostNet PriorNet DDNet	57.7 67.7 53.4	66.3 83.2 57.1	83.4 97.1 68.5	90.5 96.7 83.9	79.0 92.1 96.0	50.1 82.9 86.3	95.6 86.7 76.1	73.5 83.3 83.5	47.0 38.0 45.4	42.3 31.3 32.4	53.4 30.8 30.8	82.7 31.5 30.8
EvNet	54.8	59.0	68.5	75.9	72.6	59.8	94.9	80.9	41.5	32.5	31.1	31.1

Table 17. Attack-Detection based on precision α_0 under PGD label attacks (AUC-PR).

Att. Rad.	0.1	0.2	0.5	1.0	2.0	4.0	0.1	0.2	0.5	1.0	2.0	4.0
			MN	IST					CIFA	R10		
PostNet	63.3	75.7	92.6	95.1	75.3	39.5	63.4	66.9	42.1	32.9	31.6	31.2
PriorNet	67.6	83.2	97.1	96.9	92.7	84.7	53.3	56.0	55.6	49.2	42.2	35.4
DDNet	52.7	55.7	64.7	78.4	91.9	80.9	55.8	60.5	57.3	38.7	32.3	31.4
EvNet	49.1	48.0	45.1	42.7	41.8	39.2	48.4	46.9	46.3	46.3	44.5	42.5
			Sense	orless					Segn	nent		
PostNet	39.8	35.8	35.4	52.0	88.2	99.0	94.6	70.3	46.3	42.6	54.9	84.0
PriorNet	40.9	35.1	32.0	31.1	30.7	30.7	82.7	82.6	39.4	31.6	30.8	30.8
DDNet	47.7	40.3	35.3	32.8	31.3	30.8	80.0	86.0	43.3	33.6	31.0	30.8
EvNet	45.4	39.7	36.1	34.8	34.7	36.0	90.9	72.4	40.4	32.4	31.1	31.1

Table 18. Attack-Detection based on mutual information under PGD label attacks (AUC-PR).

Att. Rad.	0.1	0.2	0.5	1.0	2.0	4.0	0.1	0.2	0.5	1.0	2.0	4.0
			MN	IST					CIFA	R10		
PostNet	42.2	37.5	36.7	54.5	70.5	70.3	52.2	52.1	50.0	65.9	76.3	80.7
PriorNet	67.7	83.3	97.1	96.9	92.6	84.5	54.0	56.9	56.3	49.7	42.4	35.5
DDNet	53.1	56.3	66.5	81.0	94.0	82.9	56.0	60.8	57.4	38.2	32.1	31.3
EvNet	49.1	48.0	45.2	42.9	41.9	39.3	48.7	47.3	46.3	46.0	44.1	42.2
			Sense	orless					Segn	nent		
PostNet	75.3	76.6	66.5	57.7	85.6	98.7	94.8	73.5	55.9	47.9	58.0	84.0
PriorNet	40.7	35.0	32.0	31.0	30.7	30.7	83.5	82.7	39.2	31.6	30.8	30.8
DDNet	48.0	40.0	35.2	32.6	31.2	30.8	82.4	88.1	43.4	33.4	30.9	30.8
EvNet	45.5	39.7	36.1	34.8	34.7	36.0	91.7	72.9	40.5	32.4	31.1	31.1

Table 19. Attack-Detection based on differential entropy under FGSM label attacks (AUC-PR).

Att. Rad.	0.1	0.2	0.5	1.0	2.0	4.0		0.1	0.2	0.5	1.0	2.0	4.0
			MN	IST			1			CIF	AR10		
PostNet	55.9	61.8	74.8	84.0	88.9	89.9		62.1	67.2	65.7	63.1	65.4	73.8
PriorNet	67.4	82.4	96.9	98.3	98.9	99.6		58.4	63.1	68.5	70.1	68.5	62.5
DDNet	53.6	57.3	68.3	82.6	95.6	98.7		57.2	62.9	69.1	68.7	69.7	76.5
EvNet	54.1	57.4	63.8	67.6	68.6	69.9		57.8	61.7	63.3	62.9	65.7	72.5
			Sens	orless						Seg	ment		
PostNet	98.4	99.8	99.9	99.9	99.9	99.9		96.9	93.9	99.5	99.9	100.0	100.0
PriorNet	48.7	38.6	32.7	32.9	38.6	44.3	İ	89.0	80.8	46.7	37.2	33.7	32.4
DDNet	61.5	47.8	37.1	33.1	32.4	33.2		79.6	86.2	60.2	47.5	36.6	31.6
EvNet	67.3	65.5	72.3	73.4	75.3	79.1		95.7	87.2	59.3	51.7	51.1	53.5

Table 20. Attack-Detection based on differential entropy under Noise label attacks (AUC-PR).

							1 2				`	,	
Noise Std.	0.1	0.2	0.5	1.0	2.0	4.0		0.1	0.2	0.5	1.0	2.0	4.0
			Mì	NIST						CIF	AR10		
PostNet	51.3	65.3	93.8	95.1	95.2	95.2		80.8	84.5	97.6	99.5	99.3	98.2
PriorNet	32.5	36.8	88.9	99.6	99.7	92.7		34.7	32.3	34.3	60.3	95.5	100.0
DDNet	60.7	87.6	99.8	100.0	99.9	99.8		59.1	62.6	81.5	98.6	99.8	98.7
EvNet	51.2	55.7	66.9	70.3	68.0	67.1		75.7	78.6	88.2	97.8	96.4	95.6
			Sens	orless						Seg	ment		
PostNet	99.8	100.0	100.0	100.0	100.0	100.0		95.6	99.4	100.0	100.0	100.0	100.0
PriorNet	42.0	33.8	31.5	34.7	43.7	47.0		56.7	56.7	39.8	33.7	31.9	33.7
DDNet	53.4	43.5	34.3	31.6	32.5	36.1		57.0	58.9	43.1	33.7	31.5	31.3
EvNet	67.1	78.8	88.3	95.4	96.9	97.8		60.8	63.5	61.2	64.8	73.7	85.2

6.3.2. Attacking uncertainty estimation

Are uncertainty estimates a robust feature for OOD detection?

Using uncertainty estimation to distinguish between ID and OOD data is not robust as shown in the following tables.

Table 21. OOD detection based on differential entropy under PGD uncertainty attacks against differential entropy on ID data and OOD data (AUC-PR).

		ID-	Attack (non-atta	cked O	OD)				OOL)-Attack	(non-a	ttacked	ID)	
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0		0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIST	– K	MNIST						
PostNet	94.5	94.1	93.9	91.1	77.1	44.0	31.9		94.5	93.1	91.4	82.1	62.2	50.7	48.8
PriorNet	99.6	99.4	99.1	97.8	93.8	77.6	32.0		99.6	99.4	99.1	98.0	94.6	85.5	73.9
DDNet	99.3	99.1	98.9	97.8	93.5	63.3	30.7		99.3	99.1	99.0	98.3	96.7	91.3	73.8
EvNet	69.0	67.1	65.6	61.8	57.4	50.9	43.6		69.0	55.8	48.0	39.4	36.2	34.9	34.4
							Seg. – Se	eg. c	lass sky						
PostNet	99.0	80.7	53.5	38.0	34.0	41.6	49.5		99.0	88.4	69.2	45.1	36.4	42.6	75.4
PriorNet	34.8	31.4	30.9	30.8	30.8	30.8	30.8		34.8	31.8	31.0	30.8	30.8	30.8	32.1
DDNet	31.5	30.9	30.8	30.8	30.8	30.8	30.8		31.5	31.0	30.8	30.8	30.8	30.8	30.8
EvNet	92.5	67.2	43.2	31.6	30.9	30.9	31.2		92.5	86.1	82.7	48.9	32.7	30.9	30.9

Table 22. OOD detection under PGD uncertainty attacks against differential entropy on ID data and OOD data (AUC-ROC).

		ID-	Attack (non-atta	icked O	OD)				OOL	-Attack	(non-a	ttacked	ID)	
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0		0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIST -	K	MNIST						
PostNet	91.6	91.3	91.9	91.5	80.2	38.8	9.2		91.6	90.4	89.0	81.6	62.6	45.0	43.1
PriorNet	99.8	99.7	99.5	99.0	97.1	81.1	8.7		99.8	99.7	99.6	99.1	97.7	93.0	84.9
DDNet	99.2	98.9	98.6	97.3	92.1	58.2	1.2		99.2	99.0	98.8	97.9	95.8	89.1	69.3
EvNet	81.2	79.6	78.2	74.6	69.5	58.7	43.0		81.2	67.2	54.8	35.4	25.5	20.7	18.5
							CIFAR10) –	SVHN						_
PostNet	87.0	71.9	56.3	30.2	20.2	15.0	9.7		87.0	71.0	54.3	33.5	30.3	26.2	19.4
PriorNet	62.4	48.2	35.9	13.8	3.6	0.9	0.3		62.4	48.0	35.6	14.8	6.6	3.4	1.6
DDNet	87.0	76.0	63.6	29.3	6.1	1.1	0.4		87.0	78.1	66.1	26.2	5.1	0.7	0.1
EvNet	88.0	69.1	51.7	24.6	15.5	9.5	4.2		88.0	72.0	60.7	47.9	42.1	33.3	24.0
						Se	ns. – Sens	s. (class 10, 1	1					
PostNet	85.3	49.1	38.1	7.8	8.2	8.2	8.2		85.3	57.2	54.0	27.3	31.5	86.7	99.5
PriorNet	28.1	0.8	0.3	0.4	1.6	8.4	26.8		28.1	2.5	0.7	0.2	2.3	18.9	41.0
DDNet	21.0	3.0	0.9	0.4	0.6	2.1	7.3		21.0	4.4	2.1	1.9	2.2	2.2	4.1
EvNet	74.2	21.4	12.2	4.3	1.4	0.6	0.3		74.2	45.3	38.5	19.6	9.6	12.1	26.0
							Seg. – Seg	g. (class sky						
PostNet	99.2	84.7	55.5	23.0	9.7	4.4	4.7		99.2	92.1	77.1	41.5	24.9	41.0	80.8
PriorNet	17.1	4.4	1.3	0.0	0.0	0.0	0.1		17.1	5.9	1.5	0.1	0.0	0.1	5.8
DDNet	4.1	1.1	0.0	0.0	0.0	0.0	0.0		4.1	1.8	0.4	0.0	0.0	0.0	0.0
EvNet	91.2	54.5	23.3	3.9	0.9	0.4	0.2		91.2	82.9	76.4	42.2	9.7	0.8	0.6

Table 23. OOD detection (AU-PR) under PGD uncertainty attacks against precision α_0 on ID data and OOD data.

		ID-	Attack (non-atta	acked O	OD)				OOL)-Attack	k (non-a	ttacked	ID)	
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0		0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIST -	-]	KMNIST						
PostNet	98.4	97.4	96.0	88.8	70.9	39.3	31.3		98.4	97.2	95.2	82.8	52.6	34.3	32.1
PriorNet	99.6	99.5	99.2	98.0	94.1	76.0	31.1		99.6	99.5	99.2	98.2	95.3	87.5	75.6
DDNet	97.2	96.7	96.1	93.8	86.4	53.2	31.0		97.2	96.7	96.2	94.5	91.1	82.9	64.6
EvNet	39.8	39.2	38.8	37.9	37.1	36.3	35.4		39.8	34.5	32.5	31.2	31.0	30.9	31.0
							CIFAR1	0	– SVHN						
PostNet	82.4	63.8	46.1	22.3	17.4	16.7	16.4		82.4	61.8	41.5	21.8	19.8	17.5	15.8
PriorNet	37.9	25.0	19.2	15.8	15.4	15.4	15.4	İ	37.9	25.9	19.4	15.6	15.4	15.4	15.4
DDNet	81.1	70.1	58.4	30.0	16.7	15.5	15.4		81.1	71.2	59.9	27.8	16.5	15.5	15.4
EvNet	34.7	27.4	25.4	22.0	19.7	18.1	17.1		34.7	19.4	18.1	17.1	16.8	16.2	15.7
						Se	ns. – Sen	s.	class 10, 1	1					
PostNet	77.4	39.6	35.9	31.7	44.4	44.4	44.4		77.4	40.3	38.6	29.5	34.0	79.4	97.4
PriorNet	35.9	27.0	26.8	26.8	26.8	27.5	36.2		35.9	27.7	27.0	26.7	26.6	26.5	26.5
DDNet	55.6	34.4	31.7	30.4	29.5	30.2	33.4		55.6	40.9	34.1	28.0	26.9	26.6	26.5
EvNet	66.3	33.3	29.7	27.0	27.1	29.2	33.9		66.3	39.3	37.1	31.3	28.3	28.4	29.7
							Seg. – Seg	g.	class sky						
PostNet	98.4	74.8	51.0	37.2	32.8	43.5	49.9	Ŭ	98.4	84.7	66.1	42.4	34.8	40.9	71.2
PriorNet	32.1	30.9	30.8	30.8	30.8	30.8	30.8		32.1	31.0	30.8	30.8	30.8	30.8	30.8
DDNet	31.0	30.8	30.8	30.8	30.8	30.8	30.8		31.0	30.8	30.8	30.8	30.8	30.8	30.8
EvNet	98.3	83.0	60.5	34.0	31.0	30.8	30.8		98.3	94.4	88.8	65.6	37.0	31.4	30.9

Table 24. OOD detection (AUC-ROC) under PGD uncertainty attacks against precision α_0 on ID data and OOD data.

		ID-	Attack (non-atta	acked O	OD)				OOL)-Attack	(non-a	ttacked	ID)	
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0		0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIST -	K	MNIST						
PostNet	98.4	97.6	96.4	90.9	74.0	28.9	6.3		98.4	97.6	96.3	89.0	61.3	19.6	9.7
PriorNet	99.8	99.7	99.6	99.1	97.2	79.4	4.4		99.8	99.7	99.6	99.2	98.0	93.9	85.8
DDNet	96.5	95.9	95.1	92.0	82.6	44.3	3.5		96.5	95.9	95.2	92.9	88.6	78.7	59.4
EvNet	35.9	34.1	32.8	30.1	27.4	24.6	21.4		35.9	18.7	10.4	3.7	2.0	1.7	2.0
							CIFAR10) _	- SVHN						
PostNet	87.4	71.2	54.8	29.2	19.0	14.0	9.4		87.4	71.4	54.1	30.1	25.8	17.5	5.8
PriorNet	45.6	31.1	20.4	6.3	1.4	0.3	0.1		45.6	32.2	21.7	5.4	1.0	0.3	0.1
DDNet	84.9	73.8	61.8	30.2	9.3	3.0	0.8		84.9	76.6	66.2	34.6	10.4	2.3	0.3
EvNet	61.2	49.4	45.2	37.6	30.5	23.4	17.0		61.2	29.4	23.0	16.8	14.2	10.2	5.5
						Se	ns. – Sens	. (class 10, 1	1					
PostNet	87.2	48.8	37.3	4.1	0.7	0.7	0.7		87.2	50.0	45.4	16.5	27.6	81.9	98.0
PriorNet	37.3	3.5	2.4	2.2	2.9	6.3	19.2		37.3	8.0	3.6	1.4	0.6	0.1	0.0
DDNet	55.2	23.7	17.7	14.1	12.5	12.7	15.7		55.2	37.1	27.7	9.4	2.5	0.6	0.1
EvNet	75.5	30.8	18.2	5.8	1.6	0.6	0.2		75.5	47.8	41.9	24.1	10.2	10.2	15.6
							Seg. – Seg	ζ. (class sky						
PostNet	98.6	77.7	50.8	20.3	8.2	1.3	0.5		98.6	88.9	73.4	36.2	19.4	36.7	75.2
PriorNet	8.5	1.3	0.2	0.0	0.0	0.0	0.1	İ	8.5	2.0	0.4	0.0	0.0	0.0	0.0
DDNet	2.2	0.3	0.0	0.0	0.0	0.0	0.0		2.2	0.5	0.1	0.0	0.0	0.0	0.0
EvNet	97.7	78.4	47.7	9.9	1.2	0.2	0.1		97.7	93.5	86.9	62.2	21.5	3.7	1.0

Table 25. OOD detection (AU-PR) under PGD uncertainty attacks against distributional uncertainty on ID data and OOD data.

		ID-	Attack (non-atta	acked O	OD)				OOL)-Attack	(non-a	ttacked	ID)	
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0		0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIST -	K	KMNIST						
PostNet	80.5	76.2	73.4	69.1	66.6	65.4	60.2		80.5	72.1	63.9	43.9	33.0	30.9	30.8
PriorNet	99.6	99.4	99.2	98.0	94.1	76.3	31.2	١	99.6	99.4	99.2	98.2	95.2	87.2	75.2
DDNet	98.4	98.1	97.7	95.8	89.5	56.2	30.9		98.4	98.1	97.8	96.5	93.8	86.3	67.7
EvNet	40.1	39.5	39.1	38.2	37.3	36.5	35.6		40.1	34.6	32.6	31.3	31.0	31.0	31.1
							CIFAR10) -	- SVHN						
PostNet	64.2	44.7	37.5	31.1	28.5	25.0	19.3		64.2	31.0	19.5	16.3	16.4	16.5	16.3
PriorNet	40.8	27.4	20.4	15.9	15.4	15.4	15.4	İ	40.8	28.3	21.1	15.9	15.4	15.4	15.4
DDNet	82.0	71.0	59.1	29.9	16.6	15.5	15.4		82.0	72.2	60.3	26.3	16.2	15.4	15.4
EvNet	36.4	28.7	26.5	22.8	20.2	18.4	17.2		36.4	19.8	18.3	17.2	16.9	16.2	15.7
						Se	ns. – Sens	i. (class 10, 1	1					
PostNet	79.1	40.3	35.9	33.0	45.5	45.5	45.5		79.1	47.3	43.7	36.5	37.9	74.6	96.5
PriorNet	35.5	26.8	26.7	26.9	29.6	43.7	68.7		35.5	27.5	26.9	26.7	26.6	26.5	26.5
DDNet	52.9	31.7	29.8	29.1	28.4	30.1	37.6		52.9	38.4	31.5	27.5	26.8	26.6	26.5
EvNet	66.3	33.3	29.6	27.0	27.2	29.3	35.2		66.3	39.3	37.1	31.3	28.3	28.4	29.7
							Seg. – Seg	5.	class sky						
PostNet	98.0	76.3	53.1	37.4	32.9	44.6	50.2		98.0	83.5	64.8	41.8	35.4	43.1	71.3
PriorNet	32.3	30.9	30.8	30.8	30.8	32.5	45.0		32.3	31.0	30.8	30.8	30.8	30.8	30.8
DDNet	30.9	30.8	30.8	30.8	30.8	30.8	30.8		30.9	30.8	30.8	30.8	30.8	30.8	30.8
EvNet	98.1	82.1	59.1	33.8	31.0	30.8	30.8		98.1	93.8	88.2	64.5	36.4	31.3	31.0

Table 26. OOD detection (AUC-ROC) under PGD uncertainty attacks against distributional uncertainty on ID data and OOD data.

	ID-Attack (non-attacked OOD)									OOL)-Attack	(non-a	ttacked	ID)	
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0		0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIST -	K	MNIST						
PostNet	90.1	88.0	86.2	82.2	79.0	77.1	66.1		90.1	84.5	77.2	46.4	12.9	2.7	2.4
PriorNet	99.8	99.7	99.6	99.1	97.2	79.7	4.7		99.8	99.7	99.6	99.2	97.9	93.7	85.6
DDNet	98.1	97.7	97.2	94.8	87.0	48.7	3.0	İ	98.1	97.8	97.3	95.8	92.3	83.3	63.3
EvNet	36.8	35.0	33.7	30.9	28.2	25.3	22.1		36.8	19.3	10.7	3.9	2.1	1.8	2.2
							CIFAR10) -	- SVHN						
PostNet	82.9	67.7	59.2	51.3	47.7	40.1	24.2		82.9	51.9	26.2	8.9	9.5	11.1	9.9
PriorNet	48.0	33.6	22.5	7.1	1.6	0.3	0.1		48.0	34.8	24.0	6.7	1.6	0.6	0.2
DDNet	85.9	74.9	62.7	30.1	8.3	2.3	0.6	İ	85.9	77.6	66.9	32.1	8.0	1.5	0.2
EvNet	63.3	51.4	47.1	39.3	32.1	24.9	17.9		63.3	31.1	24.4	17.7	15.0	10.7	5.7
						Se	ns. – Sens	i. (class 10, 1	1					
PostNet	87.1	50.9	37.8	5.5	4.5	4.5	4.5		87.1	55.3	51.1	34.4	38.9	79. 7	97.9
PriorNet	36.5	2.9	1.8	1.8	5.2	21.5	52.8	İ	36.5	7.3	3.0	1.3	0.5	0.1	0.0
DDNet	52.3	18.7	13.1	10.3	9.3	10.8	18.4		52.3	33.1	22.0	6.7	2.2	0.6	0.1
EvNet	75.5	30.7	18.1	5.8	1.6	0.6	0.8		75.5	47.7	41.8	23.8	10.3	10.2	15.8
							Seg. – Seg	ξ. (class sky						
PostNet	98.6	78.3	51.9	20.5	8.3	2.1	1.7		98.6	88.8	73.1	35.9	21.4	39.9	75.9
PriorNet	9.4	1.6	0.3	0.0	0.0	1.8	15.4	İ	9.4	2.4	0.4	0.0	0.0	0.0	0.0
DDNet	1.3	0.2	0.0	0.0	0.0	0.0	0.0		1.3	0.2	0.0	0.0	0.0	0.0	0.0
EvNet	97.4	77.1	45.9	9.4	1.3	0.2	0.1		97.4	92.9	86.1	60.9	20.4	3.0	1.2

Table 27. OOD detection (AU-PR) under FGSM uncertainty attacks against differential entropy on ID data and OOD data.

		ID-	Attack	(non-att	acked C	OD)			00	D-Attac	ck (non-	attacked	l ID)	
Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0	4.0	0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIST -	KMNIST						
PostNet	94.5	94.2	94.1	93.5	89.9	81.2	71.6	94.5	93.3	92.0	87.6	81.1	75.7	75.7
PriorNet	99.6	99.4	99.2	98.1	95.6	90.0	65.3	99.6	99.4	99.2	98.6	97.5	95.9	94.4
DDNet	99.3	99.1	98.9	98.0	95.4	80.9	48.2	99.3	99.2	99.0	98.5	97.6	95.5	92.0
EvNet	69.0	67.4	66.2	64.0	61.9	59.8	56.70	9.0	60.1	56.5	53.4	52.7	52.9	53.5
							CIFAR10	- SVHN						
PostNet	81.8	66.2	61.6	64.2	65.7	61.3	48.4	81.8	63.1	51.9	43.4	46.6	61.7	77.0
PriorNet	54.4	40.6	33.8	27.0	25.5	27.2	35.5	54.4	42.3	36.8	30.6	28.3	29.5	32.1
DDNet	82.8	71.9	64.6	53.8	50.2	47.8	41.0	82.8	71.5	60.5	39.1	31.4	41.2	66.6
EvNet	80.3	67.8	64.0	61.9	61.6	57.4	49.6	80.3	59.2	51.5	46.7	49.0	56.3	64.6
							Sens. – Sens	. class 10,	11					
PostNet	74.5	40.6	37.2	31.4	38.1	44.9	45.9	74.5	99.6	99.8	99.9	99.9	99.9	99.9
PriorNet	32.3	35.7	57.6	83.1	88.8	79.7	70.0	32.3	28.3	28.1	27.6	28.0	32.7	38.5
DDNet	31.7	31.3	44.4	70.3	87.9	92.5	91.9	31.7	28.8	29.3	29.1	27.7	27.9	28.01
EvNet	66.5	45.7	46.8	42.3	42.0	41.4	41.8	66.5	54.7	66.5	76.2	71.1	75.3	75.8
							Seg. – Seg	. class sky						
PostNet	99.0	80.8	66.4	43.6	37.0	35.5	43.0	99.0	94.8	92.0	98.5	99.7	100.0	100.0
PriorNet	34.8	31.2	31.4	46.3	74.0	88.8	94.5	34.8	31.6	31.0	31.2	30.9	30.8	30.8
DDNet	31.5	30.8	30.8	30.9	37.9	56.2	84.3	31.5	30.9	30.8	30.8	30.8	30.8	30.8
EvNet	92.5	64.9	54.6	66.6	69.5	69.6	64.6	92.5	85.9	83.0	66.3	66.1	61.1	56.8

Table 28. OOD detection (AU-PR) under Noise uncertainty attacks against differential entropy on ID data and OOD data.

	ID-Attack (non-attacked OOD)									O	OD-Atta	ck (non-a	ttacked I	D)	
Noise Std	0.0	0.1	0.2	0.5	1.0	2.0	4.0		0.0	0.1	0.2	0.5	1.0	2.0	4.0
							MNIS	T –	KMNIS	T					
PostNet	93.0	94.2	82.3	34.4	31.6	31.0	30.9		92.2	91.8	91.5	92.3	92.7	93.2	93.5
PriorNet	99.7	99.6	96.7	40.0	40.6	45.7	55.6		99.5	97.3	96.5	99.4	100.0	99.5	72.4
DDNet	99.1	97.5	81.2	31.3	31.0	30.9	31.2	İ	99.0	98.8	99.2	99.8	99.9	99.8	99.1
EvNet	65.5	60.5	51.4	35.3	34.5	35.5	35.0		62.5	47.2	40.9	35.1	34.6	33.5	34.9
							CIFA	R10	- SVH	N					
PostNet	88.5	41.4	39.8	31.0	30.7	31.6	33.9		88.5	86.6	81.9	93.0	98.5	98.6	97.3
PriorNet	73.3	88.3	95.3	92.4	70.4	30.9	30.8		73.3	31.6	30.9	31.7	51.8	94.3	100.0
DDNet	87.3	69.3	78.4	55.2	31.6	30.7	31.4	İ	87.3	55.8	57.9	73.9	97.3	99.5	97.2
EvNet	92.4	56.8	53.8	33.4	30.9	32.9	36.6		92.4	73.7	73.5	77.7	93.7	92.5	92.1
							Sens. – S	ens.	class 10), 11					
PostNet	85.3	30.8	39.4	50.0	50.0	50.0	50.0		85.3	98.9	100.0	100.0	100.0	100.0	100.0
PriorNet	32.3	30.8	34.9	83.7	77.7	49.8	80.3	İ	32.3	30.7	30.7	32.5	40.1	49.9	47.6
DDNet	31.1	30.7	30.7	32.4	58.8	88.1	74.3		31.1	30.7	30.7	30.7	30.8	31.6	39.1
EvNet	80.3	30.8	31.2	37.9	46.3	50.0	50.0		80.3	34.6	38.4	53.9	69.3	78.8	81.5
							Seg. –	Seg	. class sl	сy					
PostNet	99.9	41.8	30.8	34.5	49.1	50.0	50.0		99.9	97.4	96.6	99.5	100.0	100.0	100.0
PriorNet	31.0	30.8	30.8	30.8	32.7	69.0	78.3		31.0	30.8	30.8	30.8	30.9	31.1	32.4
DDNet	30.8	30.8	30.8	30.8	30.8	58.2	91.3		30.8	30.8	30.8	30.8	30.8	30.8	31.9
EvNet	99.1	38.1	32.2	30.8	30.8	32.2	37.5		99.1	95.6	87.6	58.0	44.9	46.6	53.8

6.4. How to make DBU models more robust

To improve robustness of DBU models we perform median smoothing and adversarial training. Smoothing computes the smooth median, worst case and best case performance of DBU models for three tasks: distinguishing between correct and wrong predictions, attack detection, distinguishing between ID data and OOD data under label attacks and under uncertainty attacks.

Table 29. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under PGD label attacks. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$80.5 \cdot 91.5 \cdot 94.5 \\ 81.9 \cdot 86.8 \cdot 88.0 \\ 65.9 \cdot 81.2 \cdot 83.0 \\ 76.3 \cdot 90.2 \cdot 91.7$	52.8 · 71.6 · 95.2 69.6 · 78.0 · 90.1 55.8 · 70.5 · 87.2 54.7 · 74.3 · 95.7	$31.9 \cdot 51.0 \cdot 96.8 \\ 50.9 \cdot 65.8 \cdot 89.4 \\ 37.8 \cdot 56.8 \cdot 88.1 \\ 31.6 \cdot 51.5 \cdot 94.5$	$\begin{array}{c} 5.6 \cdot 11.7 \cdot 100.0 \\ 36.5 \cdot 59.9 \cdot 97.0 \\ 10.1 \cdot 21.9 \cdot 94.3 \\ 5.8 \cdot 11.9 \cdot 86.9 \end{array}$	$0.3 \cdot 0.6 \cdot 100.0$ $24.3 \cdot 39.3 \cdot 100.0$ $0.9 \cdot 1.6 \cdot 99.6$ $1.9 \cdot 7.0 \cdot 100.0$	$\begin{array}{c} 0.0 \cdot \textbf{0.0} \cdot 100.0 \\ 9.2 \cdot \textbf{17.9} \cdot 100.0 \\ 0.0 \cdot \textbf{0.0} \cdot 100.0 \\ 1.1 \cdot \textbf{4.0} \cdot 100.0 \end{array}$
Smoothed + adv. w. label attacks	PostNet PriorNet DDNet EvNet	- - -	52.1 · 71.8 · 95.6 57.6 · 71.7 · 88.9 58.6 · 78.4 · 92.2 24.3 · 34.2 · 51.8	31.2 · 47.9 · 96.1 46.1 · 64.5 · 90.1 49.4 · 66.0 · 90.5 32.6 · 49.5 · 95.5	$7.8 \cdot 14.7 \cdot 98.6$ $38.1 \cdot 59.3 \cdot 99.5$ $12.0 \cdot 21.4 \cdot 98.1$ $5.9 \cdot 13.0 \cdot 100.0$	1.8 · 4.4 · 100.0 32.3 · 51.7 · 100.0 0.8 · 1.0 · 96.6 2.6 · 5.2 · 99.9	$\begin{array}{c} 0.3 \cdot 0.5 \cdot 100.0 \\ 22.1 \cdot 41.6 \cdot 97.4 \\ 0.0 \cdot 0.0 \cdot 100.0 \\ 2.9 \cdot 5.9 \cdot 100.0 \end{array}$
Smoothed + adv. w. uncert. attacks	PostNet PriorNet DDNet EvNet	- - - -	52.8 · 74.2 · 94.6 50.6 · 68.1 · 88.6 68.8 · 84.4 · 93.2 54.2 · 73.7 · 96.1	33.0 · 49.4 · 87.5 44.4 · 66.1 · 96.0 45.1 · 60.8 · 86.8 30.5 · 50.0 · 99.5	$\begin{array}{c} 7.7 \cdot 14.2 \cdot 99.0 \\ 35.1 \cdot 57.4 \cdot 98.4 \\ 12.3 \cdot 22.0 \cdot 91.0 \\ 7.1 \cdot 13.9 \cdot 100.0 \end{array}$	$0.6 \cdot 1.2 \cdot 100.0$ $18.4 \cdot 32.2 \cdot 100.0$ $0.8 \cdot 1.7 \cdot 87.0$ $3.7 \cdot 8.7 \cdot 75.2$	$\begin{array}{c} 0.7 \cdot 1.1 \cdot 100.0 \\ 15.2 \cdot 29.3 \cdot 100.0 \\ 0.0 \cdot 0.0 \cdot 100.0 \\ 3.3 \cdot 5.8 \cdot 100.0 \end{array}$

Table 30. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under PGD label attacks. Smoothed DBU models on MNIST. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
	PostNet	$97.2 \cdot 99.4 \cdot 100.0$	95.9 · 99.1 · 99.9	94.7 · 98.9 · 99.9	89.3 · 96.8 · 99.9	75.5 · 90.2 · 100.0	35.5 · 56.7 · 100.0
Smoothed	PriorNet	$96.8 \cdot 99.2 \cdot 99.3$	$95.5 \cdot 99.1 \cdot 99.7$	$94.6 \cdot 98.8 \cdot 99.7$	$90.2 \cdot 97.2 \cdot 99.9$	$81.1 \cdot 93.4 \cdot 99.9$	$53.9 \cdot 75.2 \cdot 100.0$
models	DDNet	$97.6 \cdot 99.4 \cdot 99.5$	$96.8 \cdot 99.2 \cdot 99.4$	$95.5 \cdot 98.8 \cdot 99.4$	$90.4 \cdot 97.2 \cdot 99.8$	$77.0 \cdot 91.3 \cdot 100.0$	$29.2 \cdot 48.6 \cdot 100.0$
	EvNet	$97.3 \cdot 99.4 \cdot 99.4$	$95.4 \cdot 98.8 \cdot 99.6$	$93.9 \cdot 98.7 \cdot 99.9$	$89.0 \cdot 96.5 \cdot 100.0$	$78.9 \cdot 92.9 \cdot 100.0$	$52.2 \cdot 73.2 \cdot 100.0$
Smoothed	PostNet	-	$94.4 \cdot 98.6 \cdot 99.5$	90.6 · 97.9 · 99.9	83.4 · 93.1 · 99.9	72.1 · 91.2 · 100.0	41.8 · 65.0 · 100.0
+ adv. w.	PriorNet	-	$94.4 \cdot 98.5 \cdot 99.5$	$93.6 \cdot 98.8 \cdot 99.8$	$89.1 \cdot 96.6 \cdot 99.8$	$81.5 \cdot 94.5 \cdot 100.0$	$71.6 \cdot 88.4 \cdot 100.0$
label	DDNet	-	$94.9 \cdot 98.3 \cdot 98.7$	$94.6 \cdot 97.9 \cdot 98.9$	$88.2 \cdot 97.4 \cdot 99.8$	$72.1 \cdot 89.3 \cdot 100.0$	$28.1\cdot49.3\cdot100.0$
attacks	EvNet	-	$88.8 \cdot 95.3 \cdot 97.9$	$91.5 \cdot 97.1 \cdot 99.4$	$85.2\cdot94.9\cdot100.0$	$78.1 \cdot \textbf{91.4} \cdot 100.0$	$54.3 \cdot 75.3 \cdot 100.0$
Smoothed	PostNet	-	92.8 · 98.3 · 99.8	92.5 · 98.3 · 99.9	86.2 · 94.8 · 99.8	71.0 · 89.5 · 100.0	34.6 · 54.2 · 100.0
+ adv. w.	PriorNet	-	$95.1 \cdot 98.6 \cdot 99.6$	$94.1 \cdot 98.0 \cdot 99.4$	$87.7 \cdot 97.2 \cdot 99.9$	$80.2 \cdot 93.4 \cdot 100.0$	$68.5 \cdot 87.8 \cdot 100.0$
uncert.	DDNet	-	$96.0 \cdot 98.4 \cdot 98.8$	$95.0\cdot97.6\cdot98.7$	$87.6 \cdot 95.3 \cdot 99.7$	$73.9 \cdot 90.2 \cdot 100.0$	$32.8 \cdot 54.4 \cdot 100.0$
attacks	EvNet	-	$93.3 \cdot 98.6 \cdot 99.5$	$89.8 \cdot 97.2 \cdot 99.2$	$86.2 \cdot 95.4 \cdot 100.0$	$82.1 \cdot 93.7 \cdot 100.0$	$52.4 \cdot 73.3 \cdot 100.0$

Table 31. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under PGD label attacks. Smoothed DBU models on Sensorless. Column format: guaranteed lowest performance \cdot empirical performance \cdot guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	93.5 · 98.4 · 100.0 97.1 · 99.3 · 100.0 95.9 · 98.9 · 99.7 94.0 · 99.0 · 99.9	6.7 · 12.4 · 100.0 8.6 · 17.6 · 100.0 7.0 · 14.0 · 100.0 18.1 · 34.2 · 100.0	$\begin{array}{c} 2.9 \cdot 5.3 \cdot 100.0 \\ 3.3 \cdot 7.7 \cdot 100.0 \\ 0.8 \cdot 1.3 \cdot 100.0 \\ 9.6 \cdot 17.1 \cdot 100.0 \end{array}$	$\begin{array}{cccc} 4.1 \cdot & \textbf{4.1} \cdot & 49.1 \\ 0.7 \cdot & \textbf{1.5} \cdot 100.0 \\ 0.2 \cdot & \textbf{0.4} \cdot 100.0 \\ 4.1 \cdot & \textbf{6.8} \cdot 100.0 \end{array}$	$\begin{array}{cccc} 6.4 \cdot & 6.4 \cdot & 6.4 \\ 0.4 \cdot & 0.7 \cdot 100.0 \\ 0.2 \cdot & 0.2 \cdot 100.0 \\ 2.7 \cdot & 4.9 \cdot 100.0 \end{array}$	$\begin{array}{c} 10.6 \cdot 10.6 \cdot \ 10.6 \\ 0.1 \cdot \ 0.2 \cdot 100.0 \\ 0.2 \cdot \ 0.4 \cdot 100.0 \\ 2.4 \cdot \ 4.3 \cdot 100.0 \end{array}$
Smoothed + adv. w. label attacks	PostNet PriorNet DDNet EvNet	- - - -	$7.9 \cdot 14.9 \cdot 100.0$ $18.1 \cdot 32.1 \cdot 100.0$ $6.9 \cdot 13.4 \cdot 100.0$ $19.7 \cdot 35.7 \cdot 100.0$	$2.9 \cdot 6.3 \cdot 100.0 \\ 8.7 \cdot 16.7 \cdot 100.0 \\ 4.3 \cdot 9.0 \cdot 100.0 \\ 9.4 \cdot 16.2 \cdot 100.0$	6.6 · 6.6 · 6.6 0.1 · 0.2 · 100.0 0.2 · 0.3 · 100.0 1.6 · 3.0 · 100.0	$\begin{array}{cccc} 7.2 \cdot & \textbf{7.2} \cdot & 7.2 \\ 0.0 \cdot & \textbf{0.0} \cdot 100.0 \\ 0.2 \cdot & \textbf{0.4} \cdot 100.0 \\ 2.5 \cdot & \textbf{5.6} \cdot 100.0 \end{array}$	$\begin{array}{cccc} 9.6 \cdot & 9.6 \cdot & 9.6 \\ 0.8 \cdot & 1.0 \cdot 100.0 \\ 0.2 \cdot & 0.8 \cdot 100.0 \\ 1.0 \cdot & 1.8 \cdot 100.0 \end{array}$
Smoothed + adv. w. uncert. attacks	PostNet PriorNet DDNet EvNet	- - - -	7.9 · 14.4 · 100.0 19.1 · 32.7 · 100.0 5.4 · 10.2 · 100.0 22.3 · 38.4 · 100.0	$4.8 \cdot 9.3 \cdot 100.0 \\ 6.9 \cdot 13.7 \cdot 100.0 \\ 0.7 \cdot 1.8 \cdot 100.0 \\ 11.7 \cdot 22.4 \cdot 100.0$	$\begin{array}{c} 6.6 \cdot \ 6.6 \cdot \ 6.6 \\ 0.7 \cdot \ 1.7 \cdot 100.0 \\ 0.5 \cdot \ 0.9 \cdot 100.0 \\ 7.1 \cdot 13.1 \cdot 100.0 \end{array}$	6.7 · 6.7 · 6.7 0.0 · 0.0 · 100.0 0.3 · 1.2 · 100.0 1.8 · 3.4 · 100.0	$\begin{array}{c} 10.6 \cdot 10.6 \cdot \ 10.6 \\ 0.0 \cdot \ 0.0 \cdot 100.0 \\ 0.2 \cdot \ 0.6 \cdot 100.0 \\ 0.6 \cdot \ 1.0 \cdot 100.0 \end{array}$

Table 32. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under PGD label attacks. Smoothed DBU models on Segment. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model)..

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	94.0 · 99.1 · 99.8 97.0 · 99.8 · 99.9 96.2 · 99.5 · 99.7 95.8 · 99.6 · 99.9	63.5 · 84.7 · 100.0 75.6 · 90.8 · 100.0 75.7 · 89.8 · 99.9 80.2 · 93.7 · 100.0	$33.2 \cdot 56.1 \cdot 100.0$ $31.1 \cdot 50.8 \cdot 100.0$ $28.5 \cdot 51.6 \cdot 100.0$ $35.2 \cdot 57.2 \cdot 100.0$	$10.2 \cdot 16.9 \cdot 100.0$ $2.6 \cdot 4.7 \cdot 100.0$ $3.7 \cdot 8.2 \cdot 100.0$ $6.8 \cdot 12.0 \cdot 100.0$	$5.2 \cdot 10.3 \cdot 100.0$ $0.0 \cdot 0.0 \cdot 100.0$ $0.0 \cdot 0.0 \cdot 100.0$ $1.2 \cdot 2.1 \cdot 100.0$	$\begin{array}{cccc} 0.3 \cdot & 0.3 \cdot & 0.3 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 1.1 \cdot & 2.0 \cdot 100.0 \end{array}$
Smoothed + adv. w. label attacks	PostNet PriorNet DDNet EvNet	- - -	$66.0 \cdot 85.5 \cdot 100.0$ $79.0 \cdot 92.4 \cdot 100.0$ $76.2 \cdot 91.0 \cdot 99.6$ $82.7 \cdot 95.2 \cdot 100.0$	$22.5 \cdot 41.0 \cdot 100.0 \\ 45.2 \cdot 68.8 \cdot 100.0 \\ 27.2 \cdot 45.3 \cdot 100.0 \\ 34.0 \cdot 53.8 \cdot 100.0$	$9.0 \cdot 16.3 \cdot 100.0$ $9.2 \cdot 13.9 \cdot 100.0$ $2.3 \cdot 4.3 \cdot 100.0$ $10.9 \cdot 23.2 \cdot 100.0$	$\begin{array}{ccc} 5.2 \cdot & 9.7 \cdot 100.0 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 0.5 \cdot & 4.2 \cdot 100.0 \end{array}$	$\begin{array}{ccc} 0.6 \cdot & 0.6 \cdot & 0.6 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 2.1 \cdot & 5.1 \cdot 100.0 \end{array}$
Smoothed + adv. w. uncert. attacks	PostNet PriorNet DDNet EvNet	- - -	$71.5 \cdot 87.6 \cdot 100.0$ $82.1 \cdot 96.5 \cdot 100.0$ $77.4 \cdot 91.4 \cdot 99.9$ $76.2 \cdot 90.7 \cdot 100.0$	33.5 · 54.5 · 100.0 44.1 · 65.4 · 100.0 29.4 · 50.3 · 100.0 35.7 · 55.4 · 100.0	$12.8 \cdot 25.6 \cdot 100.0 \\ 9.0 \cdot 15.7 \cdot 100.0 \\ 4.0 \cdot 6.5 \cdot 100.0 \\ 4.2 \cdot 6.4 \cdot 100.0$	$\begin{array}{ccc} 6.5 \cdot 10.3 & \cdot & 87.2 \\ 0.0 \cdot & 0.0 \cdot & 100.0 \\ 0.0 \cdot & 0.0 \cdot & 100.0 \\ 0.8 \cdot & 1.4 \cdot & 100.0 \end{array}$	$\begin{array}{ccc} 0.0 \cdot & 0.0 \cdot 100.0 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ 0.0 \cdot & 0.0 \cdot 100.0 \\ \end{array}$

Table 33. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under FGSM label attacks. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$80.5 \cdot 91.4 \cdot 94.4$ $81.9 \cdot 87.7 \cdot 88.8$ $65.9 \cdot 84.1 \cdot 85.6$ $76.3 \cdot 90.4 \cdot 91.7$	52.3 · 73.2 · 95.4 69.6 · 78.4 · 90.3 55.3 · 69.6 · 87.0 54.1 · 74.5 · 95.5	35.8 · 57.2 · 97.5 53.3 · 70.5 · 91.7 38.6 · 55.8 · 87.2 35.5 · 54.7 · 95.1	17.0 · 29.0 · 100.0 42.1 · 62.6 · 97.2 16.3 · 28.5 · 94.6 14.6 · 29.3 · 95.6	$10.2 \cdot 18.7 \cdot 100.0$ $37.5 \cdot 55.7 \cdot 100.0$ $6.4 \cdot 12.0 \cdot 99.9$ $8.6 \cdot 16.1 \cdot 100.0$	$8.1 \cdot 14.7 \cdot 100.0$ $36.0 \cdot 59.5 \cdot 100.0$ $3.6 \cdot 7.2 \cdot 100.0$ $7.2 \cdot 13.0 \cdot 100.0$
Smoothed	PostNet	-	52.3 · 71.6 · 95.1	34.7 · 54.8 · 96.6	18.9 · 32.1 · 99.4	10.9 · 19.2 · 100.0	8.5 · 16.2 · 100.0
+ adv. w.	PriorNet	-	58.1 · 69.6 · 87.6	47.1 · 65.7 · 90.3	40.2 · 59.5 · 99.3	36.2 · 59.5 · 100.0	25.1 · 42.1 · 97.7
label	DDNet	-	57.1 · 75.2 · 91.0	49.3 · 65.3 · 90.5	18.4 · 33.6 · 98.5	7.6 · 13.5 · 99.9	3.3 · 9.6 · 100.0
attacks	EvNet	-	24.1 · 36.5 · 54.2	37.1 · 56.7 · 96.7	16.2 · 29.9 · 100.0	11.4 · 21.8 · 100.0	13.0 · 26.1 · 100.0
Smoothed	PostNet	-	52.0 · 71.8 · 94.5	35.8 · 54.6 · 89.9	18.4 · 33.6 · 99.8	$10.2 \cdot 19.1 \cdot 100.0$ $27.7 \cdot 46.2 \cdot 100.0$ $6.1 \cdot 13.1 \cdot 91.8$ $6.1 \cdot 13.5 \cdot 86.1$	12.2 · 23.0 · 100.0
+ adv. w.	PriorNet	-	50.6 · 67.3 · 88.5	46.2 · 64.3 · 95.1	39.9 · 60.8 · 98.5		28.5 · 48.6 · 100.0
uncert.	DDNet	-	67.7 · 82.2 · 92.4	45.7 · 64.7 · 88.8	20.5 · 34.8 · 93.6		4.1 · 8.4 · 100.0
attacks	EvNet	-	53.9 · 73.6 · 96.3	34.2 · 55.3 · 99.7	16.1 · 31.2 · 100.0		18.1 · 34.0 · 100.0

Table 34. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under FGSM label attacks. Smoothed DBU models on MNIST. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

		<u> </u>					
	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
	PostNet	$97.2 \cdot 99.3 \cdot 99.9$	$96.1\cdot99.2\cdot99.9$	$95.2 \cdot 98.9 \cdot 99.9$	$91.7 \cdot 98.0 \cdot 99.9$	$86.1\cdot95.9\cdot100.0$	$75.7 \cdot 91.1 \cdot 100.0$
Smoothed	PriorNet	$96.8 \cdot 99.2 \cdot 99.3$	$95.5 \cdot 99.0 \cdot 99.6$	$94.7 \cdot 98.7 \cdot 99.6$	$91.3 \cdot 97.6 \cdot 99.9$	$85.5 \cdot 95.6 \cdot 100.0$	$78.7 \cdot 92.4 \cdot 100.0$
models	DDNet	$97.6 \cdot 99.3 \cdot 99.4$	$96.8 \cdot 99.2 \cdot 99.5$	$95.6 \cdot 98.7 \cdot 99.4$	$91.7 \cdot 97.7 \cdot 99.9$	$83.4 \cdot 95.2 \cdot 100.0$	$58.3 \cdot 79.6 \cdot 100.0$
	EvNet	$97.3 \cdot 99.3 \cdot 99.4$	$95.5 \cdot 99.0 \cdot 99.6$	$94.3 \cdot 98.9 \cdot 99.9$	$92.0\cdot97.7\cdot100.0$	$87.4 \cdot 96.3 \cdot 100.0$	$78.8 \cdot 92.4 \cdot 100.0$
Smoothed	PostNet	-	95.1 · 98.9 · 99.8	91.2 · 97.2 · 99.6	87.6 · 96.3 · 99.9	81.0 · 93.3 · 100.0	69.9 · 87.2 · 100.0
+ adv. w.	PriorNet	-	$94.4 \cdot 98.7 \cdot 99.7$	$93.6 \cdot 98.2 \cdot 99.3$	$89.4 \cdot 96.3 \cdot 99.8$	$84.5 \cdot 95.1 \cdot 100.0$	$81.7 \cdot 92.5 \cdot 100.0$
label	DDNet	-	$95.5 \cdot 98.6 \cdot 99.0$	$94.6 \cdot 98.7 \cdot 99.4$	$89.7 \cdot 97.1 \cdot 99.8$	$80.0 \cdot 93.6 \cdot 100.0$	$54.4 \cdot 74.5 \cdot 100.0$
attacks	EvNet	-	$88.9 \cdot 94.8 \cdot 98.1$	$91.5 \cdot 98.4 \cdot 99.8$	$89.2 \cdot 97.0 \cdot 100.0$	$83.6 \cdot 94.7 \cdot 100.0$	$72.3 \cdot 88.0 \cdot 100.0$
Smoothed	PostNet	-	92.8 · 98.5 · 99.9	92.8 · 98.7 · 99.9	89.0 · 96.3 · 99.8	80.8 · 93.4 · 100.0	71.6 · 86.9 · 100.0
+ adv. w.	PriorNet	-	$95.1 \cdot 98.1 \cdot 98.9$	$94.3 \cdot 97.7 \cdot 99.1$	$88.5 \cdot 96.8 \cdot 99.9$	$83.4 \cdot 94.5 \cdot 100.0$	$78.9 \cdot 92.2 \cdot 100.0$
uncert.	DDNet	-	$96.0 \cdot 98.7 \cdot 99.0$	$95.5 \cdot 98.6 \cdot 99.3$	$89.5 \cdot 95.6 \cdot 99.7$	$79.6 \cdot 93.1 \cdot 100.0$	$55.9 \cdot 77.1 \cdot 100.0$
attacks	EvNet	-	$93.3 \cdot 98.9 \cdot 99.4$	$90.1\cdot97.9\cdot99.4$	$87.9 \cdot 96.3 \cdot 100.0$	$84.1\cdot94.2\cdot100.0$	$69.2 \cdot \textbf{86.9} \cdot 100.0$

Table 35. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under FGSM label attacks. Smoothed DBU models on Sensorless. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$\begin{array}{c} 94.5 \cdot 98.1 \cdot 100.0 \\ 97.1 \cdot 99.5 \cdot 100.0 \\ 95.9 \cdot 99.4 \cdot 99.8 \\ 94.0 \cdot 98.5 \cdot 99.7 \end{array}$	$10.3 \cdot 19.6 \cdot 100.0$ $13.6 \cdot 27.3 \cdot 100.0$ $8.6 \cdot 14.9 \cdot 100.0$ $26.0 \cdot 43.2 \cdot 100.0$	$\begin{array}{c} 5.1 \cdot 11.0 \cdot 100.0 \\ 6.3 \cdot 12.4 \cdot 100.0 \\ 1.6 \cdot 3.8 \cdot 100.0 \\ 15.8 \cdot 30.8 \cdot 100.0 \end{array}$	$6.4 \cdot 6.4 \cdot 6.4 2.6 \cdot 6.8 \cdot 100.0 2.6 \cdot 4.5 \cdot 100.0 11.7 \cdot 20.2 \cdot 100.0$	$\begin{array}{c} 10.4 \cdot 10.4 \cdot 10.4 \\ 3.1 \cdot 7.3 \cdot 100.0 \\ 3.4 \cdot 6.9 \cdot 100.0 \\ 8.1 \cdot 15.0 \cdot 100.0 \end{array}$	$11.4 \cdot 11.4 \cdot 11.4 \\ 3.2 \cdot 6.7 \cdot 100.0 \\ 3.3 \cdot 6.4 \cdot 100.0 \\ 7.6 \cdot 12.7 \cdot 100.0$
Smoothed + adv. w. label attacks	PostNet PriorNet DDNet EvNet	- - -	$13.1 \cdot 24.3 \cdot 100.0$ $22.4 \cdot 38.2 \cdot 100.0$ $7.3 \cdot 13.2 \cdot 100.0$ $25.5 \cdot 42.0 \cdot 100.0$	$5.7 \cdot 11.9 \cdot 100.0$ $11.8 \cdot 22.1 \cdot 100.0$ $8.5 \cdot 17.2 \cdot 100.0$ $15.6 \cdot 30.2 \cdot 100.0$	$\begin{array}{c} 9.4 \cdot 9.4 \cdot 9.4 \\ 0.2 \cdot 0.6 \cdot 100.0 \\ 3.6 \cdot 7.9 \cdot 100.0 \\ 10.4 \cdot 19.5 \cdot 100.0 \end{array}$	$11.2 \cdot 11.2 \cdot 11.2 \\ 0.0 \cdot 0.0 \cdot 100.0 \\ 3.8 \cdot 7.6 \cdot 100.0 \\ 8.6 \cdot 16.4 \cdot 100.0$	$11.8 \cdot 11.8 \cdot 11.8 \\ 0.1 \cdot 0.1 \cdot 100.0 \\ 0.8 \cdot 1.2 \cdot 100.0 \\ 7.8 \cdot 14.7 \cdot 100.0$
Smoothed + adv. w. uncert. attacks	PostNet PriorNet DDNet EvNet	- - -	10.6 · 20.3 · 100.0 25.7 · 45.0 · 100.0 7.9 · 16.4 · 100.0 27.9 · 49.2 · 100.0	5.2 · 9.9 · 100.0 12.0 · 20.5 · 100.0 1.2 · 3.8 · 100.0 18.4 · 32.9 · 100.0	$10.9 \cdot 10.9 \cdot 10.9$ $1.1 \cdot 3.7 \cdot 100.0$ $3.4 \cdot 6.3 \cdot 100.0$ $16.4 \cdot 29.3 \cdot 100.0$	$11.6 \cdot 11.6 \cdot 11.6 \\ 0.0 \cdot 0.0 \cdot 100.0 \\ 3.9 \cdot 7.9 \cdot 100.0 \\ 5.9 \cdot 10.8 \cdot 100.0$	$11.7 \cdot 11.7 \cdot 11.7 0.0 \cdot 0.0 \cdot 100.0 3.3 \cdot 8.0 \cdot 100.0 8.5 \cdot 16.1 \cdot 100.0$

Table 36. Distinguishing between correctly and wrongly labeled inputs based on differential entropy under FGSM label attacks. Smoothed DBU models on Segment. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
	PostNet	94.0 · 99.2 · 99.8	55.2 · 78.3 · 100.0	40.1 · 61.4 · 100.0	17.9 · 31.7 · 100.0	$6.8 \cdot 12.7 \cdot 100.0$	17.6 · 17.9 · 18.0
Smoothed	PriorNet	$97.0 \cdot 99.8 \cdot 99.9$	$69.2 \cdot 89.7 \cdot 100.0$	$29.7 \cdot 45.5 \cdot 100.0$	$1.7\cdot4.1\cdot100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$
models	DDNet	$96.2 \cdot 99.5 \cdot 99.6$	$70.6 \cdot 86.3 \cdot 99.8$	$22.3 \cdot 38.8 \cdot 100.0$	$6.3 \cdot 13.3 \cdot 100.0$	$1.1 \cdot \ 3.0 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$
	EvNet	$95.8 \cdot 99.1 \cdot 99.8$	$78.4 \cdot 92.5 \cdot 100.0$	$40.7 \cdot 62.1 \cdot 100.0$	$9.8 \cdot 17.6 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$	$0.0 \cdot 0.0 \cdot 100.0$
Smoothed	PostNet	-	66.0 · 83.5 · 100.0	28.8 · 44.9 · 100.0	12.3 · 24.3 · 100.0	9.3 · 17.3 · 100.0	24.8 · 24.8 · 24.8
+ adv. w.	PriorNet	-	$75.1 \cdot 91.5 \cdot 99.9$	$34.0 \cdot 60.3 \cdot 100.0$	$11.1 \cdot 24.6 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$
label	DDNet	-	$65.4 \cdot 82.8 \cdot 99.5$	$23.1 \cdot 35.3 \cdot 100.0$	$4.8 \cdot 10.4 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$
attacks	EvNet	-	$83.4 \cdot 95.3 \cdot 100.0$	$42.1\cdot \textbf{63.3}\cdot 100.0$	$15.0 \cdot 33.6 \cdot 100.0$	$0.0 \cdot 0.0 \cdot 100.0$	$0.0 \cdot 0.0 \cdot 100.0$
Smoothed	PostNet	-	67.8 · 86.5 · 100.0	34.0 · 52.5 · 100.0	16.2 · 32.8 · 100.0	14.4 · 25.2 · 92.2	7.3 · 7.3 · 7.3
+ adv. w.	PriorNet	-	$77.3 \cdot 91.2 \cdot 99.9$	$39.3 \cdot 62.7 \cdot 100.0$	$9.0 \cdot 17.8 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$
uncert.	DDNet	-	$68.8 \cdot 88.3 \cdot 99.9$	$20.4\cdot\boldsymbol{35.2}\cdot100.0$	$7.5 \cdot 12.6 \cdot 100.0$	$0.3 \cdot \ 0.9 \cdot 100.0$	$0.0 \cdot \ 0.0 \cdot 100.0$
attacks	EvNet	-	$74.0\cdot92.9\cdot100.0$	$44.1\cdot \textbf{61.8}\cdot 100.0$	$5.3 \cdot 13.0 \cdot 100.0$	$3.9 \cdot 8.2 \cdot 100.0$	$0.5 \cdot 4.2 \cdot 100.0$

Table 37. Attack detection (PGD label attacks) based on differential entropy. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	33.1 · 50.4 · 89.9 35.9 · 50.6 · 74.5 36.3 · 50.3 · 76.4 32.9 · 50.4 · 89.8	31.0 · 50.2 · 96.9 33.0 · 50.3 · 82.8 32.8 · 49.9 · 84.6 31.4 · 50.1 · 94.0	30.7 · 50.2 · 100.0 31.2 · 50.0 · 95.7 30.8 · 50.1 · 98.0 30.8 · 50.0 · 98.0	30.7 · 50.0 · 100.0 30.7 · 50.4 · 99.9 30.7 · 50.2 · 100.0 30.7 · 50.3 · 100.0	$30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 50.4 \cdot 100.0$ $30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 49.6 \cdot 100.0$
Smoothed	PostNet	32.7 · 50.1 · 90.4	31.1 · 50.2 · 96.5	$30.7 \cdot 50.2 \cdot 99.7$	30.7 · 50.3 · 100.0	$30.7 \cdot 50.2 \cdot 100.0$
+ adv. w.	PriorNet	35.2 · 51.8 · 78.6	32.8 · 51.1 · 84.4	$30.8 \cdot 50.2 \cdot 98.7$	30.7 · 50.5 · 100.0	$30.8 \cdot 50.1 \cdot 98.2$
label	DDNet	35.5 · 50.6 · 79.2	33.4 · 50.3 · 84.1	$30.8 \cdot 50.1 \cdot 99.2$	30.7 · 50.0 · 100.0	$30.7 \cdot 50.5 \cdot 100.0$
attacks	EvNet	40.3 · 50.4 · 66.8	31.4 · 50.3 · 95.8	$30.7 \cdot 50.3 \cdot 100.0$	30.7 · 50.1 · 100.0	$30.7 \cdot 50.0 \cdot 100.0$
Smoothed	PostNet	$33.3 \cdot 50.6 \cdot 88.7$	$32.5 \cdot 50.1 \cdot 87.9 \\ 31.4 \cdot 50.6 \cdot 92.8 \\ 33.4 \cdot 50.2 \cdot 83.0 \\ 30.8 \cdot 50.0 \cdot 99.6$	30.7 · 49.9 · 99.8	30.7 · 50.1 · 100.0	$30.7 \cdot 50.0 \cdot 100.0$
+ adv. w.	PriorNet	$34.5 \cdot 51.0 \cdot 80.1$		30.9 · 50.0 · 97.7	30.7 · 50.1 · 100.0	$30.7 \cdot 50.0 \cdot 100.0$
uncert.	DDNet	$37.4 \cdot 50.8 \cdot 74.5$		30.9 · 50.1 · 96.8	30.8 · 49.9 · 98.1	$30.7 \cdot 49.9 \cdot 100.0$
attacks	EvNet	$32.8 \cdot 50.1 \cdot 92.0$		30.7 · 50.1 · 100.0	31.2 · 50.2 · 96.1	$31.0 \cdot 50.0 \cdot 100.0$

Table 38. Attack detection (PGD label attacks) based on differential entropy. Smoothed DBU models on MNIST. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet	$30.9 \cdot 52.5 \cdot 95.6$	$31.5 \cdot 51.5 \cdot 90.9$	$31.1 \cdot 49.9 \cdot 97.1$	$30.7 \cdot 47.6 \cdot 100.0$	$30.7 \cdot 45.0 \cdot 100.0$
	PriorNet	$38.2 \cdot 57.8 \cdot 80.9$	$36.0 \cdot 57.2 \cdot 84.3$	$31.6 \cdot 63.4 \cdot 98.4$	$30.8 \cdot 61.0 \cdot 99.3$	$30.7 \cdot 66.8 \cdot 100.0$
	DDNet	$44.6 \cdot 51.9 \cdot 60.7$	$39.3 \cdot 52.7 \cdot 72.2$	$31.6 \cdot 50.9 \cdot 95.2$	$30.7 \cdot 47.3 \cdot 100.0$	$30.7 \cdot 45.9 \cdot 100.0$
	EvNet	$36.5 \cdot 51.8 \cdot 76.1$	$31.5 \cdot 51.1 \cdot 93.2$	$30.7 \cdot 51.1 \cdot 99.9$	$30.7 \cdot 48.7 \cdot 100.0$	$30.7 \cdot 43.8 \cdot 100.0$
Smoothed	PostNet	$33.6 \cdot 52.8 \cdot 82.3$	$31.4 \cdot 51.2 \cdot 91.6$	$30.9 \cdot 49.4 \cdot 99.1$	$30.7 \cdot 49.3 \cdot 100.0$	$30.7 \cdot 56.0 \cdot 100.0$
+ adv. w.	PriorNet	$37.3 \cdot 60.5 \cdot 84.3$	$34.3 \cdot 59.9 \cdot 87.9$	$32.1 \cdot 61.0 \cdot 97.0$	$30.7 \cdot 69.3 \cdot 100.0$	$30.7 \cdot 68.0 \cdot 100.0$
label	DDNet	$44.8 \cdot 52.2 \cdot 61.0$	$40.2 \cdot 52.6 \cdot 70.0$	$32.5 \cdot 52.4 \cdot 94.6$	$30.7 \cdot 50.3 \cdot 100.0$	$30.7 \cdot 54.6 \cdot 100.0$
attacks	EvNet	$35.8 \cdot 51.2 \cdot 76.7$	$32.9 \cdot 51.0 \cdot 88.5$	$30.7 \cdot 49.5 \cdot 100.0$	$30.7 \cdot 48.5 \cdot 100.0$	$30.7 \cdot 47.7 \cdot 100.0$
Smoothed + adv. w. uncert. attacks	PostNet PriorNet DDNet EvNet	$31.2 \cdot 52.7 \cdot 92.8 \\ 38.3 \cdot 58.2 \cdot 81.5 \\ 44.9 \cdot 52.2 \cdot 60.7 \\ 38.8 \cdot 51.9 \cdot 70.9$	$31.3 \cdot 51.7 \cdot 92.4$ $36.9 \cdot 55.5 \cdot 79.9$ $39.6 \cdot 53.3 \cdot 72.1$ $34.5 \cdot 52.3 \cdot 82.9$	$31.3 \cdot 47.3 \cdot 96.8$ $31.3 \cdot 63.5 \cdot 98.9$ $31.8 \cdot 51.7 \cdot 95.4$ $30.8 \cdot 49.9 \cdot 99.6$	$30.7 \cdot 48.9 \cdot 100.0$ $30.7 \cdot 68.6 \cdot 100.0$ $30.7 \cdot 46.1 \cdot 100.0$ $30.7 \cdot 47.7 \cdot 100.0$	$30.7 \cdot 46.3 \cdot 100.0$ $30.7 \cdot 74.6 \cdot 100.0$ $30.7 \cdot 46.0 \cdot 100.0$ $30.8 \cdot 49.4 \cdot 100.0$

Table 39. Attack detection (PGD label attacks) based on differential entropy. Smoothed DBU models on Sensorless. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$30.7 \cdot 61.9 \cdot 100.0$ $30.7 \cdot 50.1 \cdot 100.0$ $30.7 \cdot 57.5 \cdot 100.0$ $30.7 \cdot 62.0 \cdot 100.0$	30.7 · 60.1 · 100.0 30.7 · 46.5 · 100.0 30.7 · 49.9 · 100.0 30.7 · 59.6 · 100.0	$46.5 \cdot 50.0 \cdot 75.5 \\ 30.7 \cdot 42.3 \cdot 100.0 \\ 30.7 \cdot 45.5 \cdot 100.0 \\ 30.7 \cdot 55.8 \cdot 100.0$	50.0 · 50.0 · 50.0 30.7 · 66.7 · 100.0 30.7 · 50.0 · 100.0 30.7 · 48.3 · 100.0	50.0 · 50.0 · 50.0 30.9 · 79.2 · 100.0 30.7 · 59.3 · 100.0 31.8 · 50.0 · 100.0
Smoothed	PostNet	30.7 · 58.8 · 100.0	30.7 · 58.2 · 100.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0
+ adv. w.	PriorNet	30.7 · 60.2 · 100.0	30.7 · 54.6 · 100.0	30.7 · 45.0 · 100.0	30.7 · 38.0 · 100.0	33.9 · 49.9 · 100.0
label	DDNet	30.7 · 55.4 · 100.0	30.7 · 53.7 · 100.0	30.7 · 44.6 · 100.0	30.7 · 38.8 · 100.0	30.7 · 51.9 · 100.0
attacks	EvNet	30.7 · 62.1 · 100.0	30.7 · 54.3 · 100.0	30.7 · 59.9 · 100.0	30.7 · 62.1 · 100.0	30.7 · 50.0 · 100.0
Smoothed	PostNet	30.7 · 63.0 · 100.0	30.7 · 54.0 · 100.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0
+ adv. w.	PriorNet	30.7 · 58.0 · 100.0	30.7 · 55.6 · 100.0	30.7 · 44.2 · 100.0	30.7 · 53.5 · 100.0	30.7 · 78.5 · 100.0
uncert.	DDNet	30.7 · 55.1 · 100.0	30.7 · 48.2 · 100.0	30.7 · 50.1 · 100.0	30.7 · 52.6 · 100.0	30.7 · 57.0 · 100.0
attacks	EvNet	30.7 · 63.5 · 100.0	30.7 · 54.3 · 100.0	30.7 · 54.2 · 100.0	30.7 · 45.0 · 100.0	30.7 · 50.0 · 100.0

Table 40. Attack detection (PGD label attacks) based on differential entropy. Smoothed DBU models on Segment. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$\begin{array}{c} 30.8 \cdot \textbf{73.5} \cdot 100.0 \\ 30.9 \cdot \textbf{77.1} \cdot 99.9 \\ 31.4 \cdot \textbf{69.6} \cdot 99.5 \\ 30.8 \cdot \textbf{86.2} \cdot 100.0 \end{array}$	$30.8 \cdot 59.9 \cdot 100.0$ $30.8 \cdot 78.1 \cdot 100.0$ $30.8 \cdot 71.2 \cdot 100.0$ $30.8 \cdot 80.3 \cdot 100.0$	30.8 · 60.3 · 100.0 30.8 · 39.5 · 100.0 30.8 · 54.3 · 100.0 30.8 · 54.0 · 100.0	30.8 · 50.2 · 100.0 30.8 · 35.2 · 100.0 30.8 · 35.5 · 100.0 30.8 · 43.3 · 100.0	49.5 · 50.0 · 50.0 30.8 · 41.4 · 100.0 30.8 · 35.7 · 100.0 30.8 · 40.5 · 100.0
Smoothed	PostNet	30.8 · 75.6 · 100.0	30.8 · 69.7 · 100.0	30.8 · 66.5 · 100.0	30.8 · 50.0 · 100.0	50.0 · 50.0 · 50.0
+ adv. w.	PriorNet	31.0 · 74.4 · 99.2	30.8 · 74.0 · 100.0	30.8 · 59.8 · 100.0	30.8 · 56.0 · 100.0	30.8 · 38.8 · 100.0
label	DDNet	31.6 · 68.9 · 99.0	30.8 · 72.9 · 100.0	30.8 · 47.5 · 100.0	30.8 · 32.2 · 100.0	30.8 · 31.8 · 100.0
attacks	EvNet	30.8 · 83.4 · 100.0	30.8 · 87.0 · 100.0	30.8 · 61.9 · 100.0	30.8 · 39.2 · 100.0	30.8 · 41.0 · 100.0
Smoothed	PostNet	30.8 · 73.9 · 100.0	30.8 · 64.5 · 100.0	30.8 · 68.3 · 100.0	33.0 · 50.0 · 100.0	$50.0 \cdot 50.0 \cdot 50.0$
+ adv. w.	PriorNet	31.0 · 73.7 · 99.6	30.8 · 73.1 · 100.0	30.8 · 57.8 · 100.0	30.8 · 44.8 · 100.0	$30.8 \cdot 49.1 \cdot 100.0$
uncert.	DDNet	31.0 · 70.7 · 99.7	30.8 · 70.6 · 100.0	30.8 · 48.6 · 100.0	30.8 · 31.6 · 100.0	$30.8 \cdot 30.9 \cdot 100.0$
attacks	EvNet	30.8 · 85.8 · 100.0	30.8 · 86.7 · 100.0	30.8 · 54.4 · 100.0	30.8 · 45.1 · 100.0	$30.8 \cdot 34.8 \cdot 100.0$

Table 41. Attack detection (FGSM label attacks) based on differential entropy. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$33.1 \cdot 50.3 \cdot 89.9$ $36.0 \cdot 50.8 \cdot 74.6$ $36.4 \cdot 50.4 \cdot 76.4$ $32.9 \cdot 50.3 \cdot 89.7$	$31.0 \cdot 50.2 \cdot 96.9$ $33.0 \cdot 50.4 \cdot 82.8$ $32.8 \cdot 49.9 \cdot 84.6$ $31.4 \cdot 50.2 \cdot 94.0$	$\begin{array}{ccc} 30.7 \cdot \textbf{50.1} \cdot 100.0 \\ 31.2 \cdot \textbf{50.2} \cdot & 95.6 \\ 30.8 \cdot \textbf{50.1} \cdot & 97.9 \\ 30.8 \cdot \textbf{50.1} \cdot & 98.0 \end{array}$	$30.7 \cdot 49.5 \cdot 100.0$ $30.7 \cdot 50.7 \cdot 99.9$ $30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 49.7 \cdot 100.0$	$30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 51.4 \cdot 100.0$ $30.7 \cdot 49.9 \cdot 100.0$ $30.7 \cdot 49.7 \cdot 100.0$
Smoothed + adv. w. label attacks	PostNet PriorNet DDNet EvNet	32.7 · 50.1 · 90.3 35.4 · 52.3 · 78.9 35.5 · 50.6 · 79.3 40.3 · 50.4 · 66.8	$31.1 \cdot 50.3 \cdot 96.4$ $32.9 \cdot 51.3 \cdot 84.5$ $33.4 \cdot 50.3 \cdot 84.2$ $31.4 \cdot 50.3 \cdot 95.9$	$30.7 \cdot 50.1 \cdot 99.7$ $30.7 \cdot 50.3 \cdot 98.7$ $30.8 \cdot 50.1 \cdot 99.2$ $30.7 \cdot 50.2 \cdot 100.0$	$30.7 \cdot 49.8 \cdot 100.0$ $30.7 \cdot 50.7 \cdot 100.0$ $30.7 \cdot 49.9 \cdot 100.0$ $30.7 \cdot 50.1 \cdot 100.0$	$30.7 \cdot 50.5 \cdot 100.0$ $30.8 \cdot 50.2 \cdot 98.2$ $30.7 \cdot 50.1 \cdot 100.0$ $30.7 \cdot 49.6 \cdot 100.0$
Smoothed + adv. w. uncert. attacks	PostNet PriorNet DDNet EvNet	$33.3 \cdot 50.7 \cdot 88.7$ $34.6 \cdot 51.2 \cdot 80.3$ $37.4 \cdot 51.0 \cdot 74.7$ $32.8 \cdot 50.1 \cdot 92.0$	$32.5 \cdot 50.1 \cdot 87.8 \\ 31.4 \cdot 50.7 \cdot 92.8 \\ 33.4 \cdot 50.2 \cdot 83.0 \\ 30.8 \cdot 50.2 \cdot 99.6$	$\begin{array}{c} 30.7 \cdot \textbf{50.1} \cdot 99.8 \\ 30.9 \cdot \textbf{50.2} \cdot 97.7 \\ 30.9 \cdot \textbf{50.1} \cdot 96.9 \\ 30.7 \cdot \textbf{50.4} \cdot 100.0 \end{array}$	$30.7 \cdot 50.5 \cdot 100.0$ $30.7 \cdot 50.0 \cdot 100.0$ $30.8 \cdot 50.1 \cdot 98.1$ $31.2 \cdot 50.2 \cdot 96.0$	$30.7 \cdot 50.2 \cdot 100.0$ $30.7 \cdot 50.1 \cdot 100.0$ $30.7 \cdot 49.9 \cdot 100.0$ $31.0 \cdot 50.0 \cdot 100.0$

Table 42. Attack detection (FGSM label attacks) based on differential entropy. Smoothed DBU models on MNIST. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	$30.9 \cdot 52.3 \cdot 95.6$ $38.1 \cdot 57.7 \cdot 80.8$ $44.7 \cdot 52.0 \cdot 60.9$ $36.5 \cdot 51.7 \cdot 76.0$	$31.5 \cdot 51.2 \cdot 90.8$ $35.8 \cdot 56.6 \cdot 84.0$ $39.4 \cdot 52.9 \cdot 72.5$ $31.5 \cdot 51.1 \cdot 93.2$	$31.1 \cdot 49.8 \cdot 97.0$ $31.5 \cdot 61.7 \cdot 98.3$ $31.6 \cdot 50.8 \cdot 95.2$ $30.7 \cdot 50.9 \cdot 99.9$	30.7 · 48.3 · 100.0 30.8 · 58.9 · 99.2 30.7 · 47.5 · 100.0 30.7 · 48.9 · 100.0	30.7 · 46.5 · 100.0 30.7 · 62.3 · 100.0 30.7 · 46.8 · 100.0 30.7 · 46.2 · 100.0
Smoothed	PostNet	33.5 · 52.6 · 82.2	31.4 · 51.0 · 91.5	30.9 · 49.8 · 99.0	30.7 · 50.1 · 100.0	$30.7 \cdot 54.4 \cdot 100.0$
+ adv. w.	PriorNet	37.3 · 60.6 · 84.3	34.2 · 59.5 · 87.8	32.1 · 60.0 · 96.9	30.7 · 66.3 · 100.0	$30.7 \cdot 63.3 \cdot 100.0$
label	DDNet	44.9 · 52.3 · 61.0	40.3 · 52.8 · 70.2	32.5 · 52.4 · 94.6	30.7 · 50.0 · 100.0	$30.7 \cdot 57.1 \cdot 100.0$
attacks	EvNet	35.8 · 51.5 · 76.7	32.9 · 50.9 · 88.5	30.7 · 50.0 · 100.0	30.7 · 48.9 · 100.0	$30.7 \cdot 48.6 \cdot 100.0$
	PostNet	$31.2 \cdot 52.6 \cdot 92.9$	31.3 · 51.5 · 92.3	31.3 · 48.0 · 96.9	$30.7 \cdot 49.4 \cdot 100.0$	30.7 · 48.1 · 100.0
	PriorNet	$38.3 \cdot 58.3 \cdot 81.4$	36.8 · 55.2 · 79.8	31.3 · 62.5 · 98.9	$30.7 \cdot 64.5 \cdot 100.0$	30.7 · 68.7 · 100.0
	DDNet	$45.0 \cdot 52.3 \cdot 60.9$	39.7 · 53.5 · 72.4	31.8 · 51.7 · 95.4	$30.7 \cdot 46.6 \cdot 100.0$	30.7 · 44.4 · 100.0
	EvNet	$38.8 \cdot 51.8 \cdot 70.8$	34.5 · 52.0 · 82.7	30.8 · 50.0 · 99.6	$30.7 \cdot 49.3 \cdot 100.0$	30.8 · 50.3 · 100.0

Table 43. Attack detection (FGSM label attacks) based on differential entropy. Smoothed DBU models on Sensorless. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet	30.7 · 82.0 · 100.0	30.7 · 88.6 · 100.0	50.0 · 50.0 · 50.1	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0
	PriorNet	30.7 · 51.7 · 100.0	30.7 · 48.2 · 100.0	30.7 · 48.6 · 100.0	30.7 · 68.6 · 100.0	31.4 · 63.7 · 100.0
	DDNet	30.7 · 67.1 · 100.0	30.7 · 58.2 · 100.0	30.7 · 51.7 · 100.0	30.7 · 69.8 · 100.0	30.7 · 73.7 · 100.0
	EvNet	30.7 · 77.9 · 100.0	30.7 · 85.3 · 100.0	30.7 · 90.5 · 100.0	30.8 · 84.3 · 100.0	34.0 · 50.0 · 100.0
Smoothed	PostNet	30.7 · 76.7 · 100.0	30.7 · 78.7 · 100.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	$50.0 \cdot 50.0 \cdot 50.0$
+ adv. w.	PriorNet	30.7 · 63.9 · 100.0	30.7 · 58.4 · 100.0	30.7 · 45.2 · 100.0	30.7 · 43.3 · 100.0	$32.9 \cdot 35.5 \cdot 100.0$
label	DDNet	30.7 · 58.5 · 100.0	30.7 · 75.8 · 100.0	30.7 · 72.6 · 100.0	30.7 · 35.6 · 100.0	$30.7 \cdot 71.5 \cdot 100.0$
attacks	EvNet	30.7 · 80.4 · 100.0	30.7 · 71.5 · 100.0	30.7 · 75.3 · 100.0	30.7 · 78.5 · 100.0	$30.7 \cdot 50.0 \cdot 100.0$
Smoothed	PostNet	30.7 · 77.4 · 100.0	30.7 · 68.0 · 100.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	$\begin{array}{c} \mathbf{50.0 \cdot 50.0} & 50.0 & 50.0 \\ 30.7 \cdot 78.0 \cdot 100.0 \\ 30.7 \cdot 76.0 \cdot 100.0 \\ 30.9 \cdot 50.2 \cdot 100.0 \end{array}$
+ adv. w.	PriorNet	30.7 · 63.8 · 100.0	30.7 · 64.1 · 100.0	30.7 · 46.9 · 100.0	30.7 · 48.9 · 100.0	
uncert.	DDNet	30.7 · 56.5 · 100.0	30.7 · 54.6 · 100.0	30.7 · 59.4 · 100.0	30.7 · 71.8 · 100.0	
attacks	EvNet	30.7 · 71.5 · 100.0	30.7 · 75.7 · 100.0	30.7 · 90.5 · 100.0	30.7 · 54.7 · 100.0	

Table 44. Attack detection (FGSM label attacks) based on differential entropy. Smoothed DBU models on Segment. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model)..

	Att. Rad.	0.1	0.2	0.5	1.0	2.0
Smoothed models	PostNet PriorNet DDNet EvNet	30.8 · 76.9 · 100.0 30.9 · 81.3 · 99.9 31.7 · 73.8 · 99.7 30.8 · 89.1 · 100.0	30.8 · 62.5 · 100.0 30.8 · 85.0 · 100.0 30.8 · 80.5 · 100.0 30.8 · 89.5 · 100.0	30.8 · 59.2 · 100.0 30.8 · 48.7 · 100.0 30.8 · 80.4 · 100.0 30.8 · 75.3 · 100.0	30.8 · 48.7 · 100.0 30.8 · 37.1 · 100.0 30.8 · 72.7 · 100.0 30.8 · 73.1 · 100.0	$\begin{array}{c} \textbf{49.7} \cdot \textbf{50.0} \cdot 50.0 \\ 30.8 \cdot \textbf{43.7} \cdot 100.0 \\ 30.8 \cdot \textbf{70.6} \cdot 100.0 \\ 30.8 \cdot \textbf{83.1} \cdot 100.0 \end{array}$
Smoothed	PostNet	30.8 · 81.0 · 100.0	30.8 · 75.6 · 100.0	30.8 · 56.3 · 100.0	30.8 · 50.0 · 100.0	50.0 · 50.0 · 50.0
+ adv. w.	PriorNet	31.1 · 77.9 · 99.4	30.8 · 76.1 · 100.0	30.8 · 62.4 · 100.0	30.8 · 65.5 · 100.0	30.8 · 53.3 · 100.0
label	DDNet	31.9 · 72.5 · 99.3	30.8 · 82.0 · 100.0	30.8 · 65.7 · 100.0	30.8 · 53.0 · 100.0	30.8 · 61.6 · 100.0
attacks	EvNet	30.8 · 86.4 · 100.0	30.8 · 94.1 · 100.0	30.8 · 78.6 · 100.0	30.8 · 77.7 · 100.0	30.8 · 85.5 · 100.0
Smoothed	PostNet	30.8 · 76.8 · 100.0	30.8 · 64.6 · 100.0	30.8 · 82.9 · 100.0	32.2 · 50.0 · 100.0	$\begin{array}{c} 50.0 \cdot \textbf{50.0} \cdot 50.0 \\ 30.8 \cdot \textbf{61.4} \cdot 100.0 \\ 30.8 \cdot \textbf{43.5} \cdot 100.0 \\ 30.8 \cdot \textbf{96.2} \cdot 100.0 \end{array}$
+ adv. w.	PriorNet	31.1 · 77.6 · 99.7	30.8 · 76.7 · 100.0	30.8 · 69.0 · 100.0	30.8 · 53.1 · 100.0	
uncert.	DDNet	31.1 · 74.3 · 99.8	30.8 · 77.1 · 100.0	30.8 · 76.0 · 100.0	30.8 · 57.0 · 100.0	
attacks	EvNet	30.8 · 88.8 · 100.0	30.8 · 92.6 · 100.0	30.8 · 70.2 · 100.0	30.8 · 62.0 · 100.0	

Table 45. OOD detection based on differential entropy under PGD uncertainty attacks against differential entorpy on ID data and OOD data. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
				ID-	Attack		
	PostNet	$72.1 \cdot 82.7 \cdot 88.0$	$35.0\cdot56.6\cdot97.4$	$31.9 \cdot 65.6 \cdot 99.8$	$30.7 \cdot 50.6 \cdot 100.0$	$30.7 \cdot \textbf{46.9} \cdot 100.0$	$30.7 \cdot 51.6 \cdot 100.0$
Smoothed	PriorNet	$50.2 \cdot 53.1 \cdot 55.9$	$33.5 \cdot 43.3 \cdot 65.3$	$31.3 \cdot 39.7 \cdot 69.1$	$31.3 \cdot 48.3 \cdot 98.2$	$30.7 \cdot 44.4 \cdot 99.9$	$30.7 \cdot 45.4 \cdot 100.0$
models	DDNet	$72.0 \cdot 75.8 \cdot 79.8$	$35.6 \cdot 46.2 \cdot 69.8$	$32.9 \cdot 50.3 \cdot 87.1$	$31.1 \cdot 58.7 \cdot 98.6$	$30.7 \cdot 59.3 \cdot 100.0$	$30.7 \cdot 44.5 \cdot 100.0$
	EvNet	$79.5 \cdot 87.1 \cdot 92.8$	$34.1 \cdot 58.6 \cdot 95.1$	$32.5 \cdot 61.2 \cdot 96.9$	$31.7 \cdot 60.6 \cdot 98.7$	$30.7 \cdot 62.4 \cdot 100.0$	$30.7 \cdot 57.3 \cdot 100.0$
Smoothed	PostNet	-	$35.0 \cdot \textbf{58.5} \cdot 97.7$	$31.2 \cdot \textbf{46.6} \cdot 97.4$	30.8 · 57.7 · 99.7	$30.7 \cdot 49.8 \cdot 100.0$	30.7 · 50.9 · 100.0
+ adv. w.	PriorNet	-	$31.5 \cdot 36.7 \cdot 57.2$	$33.1 \cdot 51.8 \cdot 84.8$	$30.7 \cdot 57.7 \cdot 98.7$	$30.7 \cdot 40.0 \cdot 99.9$	$30.9 \cdot 53.6 \cdot 96.7$
label	DDNet	-	$36.2 \cdot \textbf{50.0} \cdot 78.6$	$32.1 \cdot 41.3 \cdot 70.2$	$30.8 \cdot 56.4 \cdot 100.0$	$30.7 \cdot 49.4 \cdot 100.0$	$30.7 \cdot 54.8 \cdot 100.0$
attacks	EvNet	-	$46.8 \cdot 61.0 \cdot 79.7$	$32.3 \cdot 58.9 \cdot 99.1$	$30.7 \cdot 45.0 \cdot 100.0$	$30.7 \cdot 63.3 \cdot 100.0$	$30.8 \cdot 38.1 \cdot 100.0$
Smoothed	PostNet	-	35.2 · 55.9 · 96.0	34.5 · 59.2 · 94.9	30.7 · 47.0 · 100.0	30.7 · 58.2 · 100.0	30.7 · 42.9 · 100.0
+ adv. w.	PriorNet	-	$31.8 \cdot 38.9 \cdot 64.1$	$31.0 \cdot 41.8 \cdot 87.9$	$30.7 \cdot 42.9 \cdot 99.2$	$30.7 \cdot 48.6 \cdot 100.0$	$30.7 \cdot \textbf{46.6} \cdot 100.0$
uncert.	DDNet	-	$39.7 \cdot 52.1 \cdot 75.7$	$36.4 \cdot \textbf{56.8} \cdot 83.8$	$31.0 \cdot 51.5 \cdot 97.4$	$31.0 \cdot 56.8 \cdot 97.8$	$30.7 \cdot 49.1 \cdot 100.0$
attacks	EvNet	-	$34.8 \cdot 64.9 \cdot 99.6$	$30.8 \cdot 48.9 \cdot 99.8$	$30.7 \cdot 66.8 \cdot 100.0$	$30.9 \cdot 41.5 \cdot 93.8$	$31.1 \cdot 55.1 \cdot 100.0$
				OOI)-Attack		
	PostNet	$72.0 \cdot 82.7 \cdot 88.0$	$35.1\cdot \textbf{56.8}\cdot 97.3$	$32.0 \cdot \textbf{65.8} \cdot 99.8$	$30.7 \cdot 50.7 \cdot 100.0$	$30.7 \cdot \textbf{46.5} \cdot 100.0$	$30.7 \cdot 51.7 \cdot 100.0$
Smoothed	PriorNet	$50.3 \cdot 53.1 \cdot 55.9$	$33.6 \cdot 43.7 \cdot 65.9$	$31.3 \cdot 39.8 \cdot 69.4$	$31.3 \cdot 48.3 \cdot 98.2$	$30.7 \cdot 44.5 \cdot 99.9$	$30.7 \cdot \textbf{46.4} \cdot 100.0$
models	DDNet	$72.0 \cdot 75.8 \cdot 79.8$	$35.6 \cdot 46.2 \cdot 70.0$	$32.9 \cdot 50.1 \cdot 86.7$	$31.1 \cdot 58.8 \cdot 98.6$	$30.7 \cdot 59.3 \cdot 100.0$	$30.7 \cdot 44.6 \cdot 100.0$
	EvNet	$79.5 \cdot 87.1 \cdot 92.8$	$34.1 \cdot 58.8 \cdot 95.2$	$32.6 \cdot 61.2 \cdot 96.9$	$31.7 \cdot 60.5 \cdot 98.7$	$30.7 \cdot 62.4 \cdot 100.0$	$30.7 \cdot 57.6 \cdot 100.0$
Smoothed	PostNet	-	$35.0\cdot \textbf{58.5}\cdot 97.8$	$31.2 \cdot \textbf{46.6} \cdot 97.2$	$30.8 \cdot 57.7 \cdot 99.7$	$30.7 \cdot \textbf{50.2} \cdot 100.0$	$30.7 \cdot 51.5 \cdot 100.0$
+ adv. w.	PriorNet	-	$31.6 \cdot 37.3 \cdot 59.3$	$33.2 \cdot 52.7 \cdot 85.8$	$30.7 \cdot 57.8 \cdot 98.7$	$30.7 \cdot 40.1 \cdot 99.9$	$30.9 \cdot 53.8 \cdot 96.8$
label	DDNet	-	$36.4 \cdot 50.2 \cdot 78.9$	$32.1 \cdot 41.5 \cdot 70.4$	$30.9 \cdot 56.2 \cdot 100.0$	$30.7 \cdot 49.3 \cdot 100.0$	$30.7 \cdot 55.1 \cdot 100.0$
attacks	EvNet	-	$47.2 \cdot 61.1 \cdot 80.0$	$32.4 \cdot 59.1 \cdot 99.1$	$30.7 \cdot 45.0 \cdot 100.0$	$30.7 \cdot 63.2 \cdot 100.0$	$30.8 \cdot 38.0 \cdot 100.0$
Smoothed	PostNet	-	$35.3 \cdot \textbf{56.4} \cdot 96.1$	$34.5 \cdot \textbf{59.0} \cdot 94.9$	$30.7 \cdot \textbf{46.8} \cdot 100.0$	$30.7 \cdot $ 57.8 $\cdot 100.0$	$30.7 \cdot 43.2 \cdot 100.0$
+ adv. w.	PriorNet	-	$31.9 \cdot 39.4 \cdot 65.5$	$31.0 \cdot 42.0 \cdot 88.6$	$30.7 \cdot 42.9 \cdot 99.2$	$30.7 \cdot \textbf{48.4} \cdot 100.0$	$30.7 \cdot 47.1 \cdot 100.0$
uncert.	DDNet	-	$40.2\cdot 52.9\cdot 76.5$	$36.4 \cdot 56.9 \cdot 83.9$	$31.1 \cdot 51.5 \cdot 97.3$	$31.0 \cdot 57.0 \cdot 97.8$	$30.7 \cdot 49.1 \cdot 100.0$
attacks	EvNet	-	34.9 · 64.8 · 99.6	$30.8 \cdot 48.8 \cdot 99.8$	$30.7 \cdot 66.1 \cdot 100.0$	30.9 · 41.6 · 93.6	$31.1 \cdot 54.7 \cdot 100.0$

Table 46. OOD detection based on differential entropy under PGD uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on MNIST. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0			
		ID-Attack								
	PostNet	$59.9 \cdot 91.1 \cdot 98.6$	$61.2 \cdot 97.7 \cdot 99.6$	$64.8 \cdot 94.7 \cdot 99.7$	$31.6 \cdot 64.9 \cdot 99.7$	$30.7\cdot\boldsymbol{63.2}\cdot100.0$	$30.7 \cdot 70.5 \cdot 100.0$			
Smoothed	PriorNet	$99.8 \cdot 99.8 \cdot 99.8$	$99.4 \cdot 99.8 \cdot 99.9$	$98.3 \cdot 99.6 \cdot 99.9$	$48.5 \cdot 91.9 \cdot 99.9$	$31.1 \cdot 74.6 \cdot 99.8$	$30.7\cdot\boldsymbol{67.3}\cdot100.0$			
models	DDNet	$98.5 \cdot 98.6 \cdot 98.7$	$95.0 \cdot 97.6 \cdot 98.9$	$74.7 \cdot 92.0 \cdot 98.2$	$31.4 \cdot 52.0 \cdot 98.5$	$30.7 \cdot 52.0 \cdot 100.0$	$30.7\cdot41.1\cdot100.0$			
	EvNet	$85.7 \cdot 87.5 \cdot 89.2$	$68.9 \cdot 90.4 \cdot 97.7$	$42.5 \cdot 90.2 \cdot 99.6$	$30.7 \cdot 69.8 \cdot 100.0$	$30.7\cdot 50.3\cdot 100.0$	$30.7\cdot45.6\cdot100.0$			
Smoothed	PostNet	-	84.3 · 96.2 · 99.3	50.4 · 89.2 · 99.5	30.9 · 46.2 · 99.4	30.7 · 46.9 · 100.0	30.7 · 62.2 · 100.0			
+ adv. w.	PriorNet	-	$99.7 \cdot 99.9 \cdot 100.0$	$98.7 \cdot 99.8 \cdot 100.0$	$83.3 \cdot 99.1 \cdot 100.0$	$30.7 \cdot 82.6 \cdot 100.0$	$30.7\cdot\boldsymbol{64.8}\cdot100.0$			
label	DDNet	-	$93.6 \cdot 96.9 \cdot 98.5$	$71.2 \cdot 89.1 \cdot 96.9$	$32.3 \cdot 50.3 \cdot 99.0$	$30.7 \cdot 50.7 \cdot 100.0$	$30.7 \cdot 55.7 \cdot 100.0$			
attacks	EvNet	-	$58.2 \cdot 84.4 \cdot 94.3$	$40.9 \cdot 87.4 \cdot 99.2$	$30.7 \cdot 59.4 \cdot 100.0$	$30.7 \cdot 40.3 \cdot 100.0$	$30.7 \cdot 53.2 \cdot 100.0$			
Smoothed	PostNet	-	$58.9 \cdot 96.1 \cdot 99.3$	$59.7 \cdot 96.1 \cdot 99.9$	$31.2 \cdot 48.2 \cdot 95.7$	$30.7 \cdot 42.0 \cdot 100.0$	$30.7 \cdot $ 56.9 \cdot 100.0			
+ adv. w.	PriorNet	-	$99.9 \cdot 100.0 \cdot 100.0$	$96.5 \cdot 99.2 \cdot 99.9$	$49.2\cdot96.9\cdot100.0$	$31.3 \cdot 88.1 \cdot 100.0$	$30.7 \cdot 77.8 \cdot 100.0$			
uncert.	DDNet	-	$95.0 \cdot 97.5 \cdot 98.8$	$80.6 \cdot 94.1 \cdot 98.7$	$31.7 \cdot 55.6 \cdot 98.6$	$30.7 \cdot 52.0 \cdot 100.0$	$30.7 \cdot 47.6 \cdot 100.0$			
attacks	EvNet	-	$66.5 \cdot 91.3 \cdot 98.1$	$48.1 \cdot 84.1 \cdot 97.6$	$30.8 \cdot 49.7 \cdot 99.9$	$30.7 \cdot 37.9 \cdot 100.0$	$30.8 \cdot 63.5 \cdot 100.0$			
		OOD-Attack								
	PostNet	$59.0 \cdot 91.2 \cdot 97.7$	$57.8 \cdot 97.2 \cdot 99.6$	$61.4 \cdot 93.8 \cdot 99.6$	$31.5 \cdot 58.9 \cdot 99.5$	$30.7 \cdot 51.5 \cdot 100.0$	$30.7 \cdot 53.5 \cdot 100.0$			
Smoothed	PriorNet	$99.7 \cdot 99.8 \cdot 99.8$	$99.4 \cdot 99.8 \cdot 99.9$	$98.4 \cdot 99.7 \cdot 100.0$	$60.7 \cdot 96.8 \cdot 100.0$	$33.0 \cdot 88.9 \cdot 100.0$	$30.7 \cdot 87.7 \cdot 100.0$			
models	DDNet	$98.4 \cdot 98.5 \cdot 98.7$	$94.2 \cdot 97.2 \cdot 98.7$	$72.1 \cdot 90.5 \cdot 97.8$	$31.6 \cdot 52.3 \cdot 98.1$	$30.7 \cdot 51.7 \cdot 100.0$	$30.7 \cdot 37.7 \cdot 100.0$			
	EvNet	$83.9 \cdot 85.7 \cdot 88.0$	$63.5 \cdot 88.6 \cdot 97.9$	$40.1 \cdot 87.7 \cdot 99.6$	$30.8 \cdot 68.9 \cdot 100.0$	$30.7 \cdot 43.3 \cdot 100.0$	$30.7 \cdot 36.8 \cdot 100.0$			
Smoothed	PostNet	-	84.7 · 96.1 · 99.4	49.7 · 89.1 · 99.5	30.9 · 45.6 · 99.3	$30.7 \cdot 45.8 \cdot 100.0$	30.7 · 69.1 · 100.0			
+ adv. w.	PriorNet	-	$99.7 \cdot 99.9 \cdot 100.0$	$98.7 \cdot 99.8 \cdot 100.0$	$86.8 \cdot 99.5 \cdot 100.0$	$30.9 \cdot 93.2 \cdot 100.0$	$30.7 \cdot 81.4 \cdot 100.0$			
label	DDNet	-	$93.9 \cdot 97.0 \cdot 98.6$	$72.0 \cdot 89.4 \cdot 97.0$	$33.0 \cdot 52.4 \cdot 98.8$	$30.7 \cdot 51.5 \cdot 100.0$	$30.7 \cdot 60.1 \cdot 100.0$			
attacks	EvNet	-	$59.5 \cdot 85.3 \cdot 94.6$	$40.7 \cdot 86.9 \cdot 99.2$	$30.7 \cdot 57.4 \cdot 100.0$	$30.7 \cdot 39.2 \cdot 100.0$	$30.7 \cdot 49.0 \cdot 100.0$			
Smoothed	PostNet	-	55.7 · 96.1 · 99.3	$58.4 \cdot 95.7 \cdot 99.8$	$31.1 \cdot 44.2 \cdot 93.1$	$30.7 \cdot 41.2 \cdot 100.0$	$30.7 \cdot 48.8 \cdot 100.0$			
+ adv. w.	PriorNet	-	$99.9 \cdot 100.0 \cdot 100.0$	$97.0 \cdot 99.3 \cdot 99.9$	$61.0 \cdot 98.4 \cdot 100.0$	$33.2 \cdot 94.4 \cdot 100.0$	$30.7 \cdot 90.2 \cdot 100.0$			
uncert.	DDNet	-	$95.3 \cdot 97.6 \cdot 98.9$	$82.2 \cdot 94.5 \cdot 98.7$	$32.1 \cdot 56.6 \cdot 98.5$	$30.7 \cdot 48.6 \cdot 100.0$	$30.7\cdot42.9\cdot100.0$			
attacks	EvNet	-	$65.2 \cdot 90.4 \cdot 98.0$	$46.8 \cdot 83.4 \cdot 97.3$	$30.8 \cdot 48.8 \cdot 99.9$	$30.7 \cdot 36.3 \cdot 100.0$	$30.8 \cdot 60.1 \cdot 100.0$			

Table 47. OOD detection based on differential entropy under PGD uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on Sensorless. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0			
		ID-Attack								
	PostNet	$49.3 \cdot 90.4 \cdot 99.8$	$30.7 \cdot 49.2 \cdot 100.0$	$30.7\cdot \textbf{36.0}\cdot 100.0$	$49.2 \cdot 50.0 \cdot 74.9$	$50.0 \cdot 50.0 \cdot 50.0$	$50.0 \cdot 50.0 \cdot 50.0$			
Smoothed	PriorNet	$31.2 \cdot 39.0 \cdot 66.9$	$30.7 \cdot 35.5 \cdot 100.0$	$30.7 \cdot 38.9 \cdot 100.0$	$30.7 \cdot 46.2 \cdot 100.0$	$30.7 \cdot 62.7 \cdot 100.0$	$30.7 \cdot 51.3 \cdot 100.0$			
models	DDNet	$31.0 \cdot 31.5 \cdot 32.7$	$30.7 \cdot 30.8 \cdot 100.0$	$30.7 \cdot 31.8 \cdot 100.0$	$30.7 \cdot 53.6 \cdot 100.0$	$30.7 \cdot 43.9 \cdot 100.0$	$30.7 \cdot 40.5 \cdot 100.0$			
	EvNet	$33.6\cdot \textbf{55.2}\cdot 91.3$	$30.7\cdot \textbf{44.2}\cdot 100.0$	$30.7 \cdot 43.8 \cdot 100.0$	$30.7\cdot 39.3\cdot 100.0$	$30.8 \cdot 51.6 \cdot 100.0$	$32.4 \cdot \textcolor{red}{\textbf{50.0}} \cdot 100.0$			
Smoothed	PostNet	-	$30.7 \cdot 62.4 \cdot 100.0$	30.7 · 39.2 · 100.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.7 \cdot 30.9 \cdot 100.0$	$30.7 \cdot 32.4 \cdot 100.0$	$30.7 \cdot 31.0 \cdot 100.0$	$30.8 \cdot 30.7 \cdot 100.0$	$38.2 \cdot \textbf{48.9} \cdot 100.0$			
label	DDNet	-	$30.7 \cdot 32.2 \cdot 100.0$	$30.7 \cdot 30.9 \cdot 100.0$	$30.7 \cdot 37.1 \cdot 100.0$	$30.7 \cdot 42.1 \cdot 100.0$	$30.7 \cdot 37.7 \cdot 100.0$			
attacks	EvNet	-	$30.7 \cdot 48.9 \cdot 100.0$	$30.7 \cdot 34.0 \cdot 100.0$	$30.7 \cdot 35.6 \cdot 100.0$	$30.7 \cdot 33.6 \cdot 100.0$	$30.7 \cdot 50.0 \cdot 100.0$			
Smoothed	PostNet	-	$30.7 \cdot \textbf{46.0} \cdot 100.0$	$30.7 \cdot \textbf{46.6} \cdot 100.0$	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.7 \cdot 35.8 \cdot 100.0$	$30.7 \cdot 32.1 \cdot 100.0$	$30.7 \cdot 81.6 \cdot 100.0$	$30.8 \cdot 41.7 \cdot 100.0$	$30.7 \cdot 61.9 \cdot 100.0$			
uncert.	DDNet	-	$30.7 \cdot 32.8 \cdot 100.0$	$30.7 \cdot 31.0 \cdot 100.0$	$30.7 \cdot 31.8 \cdot 100.0$	$30.7 \cdot 43.7 \cdot 100.0$	$30.7 \cdot 34.7 \cdot 100.0$			
attacks	EvNet	-	$30.7 \cdot 31.0 \cdot 100.0$	$30.7 \cdot 49.6 \cdot 100.0$	$30.7 \cdot 47.7 \cdot 100.0$	$30.7 \cdot 42.6 \cdot 100.0$	$30.7 \cdot $ 50.0 $\cdot 100.0$			
			OOD-Attack							
	PostNet	$49.3 \cdot 90.4 \cdot 99.8$	$30.8 \cdot 76.4 \cdot 100.0$	$30.7 \cdot 61.3 \cdot 100.0$	$47.7 \cdot 50.0 \cdot 75.1$	$50.0 \cdot 50.0 \cdot 50.0$	$50.0 \cdot 50.0 \cdot 50.0$			
Smoothed	PriorNet	$31.2 \cdot 39.0 \cdot 66.9$	$30.7 \cdot 33.9 \cdot 100.0$	$30.7 \cdot 34.3 \cdot 100.0$	$30.7 \cdot 37.0 \cdot 100.0$	$30.7 \cdot 74.0 \cdot 100.0$	$30.9 \cdot \textbf{78.1} \cdot 100.0$			
models	DDNet	$31.0\cdot 31.5 \cdot 32.7$	$30.7 \cdot 30.7 \cdot 100.0$	$30.7 \cdot 31.8 \cdot 100.0$	$30.7 \cdot 47.7 \cdot 100.0$	$30.7 \cdot 43.8 \cdot 100.0$	$30.7 \cdot 52.5 \cdot 100.0$			
	EvNet	$33.6\cdot \textbf{55.2}\cdot 91.2$	$30.7 \cdot 54.7 \cdot 100.0$	$30.7 \cdot 54.0 \cdot 100.0$	$30.7 \cdot 51.0 \cdot 100.0$	$30.7 \cdot 45.2 \cdot 100.0$	$31.7 \cdot $ 50.0 $\cdot 100.0$			
Smoothed	PostNet	-	$30.7 \cdot 82.2 \cdot 100.0$	$30.7 \cdot \textbf{61.4} \cdot 100.0$	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	$50.0 \cdot 50.0 \cdot 50.0$			
+ adv. w.	PriorNet	-	$30.7 \cdot 31.2 \cdot 100.0$	$30.7 \cdot 31.4 \cdot 99.9$	$30.7 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.7 \cdot 100.0$	$33.8 \cdot 34.0 \cdot 100.0$			
label	DDNet	-	$30.7 \cdot 32.2 \cdot 100.0$	$30.7 \cdot 30.8 \cdot 100.0$	$30.7 \cdot 33.6 \cdot 100.0$	$30.7 \cdot 46.9 \cdot 100.0$	$30.7 \cdot 40.3 \cdot 100.0$			
attacks	EvNet	-	$30.8 \cdot 75.3 \cdot 100.0$	$30.7 \cdot 31.6 \cdot 100.0$	$30.7 \cdot 42.1 \cdot 100.0$	$30.7 \cdot 38.7 \cdot 100.0$	$30.7 \cdot 50.0 \cdot 100.0$			
Smoothed	PostNet	=	$30.7\cdot \textbf{73.7}\cdot 100.0$	$30.7 \cdot \textbf{61.6} \cdot 100.0$	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	$50.0 \cdot 50.0 \cdot 50.0$			
+ adv. w.	PriorNet	-	$30.7 \cdot 35.9 \cdot 100.0$	$30.7 \cdot 30.7 \cdot 100.0$	$30.7 \cdot 39.4 \cdot 100.0$	$30.7 \cdot 36.6 \cdot 100.0$	$30.7 \cdot 97.6 \cdot 100.0$			
uncert.	DDNet	-	$30.7 \cdot 32.1 \cdot 100.0$	$30.7 \cdot 30.8 \cdot 100.0$	$30.7 \cdot 32.2 \cdot 100.0$	$30.7 \cdot 50.7 \cdot 100.0$	$30.7 \cdot 39.8 \cdot 100.0$			
attacks	EvNet	-	$30.7 \cdot 31.3 \cdot 100.0$	$30.8 \cdot 39.7 \cdot 100.0$	$30.7 \cdot 52.2 \cdot 100.0$	$30.7 \cdot 42.3 \cdot 100.0$	$30.7 \cdot $ 50.0 $\cdot 100.0$			

Table 48. OOD detection based on differential entropy under PGD uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on Segment. Column format: guaranteed lowest performance \cdot empirical performance \cdot guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0			
		ID-Attack								
	PostNet	$99.6 \cdot 99.9 \cdot 99.9$	$33.0 \cdot 83.0 \cdot 100.0$	$30.8 \cdot \textbf{43.8} \cdot 100.0$	$30.8 \cdot 31.7 \cdot 100.0$	$30.8 \cdot 40.8 \cdot 100.0$	$41.4 \cdot 50.0 \cdot 50.2$			
Smoothed	PriorNet	$30.8\cdot \textbf{31.0}\cdot 31.4$	$30.8 \cdot 30.8 \cdot 42.6$	$30.8 \cdot 30.8 \cdot 95.5$	$30.8 \cdot 33.1 \cdot 100.0$	$30.8 \cdot 76.4 \cdot 100.0$	$30.8 \cdot 78.7 \cdot 100.0$			
models	DDNet	$30.8\cdot \textbf{30.8}\cdot 30.8$	$30.8 \cdot 30.8 \cdot 32.1$	$30.8 \cdot 30.8 \cdot 69.4$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 31.0 \cdot 100.0$	$30.8 \cdot 33.4 \cdot 100.0$			
	EvNet	$94.9 \cdot 97.2 \cdot 98.3$	$31.1 \cdot 75.8 \cdot 99.9$	$30.8 \cdot 74.2 \cdot 100.0$	$30.8 \cdot 62.9 \cdot 100.0$	$30.8 \cdot 58.1 \cdot 100.0$	$30.8 \cdot \textbf{43.4} \cdot 100.0$			
Smoothed	PostNet	-	31.0 · 70.9 · 100.0	30.8 · 47.1 · 100.0	30.8 · 85.0 · 100.0	30.8 · 50.0 · 100.0	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.8 \cdot 30.8 \cdot 46.0$	$30.8 \cdot 30.8 \cdot 32.7$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.9 \cdot \textbf{30.8} \cdot 100.0$			
label	DDNet	-	$30.8 \cdot 30.8 \cdot 30.8$	$30.8 \cdot 30.8 \cdot 79.5$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 57.3 \cdot 100.0$			
attacks	EvNet	-	$36.3 \cdot 94.3 \cdot 100.0$	$30.8 \cdot 32.2 \cdot 100.0$	$30.8 \cdot 50.2 \cdot 100.0$	$30.8 \cdot 93.9 \cdot 100.0$	$30.8 \cdot 56.3 \cdot 100.0$			
Smoothed	PostNet	-	$30.8 \cdot 49.5 \cdot 100.0$	$30.8 \cdot 34.5 \cdot 100.0$	$30.8 \cdot 96.1 \cdot 100.0$	41.2 · 50.0 · 82.7	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.8 \cdot 31.2 \cdot 62.6$	$30.8 \cdot 30.8 \cdot 32.9$	$30.8 \cdot 30.8 \cdot 88.9$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8\cdot \textbf{30.8}\cdot 100.0$			
uncert.	DDNet	-	$30.8 \cdot 30.8 \cdot 31.2$	$30.8 \cdot 30.8 \cdot 68.9$	$30.8\cdot \textbf{30.8}\cdot 100.0$	$30.8 \cdot 30.9 \cdot 100.0$	$30.8 \cdot 38.6 \cdot 100.0$			
attacks	EvNet	-	$30.9 \cdot 83.5 \cdot 100.0$	$30.8 \cdot 84.0 \cdot 100.0$	$30.8 \cdot 98.6 \cdot 100.0$	$30.8 \cdot 92.8 \cdot 100.0$	$30.8 \cdot 45.6 \cdot 100.0$			
			OOD-Attack							
	PostNet	$99.6 \cdot 99.9 \cdot 99.9$	$31.3 \cdot 95.2 \cdot 100.0$	$30.8 \cdot 48.7 \cdot 100.0$	$30.8 \cdot 34.0 \cdot 100.0$	$30.8 \cdot 41.0 \cdot 100.0$	$41.8 \cdot 50.0 \cdot 50.2$			
Smoothed	PriorNet	$30.8\cdot \textbf{31.0}\cdot 31.4$	$30.8 \cdot 30.8 \cdot 44.7$	$30.8 \cdot 30.8 \cdot 86.3$	$30.8 \cdot 30.9 \cdot 100.0$	$30.8 \cdot 35.7 \cdot 100.0$	$30.8 \cdot \textbf{57.4} \cdot 100.0$			
models	DDNet	$30.8\cdot \textbf{30.8}\cdot 30.8$	$30.8 \cdot 30.8 \cdot 31.9$	$30.8 \cdot 30.8 \cdot 58.3$	$30.8\cdot \textbf{30.8}\cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8\cdot \textbf{30.8}\cdot 100.0$			
	EvNet	$94.9 \cdot 97.2 \cdot 98.3$	$31.4 \cdot 92.5 \cdot 100.0$	$30.8 \cdot 94.2 \cdot 100.0$	$30.8 \cdot 80.4 \cdot 100.0$	$30.8 \cdot 70.2 \cdot 100.0$	$30.8 \cdot \textbf{48.2} \cdot 100.0$			
Smoothed	PostNet	-	$30.8 \cdot 88.7 \cdot 100.0$	$30.8 \cdot 70.9 \cdot 100.0$	$30.8 \cdot 97.2 \cdot 100.0$	30.8 · 50.0 · 100.0	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.8 \cdot 30.9 \cdot 47.2$	$30.8 \cdot 30.8 \cdot 32.5$	$30.8 \cdot 30.8 \cdot 96.2$	$30.8 \cdot 30.8 \cdot 100.0$	$30.9 \cdot \textbf{30.8} \cdot 100.0$			
label	DDNet	-	$30.8 \cdot 30.8 \cdot 30.8$	$30.8 \cdot 30.8 \cdot 73.5$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 34.3 \cdot 100.0$			
attacks	EvNet	-	$35.9 \cdot 95.9 \cdot 100.0$	$30.8 \cdot 36.6 \cdot 100.0$	$30.8 \cdot 45.8 \cdot 100.0$	$30.8 \cdot 75.2 \cdot 100.0$	$30.8 \cdot 93.8 \cdot 100.0$			
Smoothed	PostNet	-	$30.8 \cdot \textbf{64.6} \cdot 100.0$	$30.8 \cdot \textbf{31.9} \cdot 100.0$	$30.8 \cdot \textbf{99.1} \cdot 100.0$	$37.2 \cdot \textbf{50.0} \cdot 100.0$	$49.8 \cdot 50.0 \cdot 50.0$			
+ adv. w.	PriorNet	-	$30.8 \cdot 31.3 \cdot 60.6$	$30.8 \cdot 30.8 \cdot 34.8$	$30.8 \cdot 30.8 \cdot 73.8$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8\cdot \textbf{30.8}\cdot 100.0$			
uncert.	DDNet	-	$30.8 \cdot 30.8 \cdot 31.7$	$30.8 \cdot 30.8 \cdot 64.6$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8\cdot \textbf{30.8}\cdot 100.0$			
attacks	EvNet	-	$31.1 \cdot 90.7 \cdot 100.0$	$30.8 \cdot 96.6 \cdot 100.0$	$30.8 \cdot 98.9 \cdot 100.0$	$30.8 \cdot 97.5 \cdot 100.0$	$30.8 \cdot 34.2 \cdot 100.0$			

Table 49. OOD detection based on differential entropy under FGSM uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on CIFAR10. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0		
		ID-Attack							
	PostNet	$72.2 \cdot 82.7 \cdot 88.0$	$35.0\cdot \textbf{56.5}\cdot 97.5$	$31.9 \cdot 65.5 \cdot 99.8$	$30.7 \cdot 50.6 \cdot 100.0$	$30.7\cdot46.9\cdot100.0$	$30.7\cdot 51.4\cdot 100.0$		
Smoothed	PriorNet	$50.3 \cdot 53.1 \cdot 55.9$	$33.5 \cdot 43.2 \cdot 65.0$	$31.3 \cdot 39.7 \cdot 69.1$	$31.3 \cdot 48.3 \cdot 98.2$	$30.7 \cdot 44.2 \cdot 99.9$	$30.7 \cdot 44.9 \cdot 100.0$		
models	DDNet	$72.0 \cdot 75.8 \cdot 79.8$	$35.5 \cdot \textbf{46.2} \cdot 69.7$	$32.9 \cdot 50.3 \cdot 87.0$	$31.1 \cdot 58.6 \cdot 98.6$	$30.7 \cdot 59.4 \cdot 100.0$	$30.7\cdot44.5\cdot100.0$		
	EvNet	$79.5 \cdot 87.1 \cdot 92.8$	$34.1 \cdot \textbf{58.6} \cdot 95.2$	$32.5 \cdot \textbf{61.1} \cdot 96.9$	$31.7 \cdot 60.6 \cdot 98.8$	$30.7 \cdot 62.6 \cdot 100.0$	$30.7 \cdot 57.3 \cdot 100.0$		
Smoothed	PostNet	-	$35.0 \cdot 58.5 \cdot 97.7$	$31.2 \cdot 46.6 \cdot 97.4$	30.8 · 57.7 · 99.7	$30.7 \cdot 50.1 \cdot 100.0$	$30.7 \cdot 50.6 \cdot 100.0$		
+ adv. w.	PriorNet	-	$31.5 \cdot 36.6 \cdot 56.7$	$33.1 \cdot 51.7 \cdot 84.4$	$30.7 \cdot 57.5 \cdot 98.7$	$30.7 \cdot 40.1 \cdot 99.9$	$30.9 \cdot 53.5 \cdot 96.7$		
label	DDNet	-	$36.2 \cdot 50.0 \cdot 78.5$	$32.1\cdot41.3\cdot70.1$	$30.9 \cdot 56.3 \cdot 100.0$	$30.7 \cdot 49.5 \cdot 100.0$	$30.7 \cdot 54.9 \cdot 100.0$		
attacks	EvNet	-	$46.8 \cdot 60.9 \cdot 79.6$	$32.3 \cdot 58.9 \cdot 99.1$	$30.7 \cdot 45.1 \cdot 100.0$	$30.7 \cdot 63.1 \cdot 100.0$	$30.8 \cdot 38.1 \cdot 100.0$		
Smoothed	PostNet	-	$35.2 \cdot 56.0 \cdot 95.9$	$34.5 \cdot 59.0 \cdot 94.8$	30.7 · 47.0 · 100.0	$30.7 \cdot 57.2 \cdot 100.0$	$30.7 \cdot 42.7 \cdot 100.0$		
+ adv. w.	PriorNet	-	$31.8 \cdot \textbf{38.8} \cdot 64.0$	$31.0 \cdot 41.7 \cdot 87.4$	$30.7 \cdot 42.9 \cdot 99.3$	$30.7 \cdot 48.5 \cdot 100.0$	$30.7 \cdot \textbf{46.8} \cdot 100.0$		
uncert.	DDNet	-	$39.6 \cdot 52.0 \cdot 75.6$	$36.4 \cdot \textbf{56.8} \cdot 83.8$	$31.0 \cdot 51.4 \cdot 97.3$	$31.0 \cdot 56.9 \cdot 97.7$	$30.7 \cdot 49.2 \cdot 100.0$		
attacks	EvNet	-	$34.8 \cdot 64.9 \cdot 99.7$	$30.8 \cdot 48.9 \cdot 99.8$	$30.7 \cdot 66.4 \cdot 100.0$	$30.9 \cdot 41.6 \cdot 93.6$	$31.1 \cdot 55.7 \cdot 100.0$		
				001)-Attack				
	PostNet	$72.1 \cdot 82.7 \cdot 88.0$	$35.1\cdot \textbf{56.8}\cdot 97.3$	$31.9 \cdot \textbf{65.8} \cdot 99.8$	$30.7 \cdot \textbf{50.8} \cdot 100.0$	$30.7\cdot46.5\cdot100.0$	$30.7\cdot 51.5\cdot 100.0$		
Smoothed	PriorNet	$50.3 \cdot \textbf{53.1} \cdot 55.9$	$33.6 \cdot 43.7 \cdot 65.9$	$31.3 \cdot 39.8 \cdot 69.4$	$31.3 \cdot 48.3 \cdot 98.2$	$30.7 \cdot 44.4 \cdot 99.9$	$30.7 \cdot 45.9 \cdot 100.0$		
models	DDNet	$72.0 \cdot 75.8 \cdot 79.8$	$35.6 \cdot \textbf{46.1} \cdot 70.0$	$32.9 \cdot 50.1 \cdot 86.7$	$31.1 \cdot 58.7 \cdot 98.6$	$30.7 \cdot 59.3 \cdot 100.0$	$30.7 \cdot 44.6 \cdot 100.0$		
	EvNet	$79.5 \cdot 87.1 \cdot 92.8$	$34.1 \cdot 58.8 \cdot 95.2$	$32.6 \cdot 61.3 \cdot 96.9$	$31.7 \cdot 60.5 \cdot 98.8$	$30.7 \cdot 62.2 \cdot 100.0$	$30.7 \cdot 57.7 \cdot 100.0$		
Smoothed	PostNet	-	$35.0 \cdot 58.4 \cdot 97.9$	$31.2\cdot46.6\cdot97.3$	$30.8 \cdot 57.7 \cdot 99.7$	$30.7 \cdot \textbf{50.1} \cdot 100.0$	$30.7 \cdot 51.4 \cdot 100.0$		
+ adv. w.	PriorNet	-	$31.6 \cdot 37.3 \cdot 59.2$	$33.2 \cdot 52.6 \cdot 85.8$	$30.7 \cdot 57.8 \cdot 98.7$	$30.7 \cdot 39.8 \cdot 99.9$	$30.9 \cdot 53.7 \cdot 96.8$		
label	DDNet	-	$36.4 \cdot 50.2 \cdot 78.8$	$32.1\cdot41.5\cdot70.5$	$30.8 \cdot 56.2 \cdot 100.0$	$30.7 \cdot 49.2 \cdot 100.0$	$30.7 \cdot 55.0 \cdot 100.0$		
attacks	EvNet	-	$47.2 \cdot 61.0 \cdot 79.9$	$32.4 \cdot 59.1 \cdot 99.1$	$30.7 \cdot 45.1 \cdot 100.0$	$30.7 \cdot 63.1 \cdot 100.0$	$30.8 \cdot 38.0 \cdot 100.0$		
Smoothed	PostNet	-	$35.3 \cdot \textbf{56.3} \cdot 96.1$	$34.5 \cdot \textbf{59.1} \cdot 94.9$	$30.7 \cdot 46.9 \cdot 100.0$	$30.7 \cdot $ 57.8 $\cdot 100.0$	$30.7 \cdot 43.1 \cdot 100.0$		
+ adv. w.	PriorNet	-	$31.9 \cdot 39.4 \cdot 65.4$	$31.0 \cdot 42.0 \cdot 88.7$	$30.7 \cdot 42.9 \cdot 99.2$	$30.7 \cdot 48.3 \cdot 100.0$	$30.7 \cdot 47.2 \cdot 100.0$		
uncert.	DDNet	-	$40.1 \cdot 52.8 \cdot 76.5$	$36.5 \cdot 56.9 \cdot 83.9$	$31.1 \cdot 51.5 \cdot 97.3$	$31.0 \cdot 57.0 \cdot 97.8$	$30.7 \cdot 48.7 \cdot 100.0$		
attacks	EvNet	-	$34.9 \cdot 65.0 \cdot 99.6$	$30.8 \cdot \textbf{48.8} \cdot 99.8$	$30.7 \cdot 66.6 \cdot 100.0$	$30.9 \cdot 41.1 \cdot 93.4$	$31.1 \cdot 55.3 \cdot 100.0$		

Table 50. OOD detection based on differential entropy under FGSM uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on MNIST. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0			
		ID-Attack								
	PostNet	$59.9 \cdot 91.3 \cdot 98.6$	$61.1 \cdot 97.7 \cdot 99.7$	$65.1 \cdot 94.8 \cdot 99.7$	$31.6 \cdot 64.8 \cdot 99.7$	$30.7\cdot\boldsymbol{62.4}\cdot100.0$	$30.7 \cdot \textbf{68.6} \cdot 100.0$			
Smoothed	PriorNet	$99.8 \cdot 99.8 \cdot 99.8$	$99.4 \cdot 99.8 \cdot 99.9$	$98.4 \cdot 99.7 \cdot 99.9$	$49.8 \cdot 92.7 \cdot 99.9$	$31.3 \cdot 76.6 \cdot 99.8$	$30.7 \cdot 71.8 \cdot 100.0$			
models	DDNet	$98.5 \cdot 98.6 \cdot 98.7$	$95.0 \cdot 97.6 \cdot 98.9$	$74.4 \cdot 91.9 \cdot 98.2$	$31.4 \cdot 52.0 \cdot 98.5$	$30.7 \cdot 51.8 \cdot 100.0$	$30.7\cdot40.2\cdot100.0$			
	EvNet	$85.7 \cdot 87.5 \cdot 89.2$	$69.0 \cdot 90.4 \cdot 97.7$	$42.5 \cdot 90.2 \cdot 99.6$	$30.7 \cdot 70.1 \cdot 100.0$	$30.7\cdot 50.0\cdot 100.0$	$30.7\cdot43.9\cdot100.0$			
Smoothed	PostNet	-	84.4 · 96.3 · 99.4	50.6 · 89.3 · 99.5	30.9 · 46.3 · 99.4	$30.7 \cdot 46.3 \cdot 100.0$	30.7 · 63.3 · 100.0			
+ adv. w.	PriorNet	-	$99.7 \cdot 99.9 \cdot 100.0$	$98.7 \cdot 99.8 \cdot 100.0$	$84.1 \cdot 99.2 \cdot 100.0$	$30.7 \cdot 84.6 \cdot 100.0$	$30.7\cdot\boldsymbol{68.1}\cdot100.0$			
label	DDNet	-	$93.6 \cdot 96.9 \cdot 98.5$	$71.0 \cdot 89.0 \cdot 96.9$	$32.3 \cdot 50.4 \cdot 99.0$	$30.7 \cdot 51.1 \cdot 100.0$	$30.7\cdot54.1\cdot100.0$			
attacks	EvNet	-	$58.2 \cdot 84.5 \cdot 94.3$	$40.9 \cdot 87.2 \cdot 99.2$	$30.7 \cdot 59.3 \cdot 100.0$	$30.7 \cdot 39.7 \cdot 100.0$	$30.7 \cdot 52.7 \cdot 100.0$			
Smoothed	PostNet	-	58.6 · 96.1 · 99.3	59.9 · 96.2 · 99.9	31.2 · 47.6 · 95.5	$30.7 \cdot 41.8 \cdot 100.0$	$30.7 \cdot 55.4 \cdot 100.0$			
+ adv. w.	PriorNet	-	$99.9 \cdot 100.0 \cdot 100.0$	$96.6 \cdot 99.2 \cdot 99.9$	$50.3 \cdot 97.1 \cdot 100.0$	$31.7 \cdot 89.7 \cdot 100.0$	$30.7 \cdot \textbf{81.8} \cdot 100.0$			
uncert.	DDNet	-	$95.0 \cdot 97.5 \cdot 98.8$	$80.5 \cdot 94.0 \cdot 98.6$	$31.7 \cdot 55.6 \cdot 98.6$	$30.7 \cdot 52.0 \cdot 100.0$	$30.7\cdot49.5\cdot100.0$			
attacks	EvNet	-	$66.5 \cdot 91.4 \cdot 98.1$	$48.5 \cdot \textbf{84.5} \cdot 97.6$	$30.8 \cdot 49.3 \cdot 99.9$	$30.7 \cdot 37.3 \cdot 100.0$	$30.8 \cdot 62.0 \cdot 100.0$			
		OOD-Attack								
	PostNet	$59.2 \cdot 91.3 \cdot 97.7$	$57.9 \cdot 97.2 \cdot 99.6$	$61.4 \cdot 93.8 \cdot 99.6$	$31.5 \cdot 59.1 \cdot 99.5$	$30.7\cdot 52.4\cdot 100.0$	$30.7\cdot 53.9\cdot 100.0$			
Smoothed	PriorNet	$99.7 \cdot 99.8 \cdot 99.8$	$99.4 \cdot 99.8 \cdot 99.9$	$98.3 \cdot 99.7 \cdot 100.0$	$60.4 \cdot 96.6 \cdot 100.0$	$32.8 \cdot 88.2 \cdot 99.9$	$30.7 \cdot 86.1 \cdot 100.0$			
models	DDNet	$98.4 \cdot 98.5 \cdot 98.7$	$94.3 \cdot 97.2 \cdot 98.7$	$72.2 \cdot 90.6 \cdot 97.8$	$31.6 \cdot 52.2 \cdot 98.1$	$30.7 \cdot 51.8 \cdot 100.0$	$30.7 \cdot \textbf{38.5} \cdot 100.0$			
	EvNet	$83.9 \cdot 85.7 \cdot 88.0$	$63.6 \cdot 88.6 \cdot 97.9$	$40.1 \cdot 87.6 \cdot 99.6$	$30.8 \cdot 69.2 \cdot 100.0$	$30.7\cdot43.5\cdot100.0$	$30.7\cdot\boldsymbol{37.4}\cdot100.0$			
Smoothed	PostNet	-	84.4 · 96.2 · 99.4	49.7 · 89.1 · 99.5	30.9 · 45.6 · 99.3	$30.7 \cdot 46.2 \cdot 100.0$	$30.7 \cdot 68.1 \cdot 100.0$			
+ adv. w.	PriorNet	-	$99.7 \cdot 99.9 \cdot 100.0$	$98.7 \cdot 99.8 \cdot 100.0$	$86.3 \cdot 99.4 \cdot 100.0$	$30.9 \cdot 91.9 \cdot 100.0$	$30.7\cdot \textbf{77.5}\cdot 100.0$			
label	DDNet	-	$93.9 \cdot 97.0 \cdot 98.6$	$72.1 \cdot 89.5 \cdot 97.0$	$33.0 \cdot 52.3 \cdot 98.8$	$30.7 \cdot 51.5 \cdot 100.0$	$30.7\cdot\boldsymbol{60.4}\cdot100.0$			
attacks	EvNet	-	$59.4 \cdot 85.6 \cdot 94.6$	$40.7 \cdot 86.7 \cdot 99.2$	$30.7 \cdot 57.3 \cdot 100.0$	$30.7 \cdot 39.4 \cdot 100.0$	$30.7\cdot49.0\cdot100.0$			
Smoothed	PostNet	-	55.8 · 96.1 · 99.3	58.4 · 95.7 · 99.8	31.1 · 44.6 · 93.3	$30.7 \cdot 41.4 \cdot 100.0$	$30.7 \cdot 50.1 \cdot 100.0$			
+ adv. w.	PriorNet	-	$99.9 \cdot 100.0 \cdot 100.0$	$96.9 \cdot 99.3 \cdot 99.9$	$60.3 \cdot 98.2 \cdot 100.0$	$33.0\cdot93.5\cdot100.0$	$30.7 \cdot \textbf{87.8} \cdot 100.0$			
uncert.	DDNet	-	$95.3 \cdot 97.6 \cdot 98.9$	$82.3 \cdot 94.5 \cdot 98.7$	$32.1 \cdot 56.3 \cdot 98.5$	$30.7\cdot48.9\cdot100.0$	$30.7\cdot43.4\cdot100.0$			
attacks	EvNet	-	$65.3 \cdot 90.3 \cdot 97.9$	$46.9 \cdot 83.1 \cdot 97.3$	$30.8 \cdot 48.8 \cdot 99.9$	$30.7\cdot \textbf{36.6}\cdot 100.0$	$30.8 \cdot 60.7 \cdot 100.0$			

Table 51. OOD detection based on differential entropy under FGSM uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on Sensorless. Column format: guaranteed lowest performance · empirical performance · guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model)..

	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0			
	ID-Attack									
	PostNet	$49.3 \cdot 90.4 \cdot 99.8$	$30.7 \cdot 50.3 \cdot 100.0$	$30.7\cdot36.6\cdot100.0$	$49.1 \cdot 50.0 \cdot 74.9$	$50.0 \cdot 50.0 \cdot 50.0$	$50.0 \cdot 50.0 \cdot 50.0$			
Smoothed	PriorNet	$31.2 \cdot 39.0 \cdot 66.9$	$30.7 \cdot 40.1 \cdot 100.0$	$30.7 \cdot 48.2 \cdot 100.0$	$30.7 \cdot 54.2 \cdot 100.0$	$30.7 \cdot 46.3 \cdot 100.0$	$30.7 \cdot 47.6 \cdot 100.0$			
models	DDNet	$31.0 \cdot 31.5 \cdot 32.7$	$30.7 \cdot 31.2 \cdot 100.0$	$30.7 \cdot 35.3 \cdot 100.0$	$30.7 \cdot 55.7 \cdot 100.0$	$30.7 \cdot 42.4 \cdot 100.0$	$30.7\cdot40.4\cdot100.0$			
	EvNet	$33.6 \cdot 55.1 \cdot 91.3$	$30.7 \cdot 39.1 \cdot 100.0$	$30.7 \cdot 37.1 \cdot 100.0$	$30.7 \cdot 35.4 \cdot 100.0$	$30.8 \cdot 52.1 \cdot 100.0$	$32.5 \cdot $ 50.0 $\cdot 100.0$			
Smoothed	PostNet	-	$30.7 \cdot 60.8 \cdot 100.0$	$30.7 \cdot 40.7 \cdot 100.0$	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.7 \cdot 31.3 \cdot 100.0$	$30.7 \cdot 32.9 \cdot 100.0$	$30.7\cdot40.1\cdot100.0$	$30.8 \cdot 31.1 \cdot 100.0$	$38.1 \cdot 91.0 \cdot 100.0$			
label	DDNet	-	$30.7 \cdot 34.3 \cdot 100.0$	$30.7 \cdot 33.9 \cdot 100.0$	$30.7 \cdot 38.2 \cdot 100.0$	$30.7 \cdot 63.6 \cdot 100.0$	$30.7 \cdot 41.8 \cdot 100.0$			
attacks	EvNet	-	$30.8 \cdot 41.0 \cdot 100.0$	$30.7 \cdot 34.2 \cdot 100.0$	$30.7 \cdot 38.0 \cdot 100.0$	$30.7 \cdot 39.0 \cdot 100.0$	$30.7 \cdot 50.0 \cdot 100.0$			
Smoothed	PostNet	-	$30.7\cdot \textbf{46.1}\cdot 100.0$	$30.7 \cdot \textbf{46.8} \cdot 100.0$	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.7 \cdot 36.5 \cdot 100.0$	$30.7\cdot34.4\cdot100.0$	$30.7 \cdot 77.8 \cdot 100.0$	$30.8\cdot \textbf{53.0}\cdot 100.0$	$30.7 \cdot 39.2 \cdot 100.0$			
uncert.	DDNet	-	$30.7 \cdot 36.0 \cdot 100.0$	$30.7 \cdot 37.7 \cdot 100.0$	$30.7 \cdot 41.0 \cdot 100.0$	$30.7 \cdot 42.3 \cdot 100.0$	$30.7 \cdot 39.0 \cdot 100.0$			
attacks	EvNet	=	$30.7 \cdot 31.3 \cdot 100.0$	$30.7 \cdot 43.3 \cdot 100.0$	$30.7 \cdot 36.3 \cdot 100.0$	$30.7 \cdot 43.2 \cdot 100.0$	$30.7 \cdot 50.0 \cdot 100.0$			
				00	D-Attack					
	PostNet	$49.3 \cdot 90.4 \cdot 99.8$	$30.8 \cdot 75.3 \cdot 100.0$	$30.7 \cdot 68.5 \cdot 100.0$	$46.1 \cdot 50.0 \cdot 74.8$	$50.0 \cdot 50.0 \cdot 50.0$	$50.0 \cdot 50.0 \cdot 50.0$			
Smoothed	PriorNet	$31.2 \cdot 38.9 \cdot 67.0$	$30.7 \cdot 34.1 \cdot 100.0$	$30.7 \cdot 35.7 \cdot 100.0$	$30.7 \cdot 35.0 \cdot 100.0$	$30.7 \cdot 77.6 \cdot 100.0$	$30.8 \cdot 95.3 \cdot 100.0$			
models	DDNet	$31.0 \cdot 31.5 \cdot 32.7$	$30.7 \cdot 30.8 \cdot 100.0$	$30.7 \cdot 33.1 \cdot 100.0$	$30.7 \cdot 65.7 \cdot 100.0$	$30.7 \cdot 71.8 \cdot 100.0$	$30.7 \cdot 71.5 \cdot 100.0$			
	EvNet	$33.6 \cdot 55.2 \cdot 91.4$	$30.7 \cdot 64.7 \cdot 100.0$	$30.7 \cdot 69.6 \cdot 100.0$	$30.7 \cdot 78.9 \cdot 100.0$	$30.7 \cdot 67.2 \cdot 100.0$	$32.9 \cdot 50.0 \cdot 100.0$			
Smoothed	PostNet	-	$30.7\cdot86.0\cdot100.0$	$30.7 \cdot 86.6 \cdot 100.0$	$50.0 \cdot 50.0 \cdot 50.0$	$50.0 \cdot 50.0 \cdot 50.0$	$50.0 \cdot 50.0 \cdot 50.0$			
+ adv. w.	PriorNet	-	$30.7 \cdot 31.0 \cdot 99.9$	$30.7 \cdot 31.2 \cdot 98.9$	$30.7 \cdot 30.7 \cdot 100.0$	$30.8 \cdot 30.7 \cdot 100.0$	$36.1 \cdot 35.3 \cdot 100.0$			
label	DDNet	-	$30.7 \cdot 37.2 \cdot 100.0$	$30.7 \cdot 31.1 \cdot 100.0$	$30.7 \cdot 37.1 \cdot 100.0$	$30.7 \cdot 50.5 \cdot 100.0$	$30.7 \cdot 84.6 \cdot 100.0$			
attacks	EvNet	-	$30.8 \cdot 82.5 \cdot 100.0$	$30.7 \cdot 51.7 \cdot 100.0$	$30.7 \cdot 91.5 \cdot 100.0$	$30.7 \cdot 70.0 \cdot 100.0$	$30.9 \cdot \textbf{50.0} \cdot 100.0$			
Smoothed	PostNet	-	$30.7 \cdot 78.5 \cdot 100.0$	$30.7 \cdot 67.1 \cdot 100.0$	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0	50.0 · 50.0 · 50.0			
+ adv. w.	PriorNet	-	$30.7 \cdot 35.8 \cdot 100.0$	$30.7 \cdot 30.7 \cdot 100.0$	$30.7 \cdot 39.0 \cdot 100.0$	$30.7 \cdot 58.5 \cdot 100.0$	$30.7 \cdot 100.0 \cdot 100.0$			
uncert.	DDNet	-	$30.7 \cdot 40.8 \cdot 100.0$	$30.7 \cdot 33.1 \cdot 100.0$	$30.7 \cdot 30.8 \cdot 100.0$	$30.7 \cdot 34.3 \cdot 100.0$	$30.7 \cdot 35.2 \cdot 100.0$			
attacks	EvNet	-	$30.7 \cdot 32.7 \cdot 100.0$	$30.8 \cdot 50.2 \cdot 100.0$	$30.7 \cdot 99.6 \cdot 100.0$	$30.7 \cdot 58.7 \cdot 100.0$	$30.7 \cdot 50.0 \cdot 100.0$			

Table 52. OOD detection based on differential entropy under FGSM uncertainty attacks against differential entropy on ID data and OOD data. Smoothed DBU models on Segment. Column format: guaranteed lowest performance \cdot empirical performance \cdot guaranteed highest performance (blue: normally/adversarially trained smooth classifier is more robust than the base model).

	,	-	•				
	Att. Rad.	0.0	0.1	0.2	0.5	1.0	2.0
				ID-A	Attack		
	PostNet	$99.6 \cdot 99.9 \cdot 99.9$	$33.1\cdot 78.8\cdot 100.0$	$30.8\cdot46.2\cdot100.0$	$30.8 \cdot 34.2 \cdot 100.0$	$30.8 \cdot 41.4 \cdot 100.0$	$41.5 \cdot 50.0 \cdot 50.2$
Smoothed	PriorNet	$30.9 \cdot 31.0 \cdot 31.4$	$30.8 \cdot 30.8 \cdot 39.3$	$30.8 \cdot 30.8 \cdot 94.7$	$30.8 \cdot 41.2 \cdot 100.0$	$30.8 \cdot 92.7 \cdot 100.0$	$30.8 \cdot 79.9 \cdot 100.0$
models	DDNet	$30.8\cdot \textbf{30.8}\cdot 30.8$	$30.8 \cdot 30.8 \cdot 31.8$	$30.8 \cdot 30.8 \cdot 66.8$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 32.6 \cdot 100.0$	$30.8 \cdot 38.2 \cdot 100.0$
	EvNet	$94.9 \cdot 97.2 \cdot 98.2$	$31.0 \cdot 73.1 \cdot 100.0$	$30.8 \cdot 72.3 \cdot 100.0$	$30.8 \cdot 57.1 \cdot 100.0$	$30.8\cdot\boldsymbol{63.3}\cdot100.0$	$30.8\cdot49.6\cdot100.0$
Smoothed	PostNet	-	31.0 · 62.9 · 100.0	30.8 · 47.1 · 100.0	30.8 · 90.0 · 100.0	30.8 · 50.0 · 100.0	50.0 · 50.0 · 50.0
+ adv. w.	PriorNet	-	$30.8 \cdot 30.8 \cdot 43.5$	$30.8 \cdot 30.8 \cdot 32.5$	$30.8\cdot \textbf{30.9}\cdot 100.0$	$30.8\cdot \textbf{30.9}\cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$
label	DDNet	-	$30.8 \cdot 30.8 \cdot 30.8$	$30.8 \cdot 30.8 \cdot 76.1$	$30.8 \cdot 30.9 \cdot 100.0$	$30.8 \cdot 34.8 \cdot 100.0$	$30.8 \cdot 53.0 \cdot 100.0$
attacks	EvNet	-	$35.5 \cdot 93.5 \cdot 100.0$	$30.8 \cdot 31.8 \cdot 100.0$	$30.8 \cdot 48.7 \cdot 100.0$	$30.8 \cdot 93.8 \cdot 100.0$	$30.8 \cdot 63.7 \cdot 100.0$
Smoothed	PostNet	-	$30.8 \cdot 47.5 \cdot 100.0$	$30.8 \cdot 37.5 \cdot 100.0$	$30.8 \cdot 92.9 \cdot 100.0$	41.1 · 50.0 · 97.3	50.0 · 50.0 · 50.0
+ adv. w.	PriorNet	-	$30.8 \cdot 31.1 \cdot 60.8$	$30.8 \cdot 30.8 \cdot 32.3$	$30.8 \cdot 30.8 \cdot 90.3$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 36.3 \cdot 100.0$
uncert.	DDNet	-	$30.8 \cdot 30.8 \cdot 31.0$	$30.8 \cdot 30.8 \cdot 66.8$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 31.2 \cdot 100.0$	$30.8 \cdot 57.2 \cdot 100.0$
attacks	EvNet	=	$30.9 \cdot 80.3 \cdot 100.0$	$30.8 \cdot 78.1 \cdot 100.0$	$30.8 \cdot 99.4 \cdot 100.0$	$30.8 \cdot 97.7 \cdot 100.0$	$30.8 \cdot 41.5 \cdot 100.0$
				OOD-	-Attack		
	PostNet	$99.6 \cdot 99.9 \cdot 99.9$	$31.2\cdot94.3\cdot100.0$	$30.8 \cdot 44.8 \cdot 100.0$	$30.8 \cdot \textbf{36.8} \cdot 100.0$	$30.8\cdot39.9\cdot100.0$	$44.3 \cdot 50.0 \cdot 50.0$
Smoothed	PriorNet	$30.9 \cdot 31.0 \cdot 31.4$	$30.8 \cdot 30.8 \cdot 42.0$	$30.8 \cdot 30.8 \cdot 80.4$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 37.5 \cdot 100.0$	$30.8 \cdot 94.9 \cdot 100.0$
models	DDNet	$30.8\cdot \textbf{30.8}\cdot 30.8$	$30.8 \cdot 30.8 \cdot 31.5$	$30.8 \cdot 30.8 \cdot 48.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$
	EvNet	$94.9 \cdot 97.2 \cdot 98.3$	$31.3 \cdot 92.1 \cdot 100.0$	$30.8 \cdot 90.8 \cdot 100.0$	$30.8 \cdot 89.6 \cdot 100.0$	$30.8 \cdot 89.8 \cdot 100.0$	$30.8 \cdot 87.3 \cdot 100.0$
Smoothed	PostNet	-	$30.8 \cdot 85.3 \cdot 100.0$	$30.8 \cdot \textbf{85.9} \cdot 100.0$	$30.8 \cdot \textbf{78.8} \cdot 100.0$	$30.9 \cdot \textbf{50.0} \cdot 100.0$	$50.0 \cdot 50.0 \cdot 50.0$
+ adv. w.	PriorNet	-	$30.8 \cdot 30.8 \cdot 45.0$	$30.8 \cdot 30.8 \cdot 32.1$	$30.8 \cdot 30.8 \cdot 90.3$	$30.8 \cdot 30.8 \cdot 100.0$	$31.0 \cdot 30.8 \cdot 100.0$
label	DDNet	-	$30.8 \cdot 30.8 \cdot 30.8$	$30.8 \cdot 30.8 \cdot 64.9$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 79.4 \cdot 100.0$
attacks	EvNet	-	$35.4 \cdot 95.0 \cdot 100.0$	$30.8 \cdot 35.2 \cdot 100.0$	$30.8 \cdot 51.9 \cdot 100.0$	$30.8 \cdot 80.0 \cdot 100.0$	$30.8 \cdot 99.9 \cdot 100.0$
Smoothed	PostNet	-	$30.8 \cdot 63.4 \cdot 100.0$	$30.8 \cdot 31.7 \cdot 100.0$	$30.8 \cdot 98.4 \cdot 100.0$	$33.2 \cdot 50.0 \cdot 100.0$	50.0 · 50.0 · 50.0
+ adv. w.	PriorNet	-	$30.8 \cdot 31.1 \cdot 58.0$	$30.8 \cdot 30.8 \cdot 34.1$	$30.8 \cdot 30.8 \cdot 66.8$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$
uncert.	DDNet	-	$30.8 \cdot 30.8 \cdot 31.2$	$30.8 \cdot 30.8 \cdot 61.5$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8 \cdot 30.8 \cdot 100.0$	$30.8\cdot \textbf{30.8}\cdot 100.0$
	EvNet		$31.0 \cdot 89.0 \cdot 100.0$	$30.8 \cdot 96.2 \cdot 100.0$	$30.8 \cdot 99.6 \cdot 100.0$	$30.8 \cdot 99.6 \cdot 100.0$	$30.8 \cdot 69.7 \cdot 100.0$

6.5. Visualization of differential entropy distributions on ID data and OOD data

The following Figures visualize the differential entropy distribution for ID data and OOD data for all models with standard training. We used label attacks and uncertainty attacks for CIFAR10 and MNIST. Thus, they show how well the DBU models separate on clean and perturbed ID data and OOD data.

Figures 4 and 5 visualizes the differential entropy distribution of ID data and OOD data under label attacks. On CIFAR10, PriorNet and DDNet can barely distinguish between clean ID and OOD data. We observe a better ID/OOD distinction for PostNet and EvNet for clean data. However, we do not observe for any model an increase of the uncertainty estimates on label attacked data. Even worse, PostNet, PriorNet and DDNet seem to assign higher confidence on class label attacks. On MNIST, models show a slightly better behavior. They are capable to assign a higher uncertainty to label attacks up to some attack radius.

Figures 6, 7, 8 and 9 visualizes the differential entropy distribution of ID data and OOD data under uncertainty attacks. For both CIFAR10 and MNIST data sets, we observed that uncertainty estimations of all models can be manipulated. That is, OOD uncertainty attacks can shift the OOD uncertainty distribution to more certain predictions, and ID uncertainty attacks can shift the ID uncertainty distribution to less certain predictions.

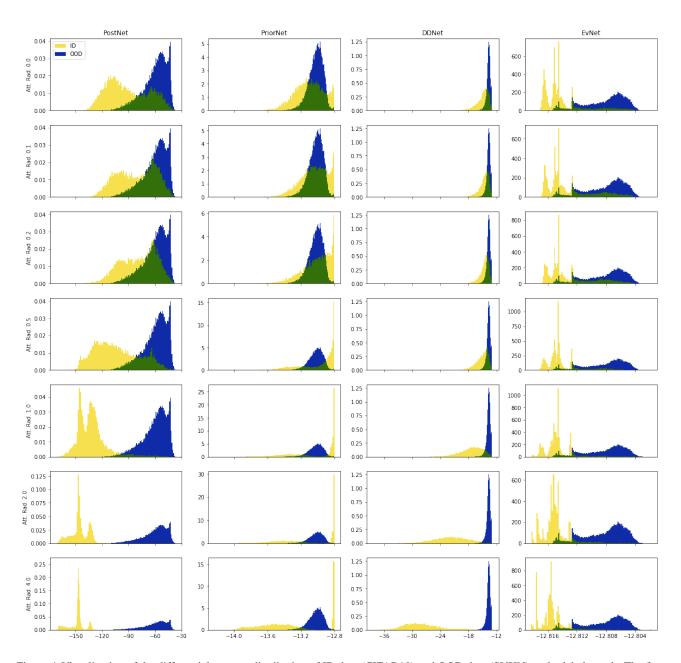


Figure 4. Visualization of the differential entropy distribution of ID data (CIFAR10) and OOD data (SVHN) under label attack. The first row corresponds to no attack. The other rows correspond do increasingly stronger attack strength.

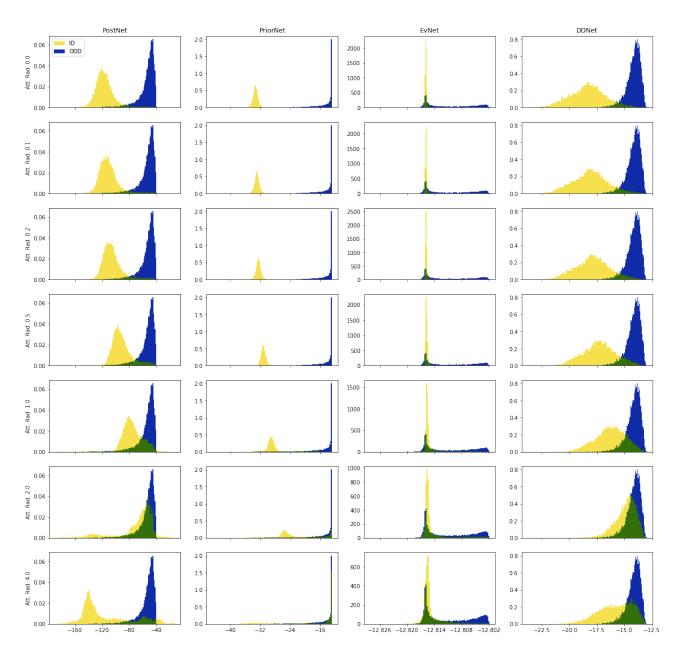


Figure 5. Visualization of the differential entropy distribution of ID data (MNIST) and OOD data (KMNIST) under label attack. The first row corresponds to no attack. The other rows correspond do increasingly stronger attack strength.

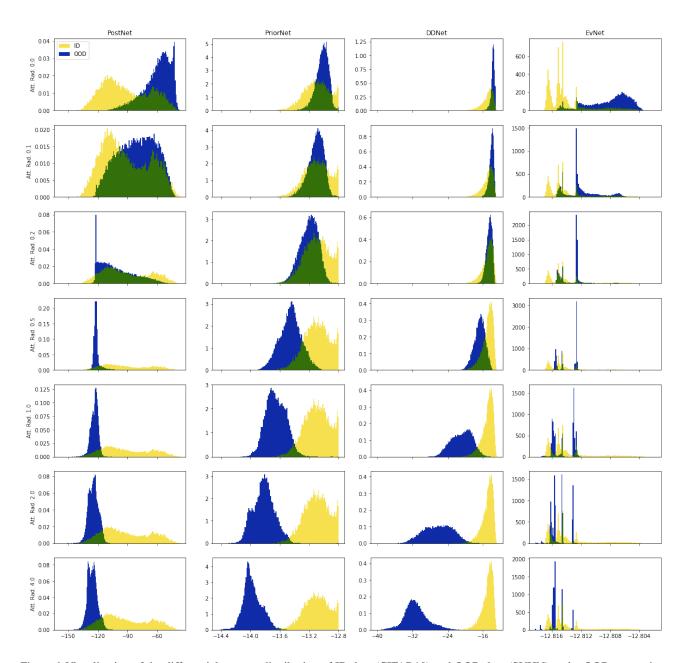


Figure 6. Visualization of the differential entropy distribution of ID data (CIFAR10) and OOD data (SVHN) under OOD uncertainty attack. The first row corresponds to no attack. The other rows correspond do increasingly stronger attack strength.

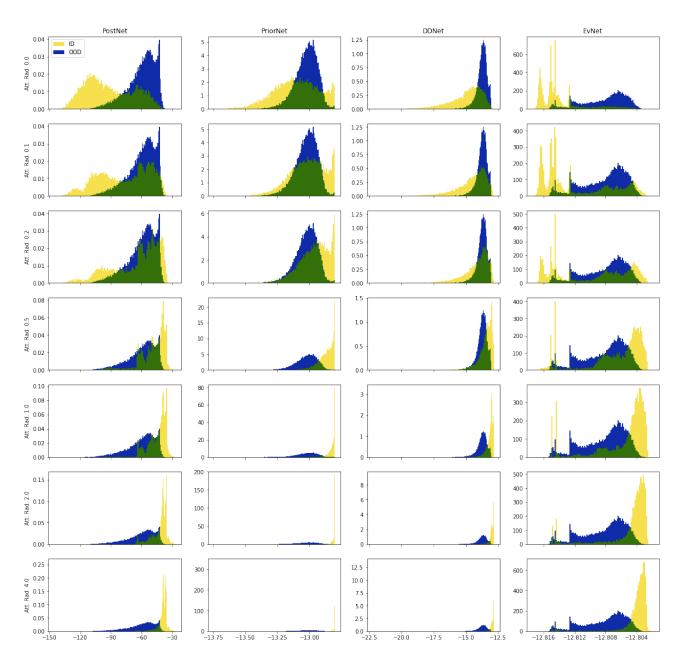


Figure 7. Visualization of the differential entropy distribution of ID data (CIFAR10) and OOD data (SVHN) under ID uncertainty attack. The first row corresponds to no attack. The other rows correspond do increasingly stronger attack strength.

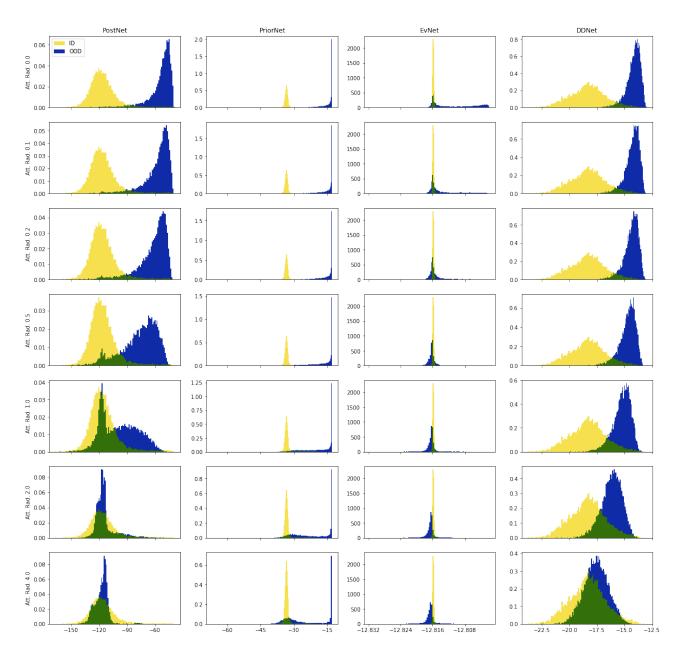


Figure 8. Visualization of the differential entropy distribution of ID data (MNIST) and OOD data (KMNIST) under OOD uncertainty attack. The first row corresponds to no attack. The other rows correspond do increasingly stronger attack strength.

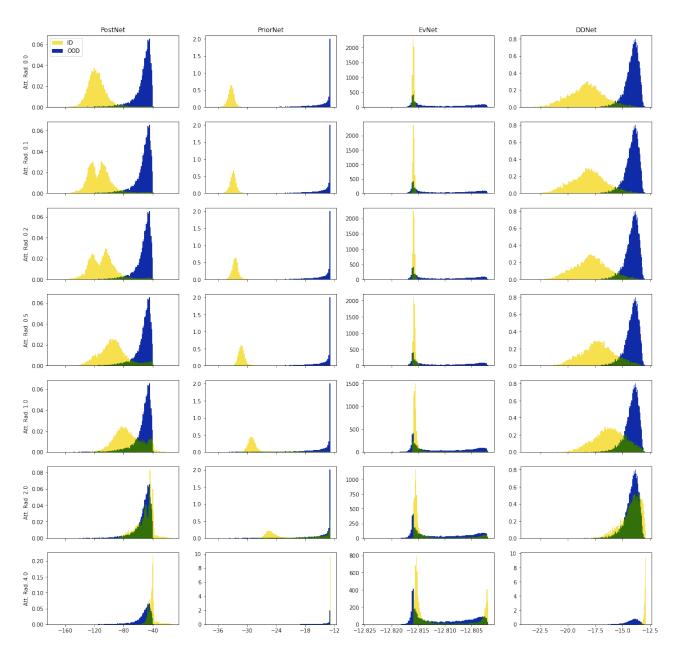


Figure 9. Visualization of the differential entropy distribution of ID data (MNIST) and OOD data (KMNIST) under ID uncertainty attack. The first row corresponds to no attack. The other rows correspond do increasingly stronger attack strength.